
The class File Format

THIS chapter describes the Java virtual machine class file format. Each class file contains the definition of a single class or interface. Although a class or interface need not have an external representation literally contained in a file (for instance, because the class is generated by a class loader), we will colloquially refer to any valid representation of a class or interface as being in the class file format.

A class file consists of a stream of 8-bit bytes. All 16-bit, 32-bit, and 64-bit quantities are constructed by reading in two, four, and eight consecutive 8-bit bytes, respectively. Multibyte data items are always stored in big-endian order, where the high bytes come first. In the Java and Java 2 platforms, this format is supported by interfaces `java.io.DataInput` and `java.io.DataOutput` and classes such as `java.io.DataInputStream` and `java.io.DataOutputStream`.

This chapter defines its own set of data types representing class file data: The types `u1`, `u2`, and `u4` represent an unsigned one-, two-, or four-byte quantity, respectively. In the Java and Java 2 platforms, these types may be read by methods such as `readUnsignedByte`, `readUnsignedShort`, and `readInt` of the interface `java.io.DataInput`.

This chapter presents the class file format using pseudostructures written in a C-like structure notation. To avoid confusion with the fields of classes and class instances, etc., the contents of the structures describing the class file format are referred to as *items*. Successive items are stored in the class file sequentially, without padding or alignment.

Tables, consisting of zero or more variable-sized items, are used in several class file structures. Although we use C-like array syntax to refer to table items, the fact that tables are streams of varying-sized structures means that it is not possible to translate a table index directly to a byte offset into the table.

Where we refer to a data structure as an array, it consists of zero or more contiguous fixed-sized items and can be indexed like an array.

4.1 Notation and Terminology

We use this font for ~~Java virtual machine instructions~~ Prolog code and ~~for class file structures~~ code fragments.

We use this font for Java virtual machine instructions and for class file structures.

Commentary, designed to clarify the specification, is given as indented text between horizontal lines:

Commentary provides intuition, motivation, rationale, examples etc.

4.2 The ClassFile Structure

A class file consists of a single ClassFile structure:

```

ClassFile {
    u4 magic;
    u2 minor_version;
    u2 major_version;
    u2 constant_pool_count;
    cp_info constant_pool[constant_pool_count-1];
    u2 access_flags;
    u2 this_class;
    u2 super_class;
    u2 interfaces_count;
    u2 interfaces[interfaces_count];
    u2 fields_count;
    field_info fields[fields_count];
    u2 methods_count;
    method_info methods[methods_count];
    u2 attributes_count;
    attribute_info attributes[attributes_count];
}

```

The items in the ClassFile structure are as follows:

`magic`

The `magic` item supplies the magic number identifying the class file format; it has the value `0xCAFEBAFE`.

`minor_version, major_version`

The values of the `minor_version` and `major_version` items are the minor and major version numbers of this class file. Together, a major and a minor version number determine the version of the class file format. If a class file has major version number M and minor version number m , we denote the version of its class file format as $M.m$. Thus, class file format versions may be ordered lexicographically, for example, $1.5 < 2.0 < 2.1$.

A Java virtual machine implementation can support a class file format of version v if and only if v lies in some contiguous range $M_i.0 \leq v \leq M_j.m$. Only Sun can specify what range of versions a Java virtual machine implementation conforming to a certain release level of the Java platform may support.¹

constant_pool_count

The value of the `constant_pool_count` item is equal to the number of entries in the `constant_pool` table plus one. A `constant_pool` index is considered valid if it is greater than zero and less than `constant_pool_count`, with the exception for constants of type `long` and `double` noted in §4.5.5.

constant_pool[]

The `constant_pool` is a table of structures (§4.5) representing various string constants, class and interface names, field names, and other constants that are referred to within the `ClassFile` structure and its substructures. The format of each `constant_pool` table entry is indicated by its first “tag” byte.

The `constant_pool` table is indexed from 1 to `constant_pool_count-1`.

access_flags

The value of the `access_flags` item is a mask of flags used to denote access permissions to and properties of this class or interface. The interpretation of each flag, when set, is as shown in Table 4.1.

¹ The Java virtual machine implementation of Sun’s JDK release 1.0.2 supports `class` file format versions 45.0 through 45.3 inclusive. Sun’s JDK releases 1.1.X can support `class` file formats of versions in the range 45.0 through 45.65535 inclusive. For $k \geq 2$ implementations of version 1. k of the Java 2 platform can support `class` file formats of versions in the range 45.0 through $44+k.0$ inclusive.

Table 4.1 Class access and property modifiers

Flag Name	Value	Interpretation
ACC_PUBLIC	0x0001	Declared <code>public</code> ; may be accessed from outside its package.
ACC_FINAL	0x0010	Declared <code>final</code> ; no subclasses allowed.
ACC_SUPER	0x0020	Treat superclass methods specially when invoked by the <i>invokespecial</i> instruction.
ACC_INTERFACE	0x0200	Is an interface, not a class.
ACC_ABSTRACT	0x0400	Declared <code>abstract</code> ; must not be instantiated.
ACC_SYNTHETIC	0x1000	Declared <code>synthetic</code> ; Not present in the source code.
ACC_ANNOTATION	0x2000	Declared as an annotation type.
ACC_ENUM	0x4000	Declared as an enum type.

A class may be marked with the ACC_SYNTHETIC flag to indicate that it was generated by the compiler and does not appear in the source code.

The ACC_ENUM bit ~~is~~ indicates that this class or its superclass is declared as an enumerated type.

An interface is distinguished by its ACC_INTERFACE flag being set. If its ACC_INTERFACE flag is not set, this class file defines a class, not an interface.

If the ACC_INTERFACE flag of this class file is set, its ACC_ABSTRACT flag must also be set (~~§2.13.1~~) ~~and its ACC_PUBLIC flag may be set (§2.13.1)~~. Such a class file must not have ~~any of the other~~ its ACC_FINAL, ACC_SUPER or ACC_ENUM flags in Table 4.1 set.

An annotation type must have its ACC_ANNOTATION flag set. If the ACC_ANNOTATION flag is set, the ACC_INTERFACE flag must be set as well.

If the ACC_INTERFACE flag of this class file is not set, it may have any of the other flags in Table 4.1 set, except the ACC_ANNOTATION flag. However, such a class file cannot

have both its ACC_FINAL and ACC_ABSTRACT flags set
 (§2.8.2)(§2.8.2).

The setting of the ACC_SUPER flag indicates which of two alternative semantics for its *invokespecial* instruction the Java virtual machine is to express; the ACC_SUPER flag exists for backward compatibility for code compiled by Sun's older compilers for the Java programming language. All new implementations of the Java virtual machine should implement the semantics for *invokespecial* documented in this specification. All new compilers to the instruction set of the Java virtual machine should set the ACC_SUPER flag. Sun's older compilers generated `ClassFile` flags with ACC_SUPER unset. Sun's older Java virtual machine implementations ignore the flag if it is set.

All bits of the `access_flags` item not assigned in Table 4.1 are reserved for future use. They should be set to zero in generated `class` files and should be ignored by Java virtual machine implementations.

`this_class`

The value of the `this_class` item must be a valid index into the `constant_pool` table. The `constant_pool` entry at that index must be a `CONSTANT_Class_info` (§4.5.1) structure representing the class or interface defined by this `class` file.

`super_class`

For a class, the value of the `super_class` item either must be zero or must be a valid index into the `constant_pool` table. If the value of the `super_class` item is nonzero, the `constant_pool` entry at that index must be a `CONSTANT_Class_info` (§4.5.1) structure representing the direct superclass of the class defined by this `class` file. Neither the direct superclass nor any of its superclasses may be a `final` class.

If the value of the `super_class` item is zero, then this `class` file must represent the class `Object`, the only class or interface without a direct superclass.

For an interface, the value of the `super_class` item must always be a valid index into the `constant_pool` table. The `constant_pool` entry at that index must be a `CONSTANT_Class_info` structure representing the class `Object`.

interfaces_count

The value of the `interfaces_count` item gives the number of direct superinterfaces of this class or interface type.

interfaces[]

Each value in the `interfaces` array must be a valid index into the `constant_pool` table. The `constant_pool` entry at each value of `interfaces[i]`, where $0 \leq i < \text{interfaces_count}$, must be a `CONSTANT_Class_info` (§4.5.1) structure representing an interface that is a direct superinterface of this class or interface type, in the left-to-right order given in the source for the type.

fields_count

The value of the `fields_count` item gives the number of `field_info` structures in the `fields` table. The `field_info` (§4.6) structures represent all fields, both class variables and instance variables, declared by this class or interface type.

fields[]

Each value in the `fields` table must be a `field_info` (§4.6) structure giving a complete description of a field in this class or interface. The `fields` table includes only those fields that are declared by this class or interface. It does not include items representing fields that are inherited from superclasses or superinterfaces.

methods_count

The value of the `methods_count` item gives the number of `method_info` structures in the `methods` table.

methods[]

Each value in the `methods` table must be a `method_info` (§4.7) structure giving a complete description of a method in this class or interface. If the method is not `native` or `abstract`, the Java virtual machine instructions implementing the method are also supplied.

The `method_info` structures represent all methods declared by this class or interface type, including instance methods, class (`static`) methods, instance initialization methods (§3.9)(§3.9), and any class or interface initialization method (§3.9)(§3.9). The

methods table does not include items representing methods that are inherited from superclasses or superinterfaces.

`attributes_count`

The value of the `attributes_count` item gives the number of attributes (§4.8) in the `attributes` table of this class.

`attributes[]`

Each value of the `attributes` table must be an ~~attribute-~~
[attribute_info](#) structure (§4.8).

The only attributes defined by this specification as appearing in the `attributes` table of a `ClassFile` structure are the `InnerClasses` (§4.8.6), `EnclosingMethod` (§4.8.7), `Synthetic` (§4.8.8), `SourceFile` (§4.8.10), `Signature`, and `Deprecated` (§4.8.15) attributes.

A Java virtual machine implementation is required to silently ignore any or all attributes in the `attributes` table of a `ClassFile` structure that it does not recognize. Attributes not defined in this specification are not allowed to affect the semantics of the `class` file, but only to provide additional descriptive information (§4.8.1).

4.3 The Internal Form of Names

4.3.1 Fully Qualified Class and Interface Names

Class and interface names that appear in `class` file structures are always represented in a fully qualified form (~~§2.7.5~~)([§2.7.5](#)). Such names are always represented as `CONSTANT_Utf8_info` (§4.5.7) structures and thus may be drawn, where not further constrained, from the entire Unicode character set. Class names and interfaces are referenced both from those `CONSTANT_NameAndType_info` (§4.5.6) structures that have such names as part of their descriptor (§4.4) and from all `CONSTANT_Class_info` (§4.5.1) structures.

For historical reasons the syntax of fully qualified class and interface names that appear in `class` file structures differs from the familiar syntax of fully qualified names documented in ~~§2.7.5~~[§2.7.5](#). In this internal form, the ASCII periods ('.') that normally separate the identifiers that make up the fully qualified name are replaced by ASCII forward slashes ('/'). The identifiers themselves must be unqualified names as discussed in section (§4.3.2) below. For example, the normal

fully qualified name of class Thread is `java.lang.Thread`. In the form used in descriptors in the class file format, a reference to the name of class Thread is implemented using a `CONSTANT_Utf8_info` structure representing the string `"java/lang/Thread"`.

4.3.2 Unqualified Names

Names of methods, fields and local variables are stored as *unqualified names*. Unqualified names must not contain the characters `'.'`, `','`, `'['` or `'/'`. Method names are further constrained so that, with the exception of the special method names [\(§3.9\)](#) [\(§3.9\)](#) `<init>` and `<clinit>`, they must not contain the characters `'<'` or `'>'`.

4.4 Descriptors and Signatures

A *descriptor* is a string representing the type of a field or method. Descriptors are represented in the class file format using modified UTF-8 strings (§4.5.7) and thus may be drawn, where not further constrained, from the entire Unicode character set.

A *signature* is a string representing the generic type of a field or method, or generic type information for a class declaration.

4.4.1 Grammar Notation

Descriptors and signatures are specified using a grammar. This grammar is a set of productions that describe how sequences of characters can form syntactically correct descriptors of various types. Terminal symbols of the grammar are shown in **bold fixed-width** font. Nonterminal symbols are shown in *italic* type. The definition of a nonterminal is introduced by the name of the nonterminal being defined, followed by a colon. One or more alternative right-hand sides for the nonterminal then follow on succeeding lines. For example, the production:

FieldType:

BaseType
ObjectType
ArrayType

states that a *FieldType* may represent either a *BaseType*, an *ObjectType*, or an *ArrayType*.

A nonterminal symbol on the right-hand side of a production that is followed by an asterisk (*) represents zero or more possibly different values produced from that nonterminal, appended without any intervening space. Similarly, a nonterminal symbol on the right-hand side of a production that is followed by a plus sign (+) represents one or more possibly different values produced from that nonterminal, appended without any intervening space. The production:

MethodDescriptor:

(*ParameterDescriptor**) *ReturnDescriptor*

states that a *MethodDescriptor* represents a left parenthesis, followed by zero or more *ParameterDescriptor* values, followed by a right parenthesis, followed by a *ReturnDescriptor*.

4.4.2 Field Descriptors

A *field descriptor* represents the type of a class, instance, or local variable. It is a series of characters generated by the grammar:

FieldDescriptor:

FieldType

ComponentType:

FieldType

FieldType:

BaseType

ObjectType

ArrayType

BaseType:

B

C

D

F

I

J**S****Z***ObjectType*:**L Classname;***ArrayType*:**[ComponentType**

The characters of *BaseType*, the **L** and **;** of *ObjectType*, and the **[** of *ArrayType* are all ASCII characters. The **C**lassname represents a fully qualified class or interface name. For historical reasons it is encoded in internal form (§4.2). A type descriptor representing an array type is valid only if it represents a type with 255 or fewer dimensions.

The interpretation of the field types is as shown in Table 4.2.

Table 4.2 Interpretation of *BaseType* characters

<i>BaseType</i> Character	Type	Interpretation
B	byte	signed byte
C	char	Unicode character
D	double	double-precision floating-point value
F	float	single-precision floating-point value
I	int	integer
J	long	long integer
L Classname;	reference	an instance of class <classname>
S	short	signed short
Z	boolean	true or false
[reference	one array dimension

For example, the descriptor of an instance variable of type `int` is simply **I**. The descriptor of an instance variable of type `Object` is **Ljava/lang/Object;**. Note that the internal form of the fully qualified name for class `Object` is used. The descriptor of an instance variable that is a multidimensional `double` array,

```
double d[][][];

is

[[[D
```

4.4.3 Method Descriptors

A *method descriptor* represents the parameters that the method takes and the value that it returns:

MethodDescriptor:

(ParameterDescriptor) ReturnDescriptor*

A *parameter descriptor* represents a parameter passed to a method:

ParameterDescriptor:

FieldType

A *return descriptor* represents the type of the value returned from a method. It is a series of characters generated by the grammar:

ReturnDescriptor:

FieldType

VoidDescriptor

VoidDescriptor:

V

The character **V** indicates that the method returns no value (its return type is `void`).

A method descriptor is valid only if it represents method parameters with a total length of 255 or less, where that length includes the contribution for `this` in the case of instance or interface method invocations. The total length is calculated by summing the contributions of the individual parameters, where a parameter of type `long` or `double` contributes two units to the length and a parameter of any other type contributes one unit.

For example, the method descriptor for the method

```
Object mymethod(int i, double d, Thread t)
```

is

(IDLjava/lang/Thread;)Ljava/lang/Object;

Note that internal forms of the fully qualified names of `Thread` and `Object` are used in the method descriptor.

The method descriptor for `mymethod` is the same whether `mymethod` is a class or an instance method. Although an instance method is passed `this`, a reference to the current class instance, in addition to its intended parameters, that fact is not reflected in the method descriptor. (A reference to `this` is not passed to a class method.) The reference to `this` is passed implicitly by the method invocation instructions of the Java virtual machine used to invoke instance methods.

4.4.4 Signatures

Signatures are used to encode Java programming language type information that is not part of the Java virtual machine type system, such as generic type and method declarations and parameterized types. See *The Java Language Specification, Third Edition*, for details about such types.

This kind of type information is needed to support reflection and debugging, and by the Java compiler.

In the following, the terminal symbol **Identifier** is used to denote an identifier for a type, field, local variable, parameter, method name or type variable, as generated by the Java compiler. Such an identifier may contain characters that must not appear in a legal identifier in the Java programming language.

ClassSignature:

*FormalTypeParametersopt SuperclassSignature SuperinterfaceSignature**

A class signature, defined by the production *ClassSignature* above, is used to encode type information about a class declaration. It ~~gives the fully qualified name of the class,~~ describes any formal type parameters ~~it the class~~ might have, and lists its (possibly parameterized) direct superclass and direct superinterfaces, if any.

FormalTypeParameters:

<FormalTypeParameter+>

FormalTypeParameter:

Identifier *ClassBound InterfaceBound**

A formal type parameter is described by its name, followed by its class and interface bounds. If the class bound does not specify a type, it is taken to be Object.

ClassBound:

: *FieldTypeSignature**opt*

InterfaceBound:

: *FieldTypeSignature*

SuperclassSignature:

ClassTypeSignature

SuperinterfaceSignature:

ClassTypeSignature

FieldTypeSignature:

ClassTypeSignature

ArrayTypeSignature

TypeVariableSignature

A field type signature, defined by the production *FieldTypeSignature* above, encodes the (possibly parameterized) type for a field, parameter or local variable.

ClassTypeSignature:

L *PackageSpecifier** *SimpleClassTypeSignature*

*ClassTypeSignatureSuffix** ;

PackageSpecifier:

Identifier / *PackageSpecifier**

SimpleClassTypeSignature:

Identifier *TypeArguments*_{opt}

ClassTypeSignatureSuffix:

. *SimpleClassTypeSignature*

TypeVariableSignature:

T ~~**Identifier**~~ **Identifier** ;

TypeArguments:

<*TypeArgument*+>

TypeArgument:

*WildcardIndicator*_{opt} *FieldTypeSignature*

*

WildcardIndicator:

+

-

ArrayTypeSignature:

[*TypeSignature*

TypeSignature:

F*FieldTypeSignature*

F*BaseType*

A class type signature gives complete type information for a class or interface type. The class type signature must be formulated such that it can be reliably

mapped to the binary name of the class it denotes by erasing any type arguments and converting ‘.’ characters in the signature to ‘\$’ characters.

MethodTypeSignature:

*FormalTypeParameters*_{opt} (*TypeSignature**) *ReturnType*
*ThrowsSignature**

ReturnType:

TypeSignature

VoidDescriptor

ThrowsSignature:

^*ClassTypeSignature*

^*TypeVariableSignature*

A method signature, defined by the production *MethodTypeSignature* above, encodes the (possibly parameterized) types of the method’s formal arguments and of the exceptions it has declared in its throws clause, its (possibly parameterized) return type, and any formal type parameters in the method declaration.

A Java compiler must output generic signature information for any class, interface, constructor or member whose generic signature would include references to type variables or parameterized types. If the throws clause of a method or constructor does not involve type variables, the *ThrowsSignature* may be elided from the *MethodTypeSignature*.

The Java virtual machine does not check the well formedness of the signatures described in this subsection during loading or linking. Instead, these checks are deferred until the signatures are used by reflective methods, as specified in the API of `Class` and members of `java.lang.reflect`. Future versions of the Java virtual machine may be required to perform some or all of these checks during loading or linking.

4.5 The Constant Pool

Java virtual machine instructions do not rely on the runtime layout of classes, interfaces, class instances, or arrays. Instead, instructions refer to symbolic information in the `constant_pool` table.

All `constant_pool` table entries have the following general format:

```
cp_info {
    u1 tag;
    u1 info[];
}
```

Each item in the `constant_pool` table must begin with a 1-byte tag indicating the kind of `cp_info` entry. The contents of the `info` array vary with the value of `tag`. The valid tags and their values are listed in Table 4.3. Each tag byte must be followed by two or more bytes giving information about the specific constant. The format of the additional information varies with the tag value.

Table 4.3 Constant pool tags

Constant Type	Value
CONSTANT_Class	7
CONSTANT_Fieldref	9
CONSTANT_Methodref	10
CONSTANT_InterfaceMethodref	11
CONSTANT_String	8
CONSTANT_Integer	3
CONSTANT_Float	4
CONSTANT_Long	5
CONSTANT_Double	6
CONSTANT_NameAndType	12
CONSTANT_Utf8	1

4.5.1 The `CONSTANT_Class_info` Structure

The `CONSTANT_Class_info` structure is used to represent a class or an interface:

```
CONSTANT_Class_info {
    u1 tag;
    u2 name_index;
}
```

The items of the `CONSTANT_Class_info` structure are the following:

tag

The tag item has the value `CONSTANT_Class` (7).

name_index

The value of the name_index item must be a valid index into the constant_pool table. The constant_pool entry at that index must be a `CONSTANT_Utf8_info` (§4.5.7) structure representing a valid fully qualified class or interface name encoded in internal form (§4.3.1).

Because arrays are objects, the opcodes *anewarray* and *multianewarray* can reference array “classes” via `CONSTANT_Class_info` (§4.5.1) structures in the constant_pool table. For such array classes, the name of the class is the descriptor of the array type. For example, the class name representing a two-dimensional int array type

```
int[][]
```

is

```
[[I
```

The class name representing the type array of class Thread

```
Thread[]
```

is

```
[Ljava/lang/Thread;
```

An array type descriptor is valid only if it represents 255 or fewer dimensions.

4.5.2 The `CONSTANT_Fieldref_info`, `CONSTANT_Methodref_info`, and `CONSTANT_InterfaceMethodref_info` Structures

Fields, methods, and interface methods are represented by similar structures:

```
CONSTANT_Fieldref_info {
    u1 tag;
    u2 class_index;
    u2 name_and_type_index;
}
```

```

CONSTANT_Methodref_info {
    u1 tag;
    u2 class_index;
    u2 name_and_type_index;
}

CONSTANT_InterfaceMethodref_info {
    u1 tag;
    u2 class_index;
    u2 name_and_type_index;
}

```

The items of these structures are as follows:

tag

The tag item of a `CONSTANT_Fieldref_info` structure has the value `CONSTANT_Fieldref` (9).

The tag item of a `CONSTANT_Methodref_info` structure has the value `CONSTANT_Methodref` (10).

The tag item of a `CONSTANT_InterfaceMethodref_info` structure has the value `CONSTANT_InterfaceMethodref` (11).

class_index

The value of the `class_index` item must be a valid index into the `constant_pool` table. The `constant_pool` entry at that index must be a `CONSTANT_Class_info` (§4.5.1) structure representing a class or interface type that has the field or method as a member.

The `class_index` item of a `CONSTANT_Methodref_info` structure must be a class type, not an interface type. The `class_index` item of a `CONSTANT_InterfaceMethodref_info` structure must be an interface type. The `class_index` item of a `CONSTANT_Fieldref_info` structure may be either a class type or an interface type.

name_and_type_index

The value of the `name_and_type_index` item must be a valid index into the `constant_pool` table. The `constant_pool` entry at that index must be a `CONSTANT_NameAndType_info` (§4.5.6) structure. This `constant_pool` entry indicates the name and descriptor of the field or method. In a `CONSTANT_Fieldref_info` the indicated descriptor must be a field descriptor (§4.4.2).

Otherwise, the indicated descriptor must be a method descriptor (§4.4.3).

If the name of the method of a `CONSTANT_Methodref_info` structure begins with a '`<`' ('`\u003c`'), then the name must be the special name `<init>`, representing an instance initialization method (§3.9)(§3.9). The return type of such a method must be `void`.

4.5.3 The `CONSTANT_String_info` Structure

The `CONSTANT_String_info` structure is used to represent constant objects of the type `String`:

```
CONSTANT_String_info {
    u1 tag;
    u2 string_index;
}
```

The items of the `CONSTANT_String_info` structure are as follows:

`tag`

The `tag` item of the `CONSTANT_String_info` structure has the value `CONSTANT_String` (8).

`string_index`

The value of the `string_index` item must be a valid index into the `constant_pool` table. The `constant_pool` entry at that index must be a `CONSTANT_Utf8_info` (§4.5.7) structure representing the sequence of characters to which the `String` object is to be initialized.

4.5.4 The `CONSTANT_Integer_info` and `CONSTANT_Float_info` Structures

The `CONSTANT_Integer_info` and `CONSTANT_Float_info` structures represent 4-byte numeric (`int` and `float`) constants:

```
CONSTANT_Integer_info {
    u1 tag;
    u4 bytes;
}
```

```

CONSTANT_Float_info {
    u1 tag;
    u4 bytes;
}

```

The items of these structures are as follows:

tag

The tag item of the `CONSTANT_Integer_info` structure has the value `CONSTANT_Integer` (3).

The tag item of the `CONSTANT_Float_info` structure has the value `CONSTANT_Float` (4).

bytes

The bytes item of the `CONSTANT_Integer_info` structure represents the value of the `int` constant. The bytes of the value are stored in big-endian (high byte first) order.

The bytes item of the `CONSTANT_Float_info` structure represents the value of the `float` constant in IEEE 754 floating-point single format (~~§3.3.2~~)(§3.3.2). The bytes of the single format representation are stored in big-endian (high byte first) order.

The value represented by the `CONSTANT_Float_info` structure is determined as follows. The bytes of the value are first converted into an `int` constant *bits*. Then:

- If *bits* is `0x7f800000`, the `float` value will be positive infinity.
- If *bits* is `0xff800000`, the `float` value will be negative infinity.
- If *bits* is in the range `0x7f800001` through `0x7fffffff` or in the range `0xff800001` through `0xffffffff`, the `float` value will be NaN.
- In all other cases, let *s*, *e*, and *m* be three values that might be computed from *bits*:

```

int s = ((bits >> 31) == 0) ? 1 : -1;
int e = ((bits >> 23) & 0xff);
int m = (e == 0) ?
        (bits & 0x7fffffff) << 1 :
        (bits & 0x7fffffff) | 0x800000;

```

Then the float value equals the result of the mathematical expression $s \cdot m \cdot 2^{e-150}$.

4.5.5 The CONSTANT_Long_info and CONSTANT_Double_info Structures

The CONSTANT_Long_info and CONSTANT_Double_info represent 8-byte numeric (long and double) constants:

```

CONSTANT_Long_info {
    u1 tag;
    u4 high_bytes;
    u4 low_bytes;
}

CONSTANT_Double_info {
    u1 tag;
    u4 high_bytes;
    u4 low_bytes;
}

```

All 8-byte constants take up two entries in the constant_pool table of the class file. If a CONSTANT_Long_info or CONSTANT_Double_info structure is the item in the constant_pool table at index n , then the next usable item in the pool is located at index $n+2$. The constant_pool index $n+1$ must be valid but is considered unusable.²

The items of these structures are as follows:

² In retrospect, making 8-byte constants take two constant pool entries was a poor choice.

tag

The `tag` item of the `CONSTANT_Long_info` structure has the value `CONSTANT_Long` (5).

The `tag` item of the `CONSTANT_Double_info` structure has the value `CONSTANT_Double` (6).

high_bytes, low_bytes

The unsigned `high_bytes` and `low_bytes` items of the `CONSTANT_Long_info` structure together represent the value of the long constant $((\text{long}) \text{high_bytes} \ll 32) + \text{low_bytes}$, where the bytes of each of `high_bytes` and `low_bytes` are stored in big-endian (high byte first) order.

The `high_bytes` and `low_bytes` items of the `CONSTANT_Double_info` structure together represent the double value in IEEE 754 floating-point double format (§3.3.2)(§3.3.2). The bytes of each item are stored in big-endian (high byte first) order.

The value represented by the `CONSTANT_Double_info` structure is determined as follows. The `high_bytes` and `low_bytes` items are first converted into the long constant *bits*, which is equal to $((\text{long}) \text{high_bytes} \ll 32) + \text{low_bytes}$. Then:

- If *bits* is `0x7ff0000000000000L`, the double value will be positive infinity.
- If *bits* is `0xfff0000000000000L`, the double value will be negative infinity.
- If *bits* is in the range `0x7ff0000000000001L` through `0x7fffffffffffffffffL` or in the range `0xfff0000000000001L` through `0xffffffffffffffffL`, the double value will be NaN.
- In all other cases, let *s*, *e*, and *m* be three values that might be computed from *bits*:

```
int s = ((bits >> 63) == 0) ? 1 : -1;
int e = (int)((bits >> 52) & 0x7ffL);
long m = (e == 0) ?
    (bits & 0xffffffffffffffffL) << 1 :
    (bits & 0xffffffffffffffffL) | 0x100000000000000L;
```

Then the floating-point value equals the `double` value of the mathematical expression $s \cdot m \cdot 2^{e - 1075}$.

4.5.6 The `CONSTANT_NameAndType_info` Structure

The `CONSTANT_NameAndType_info` structure is used to represent a field or method, without indicating which class or interface type it belongs to:

```
CONSTANT_NameAndType_info {
    u1 tag;
    u2 name_index;
    u2 descriptor_index;
}
```

The items of the `CONSTANT_NameAndType_info` structure are as follows:

`tag`

The `tag` item of the `CONSTANT_NameAndType_info` structure has the value `CONSTANT_NameAndType` (12).

`name_index`

The value of the `name_index` item must be a valid index into the `constant_pool` table. The `constant_pool` entry at that index must be a `CONSTANT_Utf8_info` (§4.5.7) structure representing either the special method name `<init>` (§3.9)(§3.9) or a valid unqualified name (§4.3.2) denoting a field or method. .

`descriptor_index`


The value of the `descriptor_index` item must be a valid index into the `constant_pool` table. The `constant_pool` entry at that index must be a `CONSTANT_Utf8_info` (§4.5.7) structure representing a valid field descriptor (§4.4.2) or method descriptor (§4.4.3).

4.5.7 The `CONSTANT_Utf8_info` Structure

The `CONSTANT_Utf8_info` structure is used to represent constant string values. String content is encoded in *modified UTF-8*.

Modified UTF-8 strings are encoded so that character sequences that contain only non-null ASCII characters can be represented using only 1 byte per character, but all Unicode characters can be represented. All

characters in the range `'\u0001'` to `'\u007F'` are represented by a single byte:




The 7 bits of data in the byte give the value of the character represented. The null character (`'\u0000'`) and characters in the range `'\u0080'` to `'\u07FF'` are represented by a pair of bytes *x* and *y*:





The bytes represent the character with the value $((x \& 0x1f) \ll 6) + (y \& 0x3f)$.

Characters in the range `'\u0800'` to `'\uFFFF'` are represented by 3 bytes *x*, *y*, and *z*:



The character with the value $((x \& 0xf) \ll 12) + ((y \& 0x3f) \ll 6) + (z \& 0x3f)$ is represented by the bytes.

Characters with code points above U+FFFF (so-called *supplementary characters*) are represented by separately encoding the two surrogate code units of their UTF-16 representation. Each of the surrogate code units is represented by three bytes. This means, supplementary characters are represented by six bytes, *u*, *v*, *w*, *x*, *y*, and *z*:



The character with the value $0x10000 + ((v \& 0x0f) \ll 16) + ((w \& 0x3f) \ll 10) + (y \& 0x0f) \ll 6 + (z \& 0x3f)$ is represented by the six bytes.

The bytes of multibyte characters are stored in the `class` file in big-endian (high byte first) order.

There are two differences between this format and the “standard” UTF-8 format. First, the null character (`(char)0`) is encoded using the 2-byte format rather than the 1-byte format, so that modified UTF-8 strings never have embedded nulls. Second, only the 1-byte, 2-byte, and 3-byte formats of standard UTF-8 are used.

The Java VM does not recognize the four-byte format of standard UTF-8; it uses its own two-times-three-byte format instead.

For more information regarding the standard UTF-8 format, see section 3.9 *Unicode Encoding Forms* of *The Unicode Standard, Version 4.0*.

The `CONSTANT_Utf8_info` structure is

```
CONSTANT_Utf8_info {  
    u1 tag;  
    u2 length;  
    u1 bytes[length];  
}
```

The items of the `CONSTANT_Utf8_info` structure are the following:

`tag`

The `tag` item of the `CONSTANT_Utf8_info` structure has the value `CONSTANT_Utf8 (1)`.

`length`

The value of the `length` item gives the number of bytes in the bytes array (not the length of the resulting string). The strings in the `CONSTANT_Utf8_info` structure are not null-terminated.

`bytes[]`

The bytes array contains the bytes of the string. No byte may have the value (byte)0 or lie in the range (byte)0xf0-(byte)0xff.

4.6 Fields

Each field is described by a `field_info` structure. No two fields in one class file may have the same name and descriptor (§4.4.2). The format of this structure is

```
field_info {  
    u2 access_flags;  
    u2 name_index;  
    u2 descriptor_index;  
    u2 attributes_count;  
    attribute_info attributes[attributes_count];  
}
```

```
}
```

The items of the `field_info` structure are as follows:

`access_flags`

The value of the `access_flags` item is a mask of flags used to denote access permission to and properties of this field. The interpretation of each flag, when set, is as shown in Table 4.4.

Table 4.4 Field access and property flags

Flag Name	Value	Interpretation
ACC_PUBLIC	0x0001	Declared <code>public</code> ; may be accessed from outside its package.
ACC_PRIVATE	0x0002	Declared <code>private</code> ; usable only within the defining class.
ACC_PROTECTED	0x0004	Declared <code>protected</code> ; may be accessed within subclasses.
ACC_STATIC	0x0008	Declared <code>static</code> .
ACC_FINAL	0x0010	Declared <code>final</code> ; no further assignment after initialization.
ACC_VOLATILE	0x0040	Declared <code>volatile</code> ; cannot be cached.
ACC_TRANSIENT	0x0080	Declared <code>transient</code> ; not written or read by a persistent object manager.
ACC_SYNTHETIC	0x1000	Declared <code>synthetic</code> ; Not present in the source code.
ACC_ENUM	0x4000	Declared as an element of an <code>enum</code> .

The ACC_ENUM bit indicates that this field is being used to hold an element of an enumerated type.

A field may be marked with the ACC_SYNTHETIC flag to indicate that it was generated by the compiler and does not appear in the source code.

Fields of classes may set any of the flags in Table 4.4. However, a specific field of a class may have at most one of its ACC_PRIVATE, ACC_PROTECTED, and ACC_PUBLIC flags set (§2.7.4)(§2.7.4) and must not have both its ACC_FINAL and ACC_VOLATILE flags set (§2.9.1)(§2.9.1).

All fields of interfaces must have their ACC_PUBLIC, ACC_STATIC, and ACC_FINAL flags set; they may have their ACC_SYNTHETIC flag set and must not have any of the other flags in Table 4.4 set (§2.13.3.1)(§2.13.3.1).

All bits of the `access_flags` item not assigned in Table 4.4 are reserved for future use. They should be set to zero in generated `class` files and should be ignored by Java virtual machine implementations.

name_index

The value of the `name_index` item must be a valid index into the `constant_pool` table. The `constant_pool` entry at that index must be a `CONSTANT_Utf8_info` (§4.5.7) structure which must represent a valid unqualified name (§4.3.2) denoting a field.

descriptor_index

The value of the `descriptor_index` item must be a valid index into the `constant_pool` table. The `constant_pool` entry at that index must be a `CONSTANT_Utf8_info` (§4.5.7) structure that must represent a valid field descriptor (§4.4.2).

attributes_count

The value of the `attributes_count` item indicates the number of additional attributes (§4.8) of this field.

attributes[]

Each value of the `attributes` table must be an attribute structure (§4.8). A field can have any number of attributes associated with it.

The attributes defined by this specification as appearing in the `attributes` table of a `field_info` structure are the `ConstantValue` (§4.8.2), `Synthetic` (§4.8.8), `Signature` (§4.8.9) and `Deprecated` (§4.8.15) attributes.

A Java virtual machine implementation must recognize and correctly read `ConstantValue` (§4.8.2) attributes found in the `attributes` table of a `field_info` structure. If a Java virtual machine recognizes class files whose major version is 49.0 or above, it must recognize and correctly read `Signature` (§4.8.9) attributes found in the `attributes` table of a `field_info` structure. A Java virtual machine implementation is required to silently ignore any or all other attributes in the `attributes` table that it does not recognize. Attributes not defined in this specification are not allowed to affect the semantics of the `class` file, but only to provide additional descriptive information (§4.8.1).

4.7 Methods

Each method, including each instance initialization method (§3.9)(§3.9) and the class or interface initialization method (§3.9)(§3.9), is described by a `method_info`

structure. No two methods in one class file may have the same name and descriptor (§4.4.3).

The structure has the following format:

```
method_info {  
    u2 access_flags;  
    u2 name_index;  
    u2 descriptor_index;  
    u2 attributes_count;  
    attribute_info attributes[attributes_count];  
}
```

The items of the method_info structure are as follows:

access_flags

The value of the `access_flags` item is a mask of flags used to denote access permission to and properties of this method. The interpretation of each flag, when set, is as shown in Table 4.5.

Table 4.5 Method access and property flags

Flag Name	Value	Interpretation
ACC_PUBLIC	0x0001	Declared <code>public</code> ; may be accessed from outside its package.
ACC_PRIVATE	0x0002	Declared <code>private</code> ; accessible only within the defining class.
ACC_PROTECTED	0x0004	Declared <code>protected</code> ; may be accessed within subclasses.
ACC_STATIC	0x0008	Declared <code>static</code> .
ACC_FINAL	0x0010	Declared <code>final</code> ; must not be overridden.
ACC_SYNCHRONIZED	0x0020	Declared <code>synchronized</code> ; invocation is wrapped in a monitor lock.
ACC_BRIDGE	0x0040	A bridge method, generated by the compiler.
ACC_VARARGS	0x0080	Declared with variable number of arguments.
ACC_NATIVE	0x0100	Declared <code>native</code> ; implemented in a language other than Java.
ACC_ABSTRACT	0x0400	Declared <code>abstract</code> ; no implementation is provided.
ACC_STRICT	0x0800	Declared <code>strictfp</code> ; floating-point mode is FP-strict
ACC_SYNTHETIC	0x1000	Declared <code>synthetic</code> ; Not present in the source code.

The `ACC_VARARGS` flag indicates that this method takes a variable number of arguments at the source code level. A method declared to take a variable number of arguments must be compiled with the `ACC_VARARGS` flag set to 1. All other methods must be compiled with the `ACC_VARARGS` flag set to

0. The ACC_BRIDGE method is used to indicate a bridge method generated by the compiler.

A method may be marked with the ACC_SYNTHETIC flag to indicate that it was generated by the compiler and does not appear in the source code.

Methods of classes may set any of the flags in Table 4.5.

However, a specific method of a class may have at most one of its ACC_PRIVATE, ACC_PROTECTED, and ACC_PUBLIC flags set (§2.7.4)(§2.7.4). If such a method has its ACC_ABSTRACT flag set it must not have any of its ACC_FINAL, ACC_NATIVE, ACC_PRIVATE, ACC_STATIC, ACC_STRICT, or ACC_SYNCHRONIZED flags set (§2.13.3.2)(§2.13.3.2).

All interface methods must have their ACC_ABSTRACT and ACC_PUBLIC flags set; they may have their ACC_VARARGS, ACC_BRIDGE and ACC_SYNTHETIC flags set and must not have any of the other flags in Table 4.5 set (§2.13.3.2)(§2.13.3.2).

A specific instance initialization method (§3.9)(§3.9) may have at most one of its ACC_PRIVATE, ACC_PROTECTED, and ACC_PUBLIC flags set and may also have its ACC_STRICT, ACC_VARARGS, and ACC_SYNTHETIC flags set, but must not have any of the other flags in Table 4.5 set.

Class and interface initialization methods (§3.9)(§3.9) are called implicitly by the Java virtual machine; the value of their access_flags item is ignored except for the settings of the ACC_STRICT flag.

All bits of the access_flags item not assigned in Table 4.5 are reserved for future use. They should be set to zero in generated class files and should be ignored by Java virtual machine implementations.

name_index

The value of the name_index item must be a valid index into the constant_pool table. The constant_pool entry at that index must be a CONSTANT_Utf8_info (§4.5.7) structure representing either one of the special method names (§3.9)(§3.9), <init> or <clinit>, or a valid unqualified name (§4.3.2) denoting a method.

descriptor_index

The value of the `descriptor_index` item must be a valid index into the `constant_pool` table. The `constant_pool` entry at that index must be a `CONSTANT_Utf8_info` (§4.5.7) structure representing a valid method descriptor (§4.4.3).

attributes_count

The value of the `attributes_count` item indicates the number of additional attributes (§4.8) of this method.

attributes[]

Each value of the `attributes` table must be an attribute structure (§4.8). A method can have any number of optional attributes associated with it.

The only attributes defined by this specification as appearing in the `attributes` table of a `method_info` structure are the `Code` (§4.8.3), `Exceptions` (§4.8.5), `Synthetic` (§4.8.8), `Signature` (§4.8.9) and `Deprecated` (§4.8.15) attributes.

A Java virtual machine implementation must recognize and correctly read `Code` (§4.8.3) and `Exceptions` (§4.8.5) attributes found in the `attributes` table of a `method_info` structure. If a Java virtual machine recognizes class files whose major version is 49.0 or above, it must recognize and correctly read `Signature` (§4.8.9) attributes found in the `attributes` table of a `method_info` structure. A Java virtual machine implementation is required to silently ignore any or all other attributes in the `attributes` table of a `method_info` structure that it does not recognize. Attributes not defined in this specification are not allowed to affect the semantics of the class file, but only to provide additional descriptive information (§4.8.1).

4.8 Attributes

Attributes are used in the `ClassFile` (§4.2), `field_info` (§4.6), `method_info` (§4.7), `Code_attribute` (§4.8.3) structures of the class file format. All attributes have the following general format:

```

attribute_info {
    u2 attribute_name_index;
    u4 attribute_length;
    u1 info[attribute_length];
}

```

For all attributes, the `attribute_name_index` must be a valid unsigned 16-bit index into the constant pool of the class. The `constant_pool` entry at `attribute_name_index` must be a `CONSTANT_Utf8_info` (§4.5.7) structure representing the name of the attribute. The value of the `attribute_length` item indicates the length of the subsequent information in bytes. The length does not include the initial six bytes that contain the `attribute_name_index` and `attribute_length` items.

Certain attributes are predefined as part of the class file specification. The predefined attributes are the `SourceFile` (§4.8.10), `ConstantValue` (§4.8.2), `Code` (§4.8.3), [StackMapTable \(§4.8.4\)](#), `Exceptions` (§4.8.5), `InnerClasses` (§4.8.6), `EnclosingMethod` (§4.8.7), `Synthetic` (§4.8.8), `Signature` (§4.8.9), `LineNumberTable` (§4.8.12), `LocalVariableTable` and `Deprecated` (§4.8.15) attributes. Within the context of their use in this specification, that is, in the attributes tables of the class file structures in which they appear, the names of these predefined attributes are reserved.

Of the predefined attributes, the `Code`, `ConstantValue` and `Exceptions` attributes must be recognized and correctly read by a class file reader for correct interpretation of the class file by a Java virtual machine implementation. The [StackMapTable attribute must be recognized and correctly interpreted by any Java virtual machine implementation that recognizes class files whose major version is 50.0 or above.](#) The `Signature` attribute must be recognized and correctly interpreted by any Java virtual machine implementation that recognizes class files whose major version is 49.0 or above. The `InnerClasses`, `EnclosingMethod` and `Synthetic` attributes must be recognized and correctly read by a class file reader in order to properly implement the Java and Java 2 platform class libraries ~~(§3.12)~~(§3.12). Use of the remaining predefined attributes is optional; a class file reader may use the information they contain, or otherwise must silently ignore those attributes.

4.8.1 Defining and Naming New Attributes

Compilers are permitted to define and emit class files containing new attributes in the attributes tables of class file structures. Java virtual machine implementations are permitted to recognize and use new attributes found in the attributes

tables of `class` file structures. However, any attribute not defined as part of this Java virtual machine specification must not affect the semantics of class or interface types. Java virtual machine implementations are required to silently ignore attributes they do not recognize.

For instance, defining a new attribute to support vendor-specific debugging is permitted. Because Java virtual machine implementations are required to ignore attributes they do not recognize, `class` files intended for that particular Java virtual machine implementation will be usable by other implementations even if those implementations cannot make use of the additional debugging information that the `class` files contain.

Java virtual machine implementations are specifically prohibited from throwing an exception or otherwise refusing to use `class` files simply because of the presence of some new attribute. Of course, tools operating on `class` files may not run correctly if given `class` files that do not contain all the attributes they require.

Two attributes that are intended to be distinct, but that happen to use the same attribute name and are of the same length, will conflict on implementations that recognize either attribute. Attributes defined other than by Sun must have names chosen according to the package naming convention defined by *The Java Language Specification*. For instance, a new attribute defined by Netscape might have the name `"com.Netscape.new-attribute"`.³

Sun may define additional attributes in future versions of this `class` file specification.

4.8.2 The `ConstantValue` Attribute

The `ConstantValue` attribute is a fixed-length attribute used in the `attributes` table of the `field_info` (§4.6) structures. A `ConstantValue` attribute represents the value of a constant field. There can be no more than one `ConstantValue` attribute in the `attributes` table of a given `field_info` structure. If the field is static (that is, the `ACC_STATIC` bit (Table 4.4) in the `flags` item of the `field_info` structure is set) then the constant field represented by the `field_info` structure is assigned the value referenced by its `ConstantValue` attribute as part of the initialization of the class or interface declaring the constant field (§2.17.4)(§2.17.4). This

³ The first edition of *The Java Language Specification* required that `"com"` be in uppercase in this example. The second edition reversed that convention and uses lowercase.

occurs immediately prior to the invocation of the class or interface initialization method (§3.9)(§3.9) of that class or interface.

If a `field_info` structure representing a non-static field has a `ConstantValue` attribute, then that attribute must silently be ignored. Every Java virtual machine implementation must recognize `ConstantValue` attributes.

The `ConstantValue` attribute has the following format:

```
ConstantValue_attribute {  
    u2 attribute_name_index;  
    u4 attribute_length;  
    u2 constantvalue_index;  
}
```

The items of the `ConstantValue_attribute` structure are as follows:

attribute_name_index

The value of the `attribute_name_index` item must be a valid index into the `constant_pool` table. The `constant_pool` entry at that index must be a `CONSTANT_Utf8_info` (§4.5.7) structure representing the string "ConstantValue".

attribute_length

The value of the `attribute_length` item of a `ConstantValue_attribute` structure must be 2.

constantvalue_index

The value of the `constantvalue_index` item must be a valid index into the `constant_pool` table. The `constant_pool` entry at that index gives the constant value represented by this attribute. The `constant_pool` entry must be of a type appropriate to the field, as shown by Table 4.6.

Table 4.6 Constant value attribute types

Field Type	Entry Type
long	CONSTANT_Long
float	CONSTANT_Float
double	CONSTANT_Double
int, short, char, byte, boolean	CONSTANT_Integer
String	CONSTANT_String

4.8.3 The Code Attribute

The Code attribute is a variable-length attribute used in the attributes table of `method_info` structures. A Code attribute contains the Java virtual machine instructions and auxiliary information for a single method, instance initialization method (§3.9)(§3.9), or class or interface initialization method (§3.9)(§3.9). Every Java virtual machine implementation must recognize Code attributes. If the method is either native or abstract, its `method_info` structure must not have a Code attribute. Otherwise, its `method_info` structure must have exactly one Code attribute.

The Code attribute has the following format:

```
Code_attribute {
    u2 attribute_name_index;
    u4 attribute_length;
    u2 max_stack;
    u2 max_locals;
    u4 code_length;
    u1 code[code_length];
    u2 exception_table_length;
    {    u2 start_pc;
        u2 end_pc;
        u2 handler_pc;
        u2 catch_type;
    }    exception_table[exception_table_length];
    u2 attributes_count;
    attribute_info attributes[attributes_count];
}
```

The items of the `Code_attribute` structure are as follows:

`attribute_name_index`

The value of the `attribute_name_index` item must be a valid index into the `constant_pool` table. The `constant_pool` entry at that index must be a `CONSTANT_Utf8_info` (§4.5.7) structure representing the string "Code".

`attribute_length`

The value of the `attribute_length` item indicates the length of the attribute, excluding the initial six bytes.

`max_stack`

The value of the `max_stack` item gives the maximum depth ~~(§3.6.2)~~(§3.6.2) of the operand stack of this method at any point during execution of the method.

`max_locals`

The value of the `max_locals` item gives the number of local variables in the local variable array allocated upon invocation of this method, including the local variables used to pass parameters to the method on its invocation.

The greatest local variable index for a value of type `long` or `double` is `max_locals-2`. The greatest local variable index for a value of any other type is `max_locals-1`.

`code_length`

The value of the `code_length` item gives the number of bytes in the code array for this method. The value of `code_length` must be greater than zero; the code array must not be empty.

`code[]`

The code array gives the actual bytes of Java virtual machine code that implement the method.

When the code array is read into memory on a byte-addressable machine, if the first byte of the array is aligned on a 4-byte boundary, the *tableswitch* and *lookupswitch* 32-bit offsets will be 4-byte aligned. (Refer to the descriptions of those

instructions for more information on the consequences of code array alignment.)

The detailed constraints on the contents of the code array are extensive and are given in a separate section (§4.10).

`exception_table_length`

The value of the `exception_table_length` item gives the number of entries in the `exception_table` table.

`exception_table[]`

Each entry in the `exception_table` array describes one exception handler in the code array. The order of the handlers in the `exception_table` array is significant. See ~~Section 3.10~~[Section 3.10](#) for more details.

Each `exception_table` entry contains the following four items:

`start_pc`, `end_pc`

The values of the two items `start_pc` and `end_pc` indicate the ranges in the code array at which the exception handler is active. The value of `start_pc` must be a valid index into the code array of the opcode of an instruction. The value of `end_pc` either must be a valid index into the code array of the opcode of an instruction or must be equal to `code_length`, the length of the code array. The value of `start_pc` must be less than the value of `end_pc`.

The `start_pc` is inclusive and `end_pc` is exclusive; that is, the exception handler must be active while the program counter is within the interval `[start_pc, end_pc)`.⁴

`handler_pc`

The value of the `handler_pc` item indicates the start of the exception handler. The value of the item must be a valid index into the code array and must be the index of the opcode of an instruction.

⁴ The fact that `end_pc` is exclusive is a historical mistake in the design of the Java virtual machines.

catch_type

If the value of the catch_type item is nonzero, it must be a valid index into the constant_pool table. The constant_pool entry at that index must be a CONSTANT_Class_info (§4.5.1) structure representing a class of exceptions that this exception handler is designated to catch. The exception handler will be called only if the thrown exception is an instance of the given class or one of its subclasses.

If the value of the catch_type item is zero, this exception handler is called for all exceptions. This is used to implement finally (see [Section 7.13, “Compiling finally”](#)).

attributes_count

The value of the attributes_count item indicates the number of attributes of the Code attribute.

attributes[]

Each value of the attributes table must be an attribute structure (§4.8). A Code attribute can have any number of optional attributes associated with it.

Currently, the LineNumberTable (§4.8.12) and LocalVariableTable (§4.8.13), attributes which contain debugging information,

[as well as the StackMapTable attribute \(§4.8.4\)](#), are defined and used with the Code attribute.

A Java virtual machine implementation is permitted to silently ignore any or all attributes in the attributes table of a Code attribute, [except for the StackMapTable attribute, which must be recognized if the class file version number is 50.0 or above](#). Attributes not defined in this specification are not allowed to affect the semantics of the class file, but only to provide additional descriptive information (§4.8.1).

4.8.4 The StackMapTable Attribute

The stack map attribute is a variable-length attribute in the attributes table of a Code attribute. The name of the attribute is `StackMapTable`. This attribute is used during the process of verification by typechecking (§4.11.1).

A stack map attribute consists of zero or more *stack map frames*. Each stack map frame specifies (either explicitly or implicitly) a bytecode offset, the verification types (§4.11.1) for the local variables, and the verification types for the operand stack.

The type checker deals with and manipulates the expected types of a method's local variables and operand stack. Throughout this section, a *location* refers to either a single local variable or to a single operand stack entry.

We will use the terms *stack map frame* and *type state* interchangeably to describe a mapping from locations in the operand stack and local variables of a method to verification types. We will usually use the term *stack map frame* when such a mapping is provided in the class file, and the term *type state* when the mapping is inferred by the type checker.

If a method's Code attribute does not have a `StackMapTable` attribute, it has an *implicit stack map attribute*. This implicit stack map attribute is equivalent to a `StackMapTable` attribute with `number_of_entries` equal to zero. A method's Code attribute may have at most one `StackMapTable` attribute, otherwise a `ClassFormatError` is thrown.

```
The format of the stack map in the class file is given below.
stack_map {
// attribute StackMapTable
    u2 attribute_name_index;
    u4 attribute_length
    u2 number_of_entries;
    stack_map_frame entries[number_of_entries];
}
```

Each `stack_map_frame` structure specifies the type state at a particular bytecode offset. Each frame type specifies (explicitly or implicitly) a value, `offset_delta`, that is used to calculate the actual bytecode offset at which it applies. The bytecode offset at which the frame applies is given by adding `1 + offset_delta` to the offset of the previous frame, unless the previous frame is the initial frame of the method, in which case the bytecode offset is `offset_delta`.

By using an offset delta rather than the actual bytecode offset we ensure, by definition, that stack map frames are in the correctly sorted order. Furthermore, by consistently using the formula `offset_delta + 1` for all explicit frames, we guarantee the absence of duplicates.

All frame types, even `full_frame`, rely on the previous frame for some of their semantics. This raises the question of what is the very first frame? The initial frame is implicit, and computed from the method descriptor. See the Prolog code for `methodInitialStackMapFrame`.

The `stack_map_frame` structure consists of a one-byte tag followed by zero or more bytes, giving more information, depending upon the tag.

A stack map frame may belong to one of several *frame types*

```
union stack_map_frame {
    same_frame;
    same_locals_1_stack_item_frame;
    same_locals_1_stack_item_frame_extended;
    chop_frame;
    same_frame_extended;
    append_frame;
    full_frame;
}
```

The frame type `same_frame` is represented by tags in the range [0-63]. If the frame type is `same_frame`, it means the frame has exactly the same locals as the previous stack map frame and that the number of stack items is zero. The `offset_delta` value for the frame is the value of the tag item, `frame_type`. The form of such a frame is then:

```
same_frame {
    u1 frame_type = SAME; /* 0-63 */
}
```

The frame type `same_locals_1_stack_item_frame` is represented by tags in the range [64, 127]. If the frame type is `same_locals_1_stack_item_frame`, it

means the frame has exactly the same locals as the previous stack map frame and that the number of stack items is 1. The offset_delta value for the frame is the value (frame_type - 64). There is a verification_type info following the frame_type for the one stack item. The form of such a frame is then:

```
same_locals_1_stack_item_frame {
    u1 frame_type = SAME_LOCALS_1_STACK_ITEM; /* 64-127 */
    verification_type_info stack[1];
}
```

Tags in the range [128-246] are reserved for future use.

The frame type same_locals_1_stack_item_frame_extended is represented by the tag 247. The frame type same_locals_1_stack_item_frame_extended indicates that the frame has exactly the same locals as the previous stack map frame and that the number of stack items is 1. The offset_delta value for the frame is given explicitly. There is a verification_type info following the frame_type for the one stack item. The form of such a frame is then:

```
same_locals_1_stack_item_frame_extended {
    u1 frame_type = SAME_LOCALS_1_STACK_ITEM_EXTENDED; /*
247 */
    u2 offset_delta;
    verification_type_info stack[1];
}
```

The frame type chop_frame is represented by tags in the range [248-250]. If the frame_type is chop_frame, it means that the operand stack is empty and the current locals are the same as the locals in the previous frame, except that the k last locals are absent. The value of k is given by the formula $251 - \text{frame_type}$.

The form of such a frame is then:

```
chop_frame {
    u1 frame_type=CHOP; /* 248-250 */
    u2 offset_delta;
}
```

The frame type same_frame_extended is represented by the tag value 251. If the frame_type is same_frame_extended, it means the frame has exactly the same locals as the previous stack map frame and that the number of stack items is zero.

The form of such a frame is then:

```

same_frame_extended {
    u1 frame_type = SAME_FRAME_EXTENDED; /* 251 */
    u2 offset_delta;
}

```

The frame type `append_frame` is represented by tags in the range [252-254]. If the frame type is `append_frame`, it means that the operand stack is empty and the current locals are the same as the locals in the previous frame, except that k additional locals are defined. The value of k is given by the formula `frame_type - 251`.

The form of such a frame is then:

```

append_frame {
    u1 frame_type = APPEND; /* 252-254 */
    u2 offset_delta;
    verification_type_info locals[frame_type - 251];
}

```

The 0th entry in `locals` represents the type of the first additional local variable. If `locals[M]` represents local variable N , then `locals[M+1]` represents local variable $N+1$ if `locals[M]` is one of `Top_variable_info`, `Integer_variable_info`, `Float_variable_info`, `Null_variable_info`, `UninitializedThis_variable_info`, `Object_variable_info`, or `Uninitialized_variable_info`, otherwise `locals[M+1]` represents local variable $N+2$. It is an error if, for any index i , `locals[i]` represents a local variable whose index is greater than the maximum number of local variables for the method.

The frame type `full_frame` is represented by the tag value 255. The form of such a frame is then:

```

full_frame {
    u1 frame_type = FULL_FRAME; /* 255 */
    u2 offset_delta;
    u2 number_of_locals;
    verification_type_info locals[number_of_locals];
    u2 number_of_stack_items;
}

```

```

        verification_type_info_stack[number_of_stack_items];
    }

```

The 0th entry in `locals` represents the type of local variable 0. If `locals[M]` represents local variable N, then `locals[M+1]` represents local variable N+1 if `locals[M]` is one of `Top_variable_info`, `Integer_variable_info`, `Float_variable_info`, `Null_variable_info`, `UninitializedThis_variable_info`, `Object_variable_info`, or `Uninitialized_variable_info`, otherwise `locals[M+1]` represents local variable N+2. It is an error if, for any index *i*, `locals[i]` represents a local variable whose index is greater than the maximum number of local variables for the method.

The 0th entry in `stack` represents the type of the bottom of the stack, and subsequent entries represent types of stack elements closer to the top of the operand stack. We shall refer to the bottom element of the stack as stack element 0, and to subsequent elements as stack element 1, 2 etc. If `stack[M]` represents stack element N, then `stack[M+1]` represents stack element N+1 if `stack[M]` is one of `Top_variable_info`, `Integer_variable_info`, `Float_variable_info`, `Null_variable_info`, `UninitializedThis_variable_info`, `Object_variable_info`, or `Uninitialized_variable_info`, otherwise `stack[M+1]` represents stack element N+2. It is an error if, for any index *i*, `stack[i]` represents a stack entry whose index is greater than the maximum operand stack size for the method.

We say that an instruction in the bytecode has a corresponding stack map frame if the offset in the offset item of the stack map frame is the same as the offset of the instruction in the bytecodes.

The `verification_type_info` structure consists of a one-byte tag followed by zero or more bytes, giving more information about the tag. Each `verification_type_info` structure specifies the verification type of one or two locations.

```

union verification_type_info {
    Top_variable_info;
    Integer_variable_info;
    Float_variable_info;
    Long_variable_info;
    Double_variable_info;
    Null_variable_info;
    UninitializedThis_variable_info;
    Object_variable_info;
}

```

```

    Uninitialized_variable_info;
}

```

The Top_variable_info type indicates that the local variable has the verification type top (\top).

```

    Top_variable_info {
        u1 tag = ITEM_Top; /* 0 */
    }

```

The Integer_variable_info type indicates that the location contains the verification type int.

```

    Integer_variable_info {
        u1 tag = ITEM_Integer; /* 1 */
    }

```

The Float_variable_info type indicates that the location contains the verification type float.

```

    Float_variable_info {
        u1 tag = ITEM_Float; /* 2 */
    }

```

The Long_variable_info type indicates that the location contains the verification type long. If the location is a local variable, then:

- It must not be the local variable with the highest index.
- The next higher numbered local variable contains the verification type \top .

If the location is an operand stack entry, then:

- The current location must not be the topmost location of the operand stack.
- the next location closer to the top of the operand stack contains the verification type \top .

This structure gives the contents of two locations in the operand stack or in the local variables.

```

    Long_variable_info {
        u1 tag = ITEM_Long; /* 4 */
    }

```

The Double_variable_info type indicates that the location contains the verification type double. If the location is a local variable, then:

- It must not be the local variable with the highest index.

- The next higher numbered local variable contains the verification type \top .

If the location is an operand stack entry, then:

- The current location must not be the topmost location of the operand stack.
- the next location closer to the top of the operand stack contains the verification type \top .

This structure gives the contents of two locations in in the operand stack or in the local variables.

```
Double_variable_info {
    u1 tag = ITEM_Double; /* 3 */
}
```

The Null_variable_info type indicates that location contains the verification type null.

```
Null_variable_info {
    u1 tag = ITEM_Null; /* 5 */
}
```

The UninitializedThis_variable_info type indicates that the location contains the verification type uninitializedThis.

```
UninitializedThis_variable_info {
    u1 tag = ITEM_UninitializedThis; /* 6 */
}
```

The Object_variable_info type indicates that the location contains an instance of the class referenced by the constant pool entry.

```
Object_variable_info {
    u1 tag = ITEM_Object; /* 7 */
    u2 cpool_index;
}
```

The Uninitialized_variable_info indicates that the location contains the verification type uninitialized(*offset*). The offset item indicates the offset of the new instruction that created the object being stored in the location.

```
Uninitialized_variable_info {
    u1 tag = ITEM_Uninitialized /* 8 */
    u2 offset;
}
```

4.8.5 The Exceptions Attribute

The Exceptions attribute is a variable-length attribute used in the attributes table of a `method_info` (§4.7) structure. The Exceptions attribute indicates which checked exceptions a method may throw. There may be at most one Exceptions attribute in each `method_info` structure.

The Exceptions attribute has the following format:

```
Exceptions_attribute {  
    u2 attribute_name_index;  
    u4 attribute_length;  
    u2 number_of_exceptions;  
    u2 exception_index_table[number_of_exceptions];  
}
```

The items of the Exceptions_attribute structure are as follows:

`attribute_name_index`

The value of the `attribute_name_index` item must be a valid index into the `constant_pool` table. The `constant_pool` entry at that index must be the `CONSTANT_Utf8_info` (§4.5.7) structure representing the string "Exceptions".

`attribute_length`

The value of the `attribute_length` item indicates the attribute length, excluding the initial six bytes.

`number_of_exceptions`

The value of the `number_of_exceptions` item indicates the number of entries in the `exception_index_table`.

`exception_index_table[]`

Each value in the `exception_index_table` array must be a valid index into the `constant_pool` table. The `constant_pool` entry referenced by each table item must be a `CONSTANT_Class_info` (§4.5.1) structure representing a class type that this method is declared to throw.

A method should throw an exception only if at least one of the following three criteria is met:

- The exception is an instance of `RuntimeException` or one of its subclasses.
- The exception is an instance of `Error` or one of its subclasses.
- The exception is an instance of one of the exception classes specified in the `exception_index_table` just described, or one of their subclasses.

These requirements are not enforced in the Java virtual machine; they are enforced only at compile time.

4.8.6 The InnerClasses Attribute

The InnerClasses attribute⁵ is a variable-length attribute in the `attributes` table of the `ClassFile` (§4.2) structure. If the constant pool of a class or interface *C* contains a `CONSTANT_Class_info` entry which represents a class or interface that is not a member of a package, then *C*'s `ClassFile` structure must have exactly one InnerClasses attribute in its `attributes` table.

The InnerClasses attribute has the following format:

```
InnerClasses_attribute {
    u2 attribute_name_index;
    u4 attribute_length;
    u2 number_of_classes;
    { u2 inner_class_info_index;
      u2 outer_class_info_index;
      u2 inner_name_index;
      u2 inner_class_access_flags;
    } classes[number_of_classes];
}
```

The items of the InnerClasses_attribute structure are as follows:

`attribute_name_index`

The value of the `attribute_name_index` item must be a valid index into the `constant_pool` table. The `constant_pool` entry at that index must be a `CONSTANT_Utf8_info` (§4.5.7) structure representing the string "InnerClasses".

⁵ The InnerClasses attribute was introduced in JDK release 1.1 to support nested classes and interfaces.

attribute_length

The value of the `attribute_length` item indicates the length of the attribute, excluding the initial six bytes.

number_of_classes

The value of the `number_of_classes` item indicates the number of entries in the `classes` array.

classes[]

Every `CONSTANT_Class_info` entry in the `constant_pool` table which represents a class or interface *C* that is not a package member must have exactly one corresponding entry in the `classes` array.

If a class has members that are classes or interfaces, its `constant_pool` table (and hence its `InnerClasses` attribute) must refer to each such member, even if that member is not otherwise mentioned by the class. These rules imply that a nested class or interface member will have `InnerClasses` information for each enclosing class and for each immediate member.

Each `classes` array entry contains the following four items:

inner_class_info_index

The value of the `inner_class_info_index` item must be zero or a valid index into the `constant_pool` table. The `constant_pool` entry at that index must be a `CONSTANT_Class_info` (§4.5.1) structure representing *C*. The remaining items in the `classes` array entry give information about *C*.

outer_class_info_index

If *C* is not a member, the value of the `outer_class_info_index` item must be zero. Otherwise, the value of the `outer_class_info_index` item must be a valid index into the `constant_pool` table, and the entry at that index must be a `CONSTANT_Class_info` (§4.5.1) structure representing the class or interface of which *C* is a member.

inner_name_index

If `C` is anonymous, the value of the `inner_name_index` item must be zero. Otherwise, the value of the `inner_name_index` item must be a valid index into the `constant_pool` table, and the entry at that index must be a `CONSTANT_Utf8_info` (§4.5.7) structure that represents the original simple name of `C`, as given in the source code from which this `class` file was compiled.

inner_class_access_flags

The value of the `inner_class_access_flags` item is a mask of flags used to denote access permissions to and properties of class or interface `C` as declared in the source code from which this `class` file was compiled. It is used by compilers to recover the original information when source code is not available. The flags are shown in Table 4.7.

Table 4.7 Nested class access and property flags

Flag Name	Value	Meaning
ACC_PUBLIC	0x0001	Marked or implicitly <code>public</code> in source.
ACC_PRIVATE	0x0002	Marked <code>private</code> in source.
ACC_PROTECTED	0x0004	Marked <code>protected</code> in source.
ACC_STATIC	0x0008	Marked or implicitly <code>static</code> in source.
ACC_FINAL	0x0010	Marked <code>final</code> in source.
ACC_INTERFACE	0x0200	Was an <code>interface</code> in source.
ACC_ABSTRACT	0x0400	Marked or implicitly <code>abstract</code> in source.
ACC_SYNTHETIC	0x1000	Declared <code>synthetic</code> ; Not present in the source code.
ACC_ANNOTATION	0x2000	Declared as an annotation type.
ACC_ENUM	0x4000	Declared as an enum type.

All bits of the `inner_class_access_flags` item not assigned in Table 4.7 are reserved for future use. They should be

set to zero in generated class files and should be ignored by Java virtual machine implementations.

The Java virtual machine does not currently check the consistency of the `InnerClasses` attribute with any class file actually representing a class or interface referenced by the attribute.

4.8.7 The EnclosingMethod Attribute

The `EnclosingMethod` attribute is an optional fixed-length attribute in the attributes table of the `ClassFile` (§4.2) structure. A class must have an `EnclosingMethod` attribute if and only if it is a local class or an anonymous class. A class may have no more than one `EnclosingMethod` attribute.

The `EnclosingMethod` attribute has the following format:

```
EnclosingMethod_attribute {  
    u2 attribute_name_index;  
    u4 attribute_length;  
    u2 class_index  
    u2 method_index;  
}
```

The items of the `EnclosingMethod_attribute` structure are as follows:

`attribute_name_index`

The value of the `attribute_name_index` item must be a valid index into the `constant_pool` table. The `constant_pool` entry at that index must be a `CONSTANT_Utf8_info` (§4.5.7) structure representing the string "EnclosingMethod".

`attribute_length`

The value of the `attribute_length` item is four.

`class_index`

The value of the `class_index` item must be a valid index into the `constant_pool` table. The `constant_pool` entry at that index must be a `CONSTANT_Class_info` (§4.5.1) structure representing the innermost class that encloses the declaration of the current class.

method_index

If the current class is not immediately enclosed by a method or constructor, then the value of the `method_index` item must be zero. Otherwise, the value of the `method_index` item must be a valid index into the `constant_pool` table. The `constant_pool` entry at that index must be a **CONSTANT_NameAndType_info** (§4.5.6) structure representing a the name and type of a method in the class referenced by the `class_index` attribute above. It is the responsibility of the Java compiler to ensure that the method identified via the `method_index` is indeed the closest lexically enclosing method of the class that contains this `EnclosingMethod` attribute.

4.8.8 The Synthetic Attribute

The `Synthetic` attribute⁶ is a fixed-length attribute in the `attributes` table of `ClassFile` (§4.2), `field_info` (§4.6), and `method_info` (§4.7) structures. A class member that does not appear in the source code must be marked using a `Synthetic` attribute, or else it must have its `ACC_SYNTHETIC` bit set. The only exceptions to this requirement are for default constructors and the class initialization method.

The `Synthetic` attribute has the following format:

```
Synthetic_attribute {
    u2 attribute_name_index;
    u4 attribute_length;
}
```

The items of the `Synthetic_attribute` structure are as follows:

attribute_name_index

The value of the `attribute_name_index` item must be a valid index into the `constant_pool` table. The `constant_pool` entry

⁶ The `Synthetic` attribute was introduced in JDK release 1.1 to support nested classes and interfaces.

at that index must be a `CONSTANT_Utf8_info` (§4.5.7) structure representing the string "Synthetic".

`attribute_length`

The value of the `attribute_length` item is zero.

4.8.9 The Signature Attribute

The Signature attribute is an optional fixed-length attribute in the `attributes` table of the `ClassFile` (§4.2), `field_info` (§4.6) and `method_info` (§4.7) structures.

The Signature attribute has the following format:

```
Signature_attribute {  
    u2 attribute_name_index;  
    u4 attribute_length;  
    u2 signature_index;  
}
```

The items of the `Signature_attribute` structure are as follows:

`attribute_name_index`

The value of the `attribute_name_index` item must be a valid index into the `constant_pool` table. The `constant_pool` entry at that index must be a `CONSTANT_Utf8_info` (§4.5.7) structure representing the string "Signature".

`attribute_length`

The value of the `attribute_length` item of a `Signature_attribute` structure must be 2.

`signature_index`

The value of the `signature_index` item must be a valid index into the `constant_pool` table. The `constant pool` entry at that index must be a `CONSTANT_Utf8_info` (§4.5.7) structure representing either a class signature, if this signature attribute is an attribute of a `ClassFile` structure, a method type signature, if this signature is an attribute of a `method_info` structure, or a field type signature otherwise.

4.8.10 The SourceFile Attribute

The SourceFile attribute is an optional fixed-length attribute in the attributes table of the ClassFile (§4.2) structure. There can be no more than one SourceFile attribute in the attributes table of a given ClassFile structure.

The SourceFile attribute has the following format:

```
SourceFile_attribute {  
    u2 attribute_name_index;  
    u4 attribute_length;  
    u2 sourcefile_index;  
}
```

The items of the SourceFile_attribute structure are as follows:

attribute_name_index

The value of the attribute_name_index item must be a valid index into the constant_pool table. The constant_pool entry at that index must be a CONSTANT_Utf8_info (§4.5.7) structure representing the string "SourceFile".

attribute_length

The value of the attribute_length item of a SourceFile_attribute structure must be 2.

sourcefile_index

The value of the sourcefile_index item must be a valid index into the constant_pool table. The constant pool entry at that index must be a CONSTANT_Utf8_info (§4.5.7) structure representing a string.

The string referenced by the sourcefile_index item will be interpreted as indicating the name of the source file from which this class file was compiled. It will not be interpreted as indicating the name of a directory containing the file or an absolute path name for the file; such platform-specific additional information must be supplied by the runtime interpreter or development tool at the time the file name is actually used.

4.8.11 The SourceDebugExtension Attribute

The SourceDebugExtension attribute is an optional attribute in the attributes table of the ClassFile (§4.2) structure. There can be no more than one SourceDebugExtension attribute in the attributes table of a given ClassFile structure.

The SourceDebugExtension attribute has the following format:

```
SourceDebugExtension_attribute {  
    u2 attribute_name_index;  
    u4 attribute_length;  
    u1 debug_extension[attribute_length];  
}
```

The items of the SourceDebugExtension_attribute structure are as follows:

attribute_name_index

The value of the attribute_name_index item must be a valid index into the constant_pool table. The constant_pool entry at that index must be a CONSTANT_Utf8_info (§4.5.7) structure representing the string "SourceDebugExtension".

attribute_length

The value of the attribute_length item indicates the length of the attribute, excluding the initial six bytes. The value of the attribute_length item is thus the number of bytes in the debug_extension[] item.

debug_extension[]

The debug_extension array holds a string, which must be in UTF-8 format. There is no terminating zero byte. The string in the debug_extension item will be interpreted as extended debugging information. The content of this string has no semantic effect on the Java Virtual Machine.

4.8.12 The LineNumberTable Attribute

The `LineNumberTable` attribute is an optional variable-length attribute in the `attributes` table of a `Code` (§4.8.3) attribute. It may be used by debuggers to determine which part of the Java virtual machine code array corresponds to a given line number in the original source file. If `LineNumberTable` attributes are present in the `attributes` table of a given `Code` attribute, then they may appear in any order. Furthermore, multiple `LineNumberTable` attributes may together represent a given line of a source file; that is, `LineNumberTable` attributes need not be one-to-one with source lines.

The `LineNumberTable` attribute has the following format:

```
LineNumberTable_attribute {
    u2 attribute_name_index;
    u4 attribute_length;
    u2 line_number_table_length;
    { u2 start_pc;
      u2 line_number;
    } line_number_table[line_number_table_length];
}
```

The items of the `LineNumberTable_attribute` structure are as follows:

`attribute_name_index`

The value of the `attribute_name_index` item must be a valid index into the `constant_pool` table. The `constant_pool` entry at that index must be a `CONSTANT_Utf8_info` (§4.5.7) structure representing the string "`LineNumberTable`".

`attribute_length`

The value of the `attribute_length` item indicates the length of the attribute, excluding the initial six bytes.

`line_number_table_length`

The value of the `line_number_table_length` item indicates the number of entries in the `line_number_table` array.

`line_number_table[]`

Each entry in the `line_number_table` array indicates that the line number in the original source file changes at a given point in

the code array. Each `line_number_table` entry must contain the following two items:

`start_pc`

The value of the `start_pc` item must indicate the index into the code array at which the code for a new line in the original source file begins. The value of `start_pc` must be less than the value of the `code_length` item of the `Code` attribute of which this `LineNumberTable` is an attribute.

`line_number`

The value of the `line_number` item must give the corresponding line number in the original source file.

4.8.13 The `LocalVariableTable` Attribute

The `LocalVariableTable` attribute is an optional variable-length attribute of a `Code` (§4.8.3) attribute. It may be used by debuggers to determine the value of a given local variable during the execution of a method. If `LocalVariableTable` attributes are present in the attributes table of a given `Code` attribute, then they may appear in any order. There may be no more than one `LocalVariableTable` attribute per local variable in the `Code` attribute.

The `LocalVariableTable` attribute has the following format:

```
LocalVariableTable_attribute {
    u2 attribute_name_index;
    u4 attribute_length;
    u2 local_variable_table_length;
    { u2 start_pc;
      u2 length;
      u2 name_index;
      u2 descriptor_index;
      u2 index;
    } local_variable_table[
        local_variable_table_length];
}
```

The items of the `LocalVariableTable_attribute` structure are as follows:

attribute_name_index

The value of the `attribute_name_index` item must be a valid index into the `constant_pool` table. The `constant_pool` entry at that index must be a `CONSTANT_Utf8_info` (§4.5.7) structure representing the string "LocalVariableTable".

attribute_length

The value of the `attribute_length` item indicates the length of the attribute, excluding the initial six bytes.

local_variable_table_length

The value of the `local_variable_table_length` item indicates the number of entries in the `local_variable_table` array.

local_variable_table[]

Each entry in the `local_variable_table` array indicates a range of code array offsets within which a local variable has a value. It also indicates the index into the local variable array of the current frame at which that local variable can be found. Each entry must contain the following five items:

start_pc, length

The given local variable must have a value at indices into the code array in the interval `[start_pc, start_pc+length)`, that is, between `start_pc` and `start_pc+length` exclusive. The value of `start_pc` must be a valid index into the code array of this Code attribute and must be the index of the opcode of an instruction. The value of `start_pc+length` must either be a valid index into the code array of this Code attribute and be the index of the opcode of an instruction, or it must be the first index beyond the end of that code array.

name_index, descriptor_index

The value of the `name_index` item must be a valid index into the `constant_pool` table. The `constant_pool` entry at that index must contain a `CONSTANT_Utf8_info` (§4.5.7)

structure representing a valid unqualified name (§4.3.2) denoting a local variable.

The value of the `descriptor_index` item must be a valid index into the `constant_pool` table. The `constant_pool` entry at that index must contain a `CONSTANT_Utf8_info` (§4.5.7) structure. That `CONSTANT_Utf8_info` structure must represent a field descriptor (§4.4.2) encoding the type of a local variable in the source program.

`index`

The given local variable must be at `index` in the local variable array of the current frame. If the local variable at `index` is of type `double` or `long`, it occupies both `index` and `index+1`.

4.8.14 The `LocalVariableTypeTable` Attribute

The `LocalVariableTypeTable` attribute is an optional variable-length attribute of a `Code` (§4.8.3) attribute. It may be used by debuggers to determine the value of a given local variable during the execution of a method. If `LocalVariableTypeTable` attributes are present in the `attributes` table of a given `Code` attribute, then they may appear in any order. There may be no more than one `LocalVariableTypeTable` attribute per local variable in the `Code` attribute.

The `LocalVariableTypeTable` attribute differs from the `LocalVariableTable` attribute in that it provides signature information rather than descriptor information. This difference is only significant for variables whose type is a generic reference type. Such variables will appear in both tables, while variables of other types will appear only in `LocalVariableTable`.

The `LocalVariableTypeTable` attribute has the following format:

```

LocalVariableTypeTable_attribute {
    u2 attribute_name_index;
    u4 attribute_length;
    u2 local_variable_type_table_length;
    { u2 start_pc;
      u2 length;
      u2 name_index;
      u2 signature_index;
      u2 index;
    } local_variable_type_table[
        local_variable_type_table_length];
}

```

The items of the LocalVariableTypeTable_attribute structure are as follows:

attribute_name_index

The value of the attribute_name_index item must be a valid index into the constant_pool [table-a124table](#). The constant_pool entry at that index must be a CONSTANT_Utf8_info (§4.5.7) [a122](#)-structure representing the string "LocalVariableTypeTable"[-a123](#)".

attribute_length

The value of the attribute_length item indicates the length of the attribute, excluding the initial six bytes.

[local_variable_table_length](#)

[local_variable_type_table_length](#)

The value of the [local_variable_table_length-local_variable_type_table_length](#) item indicates the number of entries in the local_variable_table array.

[local_variable_tablelocal_variable_type_table\[\]](#)

Each entry in the [local_variable_table-local_variable_type_table](#) array indicates a range of code array offsets within which a local variable has a value. It also indicates the index into the local variable array of the current frame at which that local variable can be found. Each entry must contain the following five items:

`start_pc, length`

The given local variable must have a value at indices into the code array in the interval `[start_pc, start_pc+length)`, that is, between `start_pc` and `start_pc+length` exclusive. The value of `start_pc` must be a valid index into the code array of this Code attribute and must be the index of the opcode of an ~~instruction~~^{a124}~~instruction~~^a. The value of `start_pc+length` must either be a valid index into the code array of this Code attribute and be the index of the opcode of an instruction, or it must be the first index beyond the end of that code ~~array~~^{a125}~~array~~.

`name_index, signature_index`

The value of the `name_index` item must be a valid index into the `constant_pool` ~~table~~^{a127}~~table~~. The `constant_pool` entry at that index must contain a `CONSTANT_Utf8_info` (§4.5.7) structure ~~a128~~ representing a valid unqualified name (§4.3.2) denoting a local ~~variable~~^{a128}~~variable~~. ~~Careful here — do we want any restrictions at all?~~

The value of the `signature_index` item must be a valid index into the `constant_pool` table. The `constant_pool` entry at that index must contain a `CONSTANT_Utf8_info` (§4.5.7) structure representing a field type signature (§4.4.4) encoding the type of a local variable in the source program.

`index`

The given local variable must be at `index` in the local variable array of the current frame. If the local variable at `index` is of type `double` or `long`, it occupies both `index` and `index+1`.

4.8.15 ~~The Deprecated Attribute~~

The `Deprecated` attribute⁷ is an optional fixed-length attribute in the `attributes` table of `ClassFile` (§4.2), `field_info` (§4.6), and `method_info` (§4.7) structures. A class, interface, method, or field may be marked using a `Deprecated` attribute to indicate that the class, interface, method, or field has been superseded. A

runtime interpreter or tool that reads the class file format, such as a compiler, can use this marking to advise the user that a superseded class, interface, method, or field is being referred to. The presence of a `Deprecated` attribute does not alter the semantics of a class or interface.

The `Deprecated` attribute has the following format:

```
Deprecated_attribute {
    u2 attribute_name_index;
    u4 attribute_length;
}
```

The items of the `Deprecated_attribute` structure are as follows:

`attribute_name_index`

The value of the `attribute_name_index` item must be a valid index into the `constant_pool` table. The `constant_pool` entry at that index must be a `CONSTANT_Utf8_info` (§4.5.7) structure representing the string "Deprecated".

`attribute_length`

The value of the `attribute_length` item is zero.

4.8.16 The `RuntimeVisibleAnnotations` attribute

The `RuntimeVisibleAnnotations` attribute is a variable length attribute in the attributes table of the `ClassFile`, `field_info`, and `method_info` structures. The `RuntimeVisibleAnnotations` attribute records runtime-visible Java programming language annotations on the corresponding class, method, or field. Each `ClassFile`, `field_info`, and `method_info` structure may contain at most one `RuntimeVisibleAnnotations` attribute, which records all the runtime-visible Java programming language annotations on the corresponding program element. The JVM must make these annotations available so they can be returned by the appropriate reflective APIs.

The `RuntimeVisibleAnnotations` attribute has the following format:

⁷ The `Deprecated` attribute was introduced in JDK release 1.1 to support the `@deprecated` tag in documentation comments.

```

RuntimeVisibleAnnotations_attribute {
    u2 attribute_name_index;
    u4 attribute_length;
    u2 num_annotations;
    annotation annotations[num_annotations];
}

```

The items of the `RuntimeVisibleAnnotations` structure are as follows:

`attribute_name_index`

The value of the `attribute_name_index` item must be a valid index into the `constant_pool` table. The `constant_pool` entry at that index must be a `CONSTANT_Utf8_info` structure representing the string "RuntimeVisibleAnnotations".

`attribute_length`

The value of the `attribute_length` item indicates the length of the attribute, excluding the initial six bytes. The value of the `attribute_length` item is thus dependent on the number of runtime-visible annotations represented by the structure, and their values.

`num_annotations`

The value of the `num_annotations` item gives the number of runtime-visible annotations represented by the structure. Note that a maximum of 65535 runtime-visible Java programming language annotations may be directly attached to a program element.

`annotations`

Each value of the `annotations` table represents a single runtime-visible annotation on a program element.

The annotation structure has the following format:

```

annotation {
    u2    type_index;
    u2 num_element_value_pairs;
    {    u2 element_name_index;
    element_value value;
    }    element_value_pairs[num_element_value_pairs]
}

```

The items of the annotation structure are as follows:

`type_index`

The value of the `type_index` item must be a valid index into the `constant_pool` table. The `constant_pool` entry at that index must be a `CONSTANT_Utf8_info` structure representing a field descriptor representing the annotation type corresponding to the annotation represented by this annotation structure.

`num_element_value_pairs`

The value of the `num_element_value_pairs` item gives the number of element-value pairs of the annotation represented by this annotation structure. Note that a maximum of 65535 element-value pairs may be contained in a single annotation.

`element_value_pairs`

Each value of the `element_value_pairs` table represents a single element-value pair in the annotation represented by this annotation structure. Each `element_value_pairs` entry contains the following two items:

`element_name_index`

The value of the `element_name_index` item must be a valid index into the `constant_pool` table. The `constant_pool` entry at that index must be a `CONSTANT_Utf8_info` structure representing the name of the annotation type element represented by this `element_value_pairs` entry.

`value`

The value of the `value` item represents the value of the element-value pair represented by this `element_value_pairs` entry.

4.8.16.1 The `element_value` structure

The `element_value` structure is a discriminated union representing the value of an element-value pair. It is used to represent element values in all attributes that describe annotations (`RuntimeVisibleAnnotations`, `RuntimeInvisibleAnnotations`, `RuntimeVisibleParameterAnnotations`, and `RuntimeInvisibleParameterAnnotations`).

The `element_value` structure has the following format:

```

element_value {
    u1 tag;
    union {
        u2  const_value_index;
    {
        u2 type_name_index;
        u2 const_name_index;
    } enum_const_value;
        u2  class_info_index;
        annotation annotation_value;
    {
        u2  num_values;
        element_value values[num_values];
        } array_value;
    } value;
}

```

The items of the `element_value` structure are as follows:

tag

The `tag` item indicates the type of this annotation element-value pair. The letters 'B', 'C', 'D', 'F', 'I', 'J', 'S', and 'Z' indicate a primitive type. These letters are interpreted as `BaseType` characters (§Table 4.2). The other legal values for `tag` are listed with their interpretations in this table:

Table 4.8

tag value	Element Type
s	String
e	enum constant
c	class
@	annotation type
[array

value

The `value` item represents the value of this annotation element. This item is a union. The `tag` item, above, determines which item of the union is to be used:

const_value_index

The `const_value_index` item is used if the `tag` item is one of 'B', 'C', 'D', 'F', 'I', 'J', 'S', 'Z', or 's'. The value of the `const_value_index` item must be a valid index into the `constant_pool` table. The `constant_pool` entry at that index must be of the correct entry type for the field type designated by the `tag` item, as specified in Table 4.8.

enum_const_value

The `enum_const_value` item is used if the `tag` item is 'e'. The `enum_const_value` item consists of the following two items:

type_name_index

The value of the `type_name_index` item must be a valid index into the `constant_pool` table. The `constant_pool` entry at that index must be a `CONSTANT_Utf8_info` structure representing the binary name (JLS 13.1) of the type of the enum constant represented by this `element_value` structure.

const_name_index

The value of the `const_name_index` item must be a valid index into the `constant_pool` table. The `constant_pool` entry at that index must be a `CONSTANT_Utf8_info` structure representing the simple name of the enum constant represented by this `element_value` structure.

class_info_index

The `class_info_index` item is used if the `tag` item is 'c'. The `class_info_index` item must be a valid index into the `constant_pool` table. The `constant_pool` entry at that index must be a `CONSTANT_Utf8_info` structure representing the return descriptor (§4.4.3) of the type that is reified by the class represented by this `element_value` structure (e.g., 'V' for `Void.class`, 'Ljava/lang/Object;' for `Object`, etc.)

annotation_value

The `annotation_value` item is used if the `tag` item is '@'. The `element_value` structure represents a "nested" annotation.

array_value

The `array_value` item is used if the `tag` item is '['. The `array_value` item consists of the following two items:

num_values

The value of the `num_values` item gives the number of elements in the array-typed value represented by this `element_value` structure. Note that a maximum of 65535 elements are permitted in an array-typed element value.

`values`

Each value of the `values` table gives the value of an element of the array-typed value represented by this `element_value` structure.

4.8.17 The `RuntimeInvisibleAnnotations` attribute

The `RuntimeInvisibleAnnotations` attribute is similar to the `RuntimeVisibleAnnotations` attribute, except that the annotations represented by a `RuntimeInvisibleAnnotations` attribute must not be made available for return by reflective APIs, unless the JVM has been instructed to retain these annotations via some implementation-specific mechanism such as a command line flag. In the absence of such instructions, the JVM ignores this attribute.

The `RuntimeInvisibleAnnotations` attribute is a variable length attribute in the attributes table of the `ClassFile`, `field_info`, and `method_info` structures. The `RuntimeInvisibleAnnotations` attribute records runtime-invisible Java programming language annotations on the corresponding class, method, or field. Each `ClassFile`, `field_info`, and `method_info` structure may contain at most one `RuntimeInvisibleAnnotations` attribute, which records all the runtime-invisible Java programming language annotations on the corresponding program element.

The `RuntimeInvisibleAnnotations` attribute has the following format:

```
RuntimeInvisibleAnnotations_attribute {
    u2 attribute_name_index;
    u4 attribute_length;
    u2 num_annotations;
    annotation annotations[num_annotations];
}
```

The items of the `RuntimeInvisibleAnnotations` structure are as follows:

`attribute_name_index`

The value of the `attribute_name_index` item must be a valid index into the `constant_pool` table. The `constant_pool` entry at that index must be a

CONSTANT_Utf8_info structure representing the string "RuntimeInvisibleAnnotations".

`attribute_length`

The value of the `attribute_length` item indicates the length of the attribute, excluding the initial six bytes. The value of the `attribute_length` item is thus dependent on the number of runtime-invisible annotations represented by the structure, and their values.

`num_annotations`

The value of the `num_annotations` item gives the number of runtime-invisible annotations represented by the structure. Note that a maximum of 65535 runtime-invisible Java programming language annotations may be directly attached to a program element.

`annotations`

Each value of the `annotations` table represents a single runtime-invisible annotation on a program element.

4.8.18 The RuntimeVisibleParameterAnnotations attribute

The `RuntimeVisibleParameterAnnotations` attribute is a variable length attribute in the attributes table of the `method_info` structure. The `RuntimeVisibleParameterAnnotations` attribute records runtime-visible Java programming language annotations on the parameters of the corresponding method. Each `method_info` structure may contain at most one `RuntimeVisibleParameterAnnotations` attribute, which records all the runtime-visible Java programming language annotations on the parameters of the corresponding method. The JVM must make these annotations available so they can be returned by the appropriate reflective APIs.

The `RuntimeVisibleParameterAnnotations` attribute has the following format:

```

RuntimeVisibleParameterAnnotations_attribute {
    u2 attribute_name_index;
    u4 attribute_length;
    u1 num_parameters;
    {
        u2 num_annotations;
        annotation annotations[num_annotations];
    } parameter_annotations[num_parameters];
}

```

The items of the `RuntimeVisibleParameterAnnotations` structure are as follows:

`attribute_name_index`

The value of the `attribute_name_index` item must be a valid index into the `constant_pool` table. The `constant_pool` entry at that index must be a `CONSTANT_Utf8_info` structure representing the string `"RuntimeVisibleParameterAnnotations"`.

`attribute_length`

The value of the `attribute_length` item indicates the length of the attribute, excluding the initial six bytes. The value of the `attribute_length` item is thus dependent on the number of parameters, the number of runtime-visible annotations on each parameter, and their values.

`num_parameters`

The value of the `num_parameters` item gives the number of parameters of the method represented by the `method_info` structure on which the annotation occurs. (This duplicates information that could be extracted from the method descriptor.)

`parameter_annotations`

Each value of the `parameter_annotations` table represents all of the runtime-visible annotations on a single parameter. The sequence of values in the table corresponds to the sequence of parameters in the method signature. Each `parameter_annotations` entry contains the following two items:

`num_annotations`

The value of the `num_annotations` item indicates the number of runtime-visible annotations on the parameter corresponding to the sequence number of this `parameter_annotations` element.

`annotations`

Each value of the annotations table represents a single runtime-visible annotation on the parameter corresponding to the sequence number of this `parameter_annotations` element.

4.8.19 The `RuntimeInvisibleParameterAnnotations` attribute

The `RuntimeInvisibleParameterAnnotations` attribute is similar to the `RuntimeVisibleParameterAnnotations` attribute, except that the annotations represented by a `RuntimeInvisibleParameterAnnotations` attribute must not be made available for return by reflective APIs, unless the JVM has specifically been instructed to retain these annotations via some implementation-specific mechanism such as a command line flag. In the absence of such instructions, the JVM ignores this attribute.

The `RuntimeInvisibleParameterAnnotations` attribute is a variable length attribute in the attributes table of the `method_info` structure. The `RuntimeInvisibleParameterAnnotations` attribute records runtime-invisible Java programming language annotations on the parameters of the corresponding method. Each `method_info` structure may contain at most one `RuntimeInvisibleParameterAnnotations` attribute, which records all the runtime-invisible Java programming language annotations on the parameters of the corresponding method.

The `RuntimeInvisibleParameterAnnotations` attribute has the following format:

```
RuntimeInvisibleParameterAnnotations_attribute {  
    u2 attribute_name_index;  
    u4 attribute_length;  
    u1 num_parameters;  
    {  
        u2 num_annotations;  
        annotation annotations[num_annotations];  
    } parameter_annotations[num_parameters];  
}
```

The items of the `RuntimeInvisibleParameterAnnotations` structure are as follows:

attribute_name_index

The value of the `attribute_name_index` item must be a valid index into the `constant_pool` table. The `constant_pool` entry at that index must be a `CONSTANT_Utf8_info` structure representing the string "RuntimeInvisibleParameterAnnotations".

attribute_length

The value of the `attribute_length` item indicates the length of the attribute, excluding the initial six bytes. The value of the `attribute_length` item is thus dependent on the number of parameters, the number of runtime-invisible annotations on each parameter, and their values.

num_parameters

The value of the `num_parameters` item gives the number of parameters of the method represented by the `method_info` structure on which the annotation occurs. (This duplicates information that could be extracted from the method descriptor.)

parameter_annotations

Each value of the `parameter_annotations` table represents all of the runtime-invisible annotations on a single parameter. The sequence of values in the table corresponds to the sequence of parameters in the method signature. Each `parameter_annotations` entry contains the following two items:

num_annotations

The value of the `num_annotations` item indicates the number of runtime-invisible annotations on the parameter corresponding to the sequence number of this `parameter_annotations` element.

annotations

Each value of the `annotations` table represents a single runtime-invisible annotation on the parameter corresponding to the sequence number of this `parameter_annotations` element.

4.8.20 The AnnotationDefault attribute

The `AnnotationDefault` attribute is a variable length attribute in the attributes table of certain `method_info` structures, namely those representing elements of annotation types. The `AnnotationDefault` attribute records the default value for the element represented by the `method_info` structure. Each `method_info` structures representing an element of an annotation types may contain at most one Anno-

`AnnotationDefault` attribute. The JVM must make this default value available so it can be applied by appropriate reflective APIs.

The `AnnotationDefault` attribute has the following format:

```
AnnotationDefault_attribute {
    u2 attribute_name_index;
    u4 attribute_length;
    element_value default_value;
}
```

The items of the `AnnotationDefault` structure are as follows:

`attribute_name_index`

The value of the `attribute_name_index` item must be a valid index into the `constant_pool` table. The `constant_pool` entry at that index must be a `CONSTANT_Utf8_info` structure representing the string "AnnotationDefault".

`attribute_length`

The value of the `attribute_length` item indicates the length of the `attribute`, excluding the initial six bytes. The value of the `attribute_length` item is thus dependent on the default value.

`default_value`

The `default_value` item represents the default value of the `annotation type element` whose default value is represented by this `AnnotationDefault` attribute.

4.9 Format Checking

When a prospective `class` file is loaded (§2.17.2)(§2.17.2) by the Java virtual machine, the Java virtual machine first ensures that the file has the basic format of a `class` file. This process is known as *format checking*. The first four bytes must contain the right magic number. All recognized attributes must be of the proper length. The `class` file must not be truncated or have extra bytes at the end. The constant pool must not contain any superficially unrecognizable information.

This check for basic `class` file integrity is necessary for any interpretation of the `class` file contents.

However, format checking is distinct from verification. Historically, the two have been confused, because both are a form of integrity check.

4.10 Constraints on Java Virtual Machine Code

The Java virtual machine code for a method, instance initialization method (§3.9)(§3.9), or class or interface initialization method (§3.9)(§3.9) is stored in the code array of the `Code` attribute of a `method_info` structure of a `class` file. This section describes the constraints associated with the contents of the `Code_attribute` structure.

4.10.1 Static Constraints

The *static constraints* on a `class` file are those defining the well-formedness of the file. With the exception of the static constraints on the Java virtual machine code of the `class` file, these constraints have been given in the previous ~~section~~sections. The static constraints on the Java virtual machine code in a `class` file specify how Java virtual machine instructions must be laid out in the code array and what the operands of individual instructions must be.

The static constraints on the instructions in the code array are as follows:

- The code array must not be empty, so the `code_length` item cannot have the value 0.
- The value of the `code_length` item must be less than 65536.
- The opcode of the first instruction in the code array begins at index 0.
- Only instances of the instructions documented in ~~Section 6.4~~Section 6.4 may appear in the code array. Instances of instructions using the reserved opcodes (§6.2)(§6.2) or any opcodes not documented in this specification must not appear in the code array.
- If the class file version number is 51.0 or above, then neither the `jsr` opcode or the `jsr_w` opcode may appear in the code array.

- For each instruction in the code array except the last, the index of the opcode of the next instruction equals the index of the opcode of the current instruction plus the length of that instruction, including all its operands. The *wide* instruction is treated like any other instruction for these purposes; the opcode specifying the operation that a *wide* instruction is to modify is treated as one of the operands of that *wide* instruction. That opcode must never be directly reachable by the computation.
- The last byte of the last instruction in the code array must be the byte at index `code_length-1`.

The static constraints on the operands of instructions in the code array are as follows:

- The target of each jump and branch instruction (*jsr*, *jsr_w*, *goto*, *goto_w*, *ifeq*, *ifne*, *ifle*, *iflt*, *ifge*, *ifgt*, *ifnull*, *ifnonnull*, *if_icmpeq*, *if_icmpne*, *if_icmple*, *if_icmplt*, *if_icmpge*, *if_icmpgt*, *if_acmpeq*, *if_acmpne*) must be the opcode of an instruction within this method. The target of a jump or branch instruction must never be the opcode used to specify the operation to be modified by a *wide* instruction; a jump or branch target may be the *wide* instruction itself.
- Each target, including the default, of each *tableswitch* instruction must be the opcode of an instruction within this method. Each *tableswitch* instruction must have a number of entries in its jump table that is consistent with the value of its *low* and *high* jump table operands, and its *low* value must be less than or equal to its *high* value. No target of a *tableswitch* instruction may be the opcode used to specify the operation to be modified by a *wide* instruction; a *tableswitch* target may be a *wide* instruction itself.
- Each target, including the default, of each *lookupswitch* instruction must be the opcode of an instruction within this method. Each *lookupswitch* instruction must have a number of *match-offset* pairs that is consistent with the value of its *npairs* operand. The *match-offset* pairs must be sorted in increasing numerical order by signed *match* value. No target of a *lookupswitch* instruction may be the opcode used to specify the operation to be modified by a *wide* instruction; a *lookupswitch* target may be a *wide* instruction itself.
- The operand of each *ldc* instruction must be a valid index into the `constant_pool` table. The operands of each *ldc_w* instruction must represent a valid index into the `constant_pool` table. In both cases the constant pool

entry referenced by that index must be of type ~~CONSTANT_Class~~, ~~CONSTANT_Integer~~, ~~CONSTANT_Float~~, or ~~CONSTANT_String~~ if the class file version number is less than 49.0.0. If the class file version is 49.0 or above, then the constant pool entry referenced by the entry must be of type ~~CONSTANT_Class~~, ~~CONSTANT_Integer~~, ~~CONSTANT_Float~~, ~~or~~ ~~CONSTANT_String~~ or ~~CONSTANT_Class~~.

- The operands of each *ldc2_w* instruction must represent a valid index into the `constant_pool` table. The constant pool entry referenced by that index must be of type `CONSTANT_Long` or `CONSTANT_Double`. In addition, the subsequent constant pool index must also be a valid index into the constant pool, and the constant pool entry at that index must not be used.
- The operands of each *getfield*, *putfield*, *getstatic*, and *putstatic* instruction must represent a valid index into the `constant_pool` table. The constant pool entry referenced by that index must be of type `CONSTANT_Fieldref`.
- The indexbyte operands of each *invokevirtual*, *invokespecial*, and *invokestatic* instruction must represent a valid index into the `constant_pool` table. The constant pool entry referenced by that index must be of type `CONSTANT_Methodref`.
- Only the *invokespecial* instruction is allowed to invoke an instance initialization method (§3.9)(§3.9). No other method whose name begins with the character '<' ('\u003c') may be called by the method invocation instructions. In particular, the class or interface initialization method specially named `<clinit>` is never called explicitly from Java virtual machine instructions, but only implicitly by the Java virtual machine itself.
- The indexbyte operands of each *invokeinterface* instruction must represent a valid index into the `constant_pool` table. The constant pool entry referenced by that index must be of type `CONSTANT_InterfaceMethodref`. The value of the *count* operand of each *invokeinterface* instruction must reflect the number of local variables necessary to store the arguments to be passed to the interface method, as implied by the descriptor of the `CONSTANT_NameAndType_info` structure referenced by the `CONSTANT_InterfaceMethodref` constant pool entry. The fourth operand byte of each *invokeinterface* instruction must have the value zero.
- The operands of each *instanceof*, *checkcast*, *new*, and *newarray* instruction and the indexbyte operands of each *multianewarray* instruction must repre-

sent a valid index into the `constant_pool` table. The constant pool entry referenced by that index must be of type `CONSTANT_Class`.

- No *anewarray* instruction may be used to create an array of more than 255 dimensions.
- No *new* instruction may reference a `CONSTANT_Class` `constant_pool` table entry representing an array class. The *new* instruction cannot be used to create an array.
- A *multianewarray* instruction must be used only to create an array of a type that has at least as many dimensions as the value of its *dimensions* operand. That is, while a *multianewarray* instruction is not required to create all of the dimensions of the array type referenced by its *indexbyte* operands, it must not attempt to create more dimensions than are in the array type. The *dimensions* operand of each *multianewarray* instruction must not be zero.
- The *atype* operand of each *newarray* instruction must take one of the values `T_BOOLEAN` (4), `T_CHAR` (5), `T_FLOAT` (6), `T_DOUBLE` (7), `T_BYTE` (8), `T_SHORT` (9), `T_INT` (10), or `T_LONG` (11).
- The index operand of each *iload*, *fload*, *aload*, *istore*, *fstore*, *astore*, *iinc*, and *ret* instruction must be a nonnegative integer no greater than `max_locals-1`.
- The implicit index of each *iload*_*<n>*, *fload*_*<n>*, *aload*_*<n>*, *istore*_*<n>*, *fstore*_*<n>*, and *astore*_*<n>* instruction must be no greater than the value of `max_locals-1`.
- The index operand of each *lload*, *dload*, *lstore*, and *dstore* instruction must be no greater than the value of `max_locals-2`.
- The implicit index of each *lload*_*<n>*, *dload*_*<n>*, *lstore*_*<n>*, and *dstore*_*<n>* instruction must be no greater than the value of `max_locals-2`.
- The *indexbyte* operands of each *wide* instruction modifying an *iload*, *fload*, *aload*, *istore*, *fstore*, *astore*, *ret*, or *iinc* instruction must represent a nonnegative integer no greater than `max_locals-1`. The *indexbyte* operands of each *wide* instruction modifying an *lload*, *dload*, *lstore*, or *dstore* instruction must represent a nonnegative integer no greater than `max_locals-2`.

4.10.2 Structural Constraints

The structural constraints on the code array specify constraints on relationships between Java virtual machine instructions. The structural constraints are as follows:

- Each instruction must only be executed with the appropriate type and number of arguments in the operand stack and local variable array, regardless of the execution path that leads to its invocation. An instruction operating on values of type `int` is also permitted to operate on values of type `boolean`, `byte`, `char`, and `short`. (As noted in §3.3.4 and §3.11.1, the Java virtual machine internally converts values of types `boolean`, `byte`, `char`, and `short` to type `int`.)
- If an instruction can be executed along several different execution paths, the operand stack must have the same depth (§3.6.2) prior to the execution of the instruction, regardless of the path taken.
- At no point during execution can the order of the local variable pair holding a value of type `long` or `double` be reversed or the pair split up. At no point can the local variables of such a pair be operated on individually.
- No local variable (or local variable pair, in the case of a value of type `long` or `double`) can be accessed before it is assigned a value.
- At no point during execution can the operand stack grow to a depth (§3.6.2) greater than that implied by the `max_stack` item.
- At no point during execution can more values be popped from the operand stack than it contains.
- Each *invokespecial* instruction must name an instance initialization method (§3.9), a method in the current class, or a method in a superclass of the current class.
- When the instance initialization method (§3.9) is invoked, an uninitialized class instance must be in an appropriate position on the operand stack. An instance initialization method must never be invoked on an initialized class instance.
- When any instance method is invoked or when any instance variable is accessed, the class instance that contains the instance method or instance variable must already be initialized.

- There must never be an uninitialized class instance on the operand stack or in a local variable when any backwards branch is taken.
- There must never be an uninitialized class instance on the operand stack or in a local variable when a *jsr* or *jsr_w* instruction is executed.
- Each instance initialization method (§3.9)(§3.9), except for the instance initialization method derived from the constructor of class `Object`, must call either another instance initialization method of `this` or an instance initialization method of its direct superclass `super` before its instance members are accessed. However, instance fields of `this` that are declared in the current class may be assigned before calling any instance initialization method.
- The arguments to each method invocation must be method invocation compatible (§2.6.8)(§2.6.8) with the method descriptor (§4.4.3).
- The type of every class instance that is the target of a method invocation instruction must be assignment compatible (§2.6.7) with the class or interface type specified in the instruction. In addition, the type of the target of an *invokespecial* instruction must be assignment compatible with the current class, unless an instance initialization method is being invoked.
- ~~The type of every class instance that is the target of a method invocation instruction must be assignment compatible (§2.6.7) with the class or interface type specified in the instruction. In addition, the type of the target of an *invokespecial* instruction must be assignment compatible with the current class, unless an instance initialization method is being invoked.~~ Each return instruction must match its method's return type. If the method returns a `boolean`, `byte`, `char`, `short`, or `int`, only the *ireturn* instruction may be used. If the method returns a `float`, `long`, or `double`, only an *freturn*, *lreturn*, or *dreturn* instruction, respectively, may be used. If the method returns a reference type, it must do so using an *areturn* instruction, and the type of the returned value must be assignment compatible (§2.6.7)(§2.6.7) with the return descriptor (§4.4.3) of the method. All instance initialization methods, class or interface initialization methods, and methods declared to return `void` must use only the *return* instruction.
- If *getfield* or *putfield* is used to access a protected field of a superclass that is a member of a different runtime package than the current class, then the type of the class instance being accessed must be the same as or a subclass of the current class. If *invokevirtual* or *invokespecial* is used to access a

protected method of a superclass that is a member of different runtime package than the current class, then the type of the class instance being accessed must be the same as or a subclass of the current class

- The type of every class instance accessed by a *getfield* instruction or modified by a *putfield* instruction must be assignment compatible (§2.6.7)(§2.6.7) with the class type specified in the instruction.
- The type of every value stored by a *putfield* or *putstatic* instruction must be compatible with the descriptor of the field (§4.4.2) of the class instance or class being stored into. If the descriptor type is `boolean`, `byte`, `char`, `short`, or `int`, then the value must be an `int`. If the descriptor type is `float`, `long`, or `double`, then the value must be a `float`, `long`, or `double`, respectively. If the descriptor type is a reference type, then the value must be of a type that is assignment compatible (§2.6.7)(§2.6.7) with the descriptor type.
- The type of every value stored into an array by an *aastore* instruction must be a reference type. The component type of the array being stored into by the *aastore* instruction must also be a reference type.
- Each *athrow* instruction must throw only values that are instances of class `Throwable` or of subclasses of `Throwable`. Each class mentioned in a `catch_type` item of a method's exception table must be `Throwable` or of subclasses of `Throwable`.
- Execution never falls off the bottom of the code array.
- No return address (a value of type `returnAddress`) may be loaded from a local variable.
- The instruction following each *jsr* or *jsr_w* instruction may be returned to only by a single *ret* instruction.
- No *jsr* or *jsr_w* instruction may be used to recursively call a subroutine if that subroutine is already present in the subroutine call chain. (Subroutines can be nested when using try-finally constructs from within a finally clause. ~~For more information on Java virtual machine subroutines, see §4.11.1.6.)~~
- Each instance of type `returnAddress` can be returned to at most once. If a *ret* instruction returns to a point in the subroutine call chain above the *ret* instruction corresponding to a given instance of type `returnAddress`, then that instance can never be used as a return address.

4.11 Verification of *class* Files

Even though any compiler for the Java programming language must only produce *class* files that satisfy all the static [and structural](#) constraints in the previous sections, the Java virtual machine has no guarantee that any file it is asked to load was generated by that compiler or is properly formed. Applications such as web browsers do not download source code, which they then compile; these applications download already-compiled *class* files. The browser needs to determine whether the *class* file was produced by a trustworthy compiler or by an adversary attempting to exploit the virtual machine.

An additional problem with compile-time checking is version skew. A user may have successfully compiled a class, say `PurchaseStockOptions`, to be a subclass of `TradingClass`. But the definition of `TradingClass` might have changed since the time the class was compiled in a way that is not compatible with preexisting binaries. Methods might have been deleted or had their return types or modifiers changed. Fields might have changed types or changed from instance variables to class variables. The access modifiers of a method or variable may have changed from `public` to `private`. For a discussion of these issues, see Chapter 13, “Binary Compatibility,” in the *The JavaTM Language Specification*.

Because of these potential problems, the Java virtual machine needs to verify for itself that the desired constraints are satisfied by the *class* files it attempts to incorporate. A Java virtual machine implementation verifies that each *class* file satisfies the necessary constraints at linking time (~~§2.17.3~~)([§2.17.3](#)).

Linking-time verification enhances the performance of the interpreter. Expensive checks that would otherwise have to be performed to verify constraints at run time for each interpreted instruction can be eliminated. The Java virtual machine can assume that these checks have already been performed. For example, the Java virtual machine will already know the following:

- There are no operand stack overflows or underflows.
- All local variable uses and stores are valid.
- The arguments to all the Java virtual machine instructions are of valid types.

The verifier also performs verification that can be done without looking at the code array of the Code attribute (§4.8.3). The checks performed include the following:

- Ensuring that `final` classes are not subclassed and that `final` methods are not overridden.
- Checking that every class (except `Object`) has a direct superclass.
- Ensuring that the constant pool satisfies the documented static constraints: for example, that each `CONSTANT_Class_info` structure in the constant pool contains in its `name_index` item a valid constant pool index for a `CONSTANT_Utf8_info` structure.
- Checking that all field references and method references in the constant pool have valid names, valid classes, and a valid type descriptor.

Note that these ~~check~~-checks do not ensure that the given field or method actually exists in the given class, nor ~~does it do they~~ check that the type descriptors given refer to real classes. They ensure only that these items are well formed. More detailed checking is performed when the ~~byte-codes~~-bytecodes themselves are verified, and during resolution.

There are two strategies that Java virtual machines may use for verification. Verification by type checking must be used to verify class files whose version number is greater than or equal to 50.0.

Verification by type inference must be supported by all Java virtual machines, except those conforming to the JavaCard and J2ME CLDC profiles, in order to verify class files whose version number is less than 50.0. Verification on virtual machines supporting the J2ME CLDC and JavaCard profiles is governed by their respective specifications.

4.11.1 Verification by Type ~~Inference~~Checking

Class files whose version number is greater than or equal to 50.0, must be verified using the typechecking rules given in this section. If, and only if, the class file's major version number equals 50.0, then if the typechecking fails, a virtual machine implementation may choose to attempt to perform verification by type inference.

This is a pragmatic adjustment, designed to ease the transition to the new verification discipline.

Many tools that manipulate class files may alter the bytecodes of a method in a manner that requires adjustment of the method's stack frame maps. If a tool does not make the necessary adjustments to the stack frame maps, typechecking may fail, even though the bytecode is in principle valid (and would consequently verify under the old type inference scheme).

The need to adjust stack frame maps is a new requirement for tools that operate on class files and/or bytecode streams. To allow the implementors of such tools more time to adapt, virtual machines are allowed to fall back to the older verification discipline, but only for a limited time.

In cases where typechecking fails but type inference is invoked and succeeds, a certain performance penalty is expected. Such a penalty is unavoidable. It also should serve as a signal to tool vendors that their output needs to be adjusted, and provides vendors with additional incentive to make these adjustments.

If a virtual machine implementation ever attempts to perform verification by type inference on version 50.0 classfiles, it must do so in all cases where verification by typechecking fails.

This means that that a virtual machine cannot choose to resort to type inference in once case and not in another. It must either reject class files that do not verify via typechecking, or else consistently failover to the type inferencing verifier whenever typechecking fails.

The type checker requires a list of stack frame maps for each method with a Code attribute. The type checker reads the stack frame maps for each such method and uses these maps to generate a proof of the type safety of the instructions in the Code attribute. The list of stack frame maps is given by the StackMapTable (§4.8.4) attribute of the Code attribute.

The intent is that a stack frame map must appear at the beginning of each basic block in a method. The stack frame map specifies the verification type of each operand stack entry and of each local variable at the start of each basic block.

The type rules that the typechecker enforces are specified by means of Prolog clauses. English language text is used to describe the type rules in an informal way, while the Prolog code provides a formal specification.

If the predicate `classIsTypeSafe` is not true, the type checker must throw the exception `VerifyError` to indicate that the class file is malformed. Otherwise, the class file has type checked successfully and bytecode verification has completed successfully.

`classIsTypeSafe(Class) :-`

`classClassName(Class, Name),`
`classDefiningLoader(Class, L),`
`superclassChain(Name, L, Chain),`
`Chain \= [].`
`classSuperClassName(Class, SuperclassName),`
`loadedClass(SuperclassName, L, Superclass),`
`classIsNotFinal(Superclass),`
`classMethods(Class, Methods),`
`checklist(methodIsTypeSafe(Class), Methods).`

`classIsTypeSafe(Class) :-`

`classClassName(Class, 'java/lang/Object'),`
`classDefiningLoader(Class, L),`
`isBootstrapClassLoader(L),`
`classMethods(Class, Methods),`
`checklist(methodIsTypeSafe(Class), Methods).`

Thus, a class is type safe if all its methods are type safe, and it does not subclass a final class.

The predicate `classIsTypeSafe` assumes that `Class` is a Prolog term representing a binary class that has been successfully parsed and loaded. This specification does not mandate the precise structure of this term, but does require that certain predicates (e.g., `classMethods`) be defined upon it, as specified in Section 4.11.1.3.1.

For example, we assume a predicate `classMethods(Class, Methods)` that, given a term representing a class as described above as its first argument, binds its second argument to a list comprising all the methods of the class, represented in a convenient form described below.

We also require the existence of a predicate `loadedClass(Name, InitiatingLoader, ClassDefinition)` which asserts that there exists a class named `Name` whose representation (in accordance with this specification) when loaded by the class loader `InitiatingLoader` is `ClassDefinition`. additional required predicates are discussed in §4.11.1.3.1.

Individual instructions are presented as terms whose functor is the name of the instruction and whose arguments are its parsed operands.

For example, an `aload` instruction is represented as the term `aload(N)`, which includes the index `N` that is the operand of the instruction.

A few instructions have operands that are constant pool entries representing methods or fields. As specified in §4.4.3, methods are represented by `CONSTANT_InterfaceMethodref_info` (for interface methods) or `CONSTANT_Methodref_info` (for other methods) structures in the constant pool. Such structures are represented here as functor applications of the form `imethod(MethodClassName, MethodName, MethodDescriptor)` (for interface methods) or `method(MethodClassName, MethodName, MethodDescriptor)` (for other methods), where `MethodClassName` is the name of the class referenced by the `class_index` item for the structure, and `MethodName` and `MethodDescriptor` correspond to the name and type descriptor referenced by the `name_and_type_index` of the structure.

Similarly, fields are represented by `CONSTANT_Fieldref_info` structures in the class file. These structures are represented here as functor applications of the form `field(FieldClassName, FieldName, FieldDescriptor)` where `FieldClassName` is the name of the class referenced by the `class_index` item in the structure, and `FieldName` and `FieldDescriptor` correspond to the name and type descriptor referenced by the `name_and_type_index` item of the structure. For clarity, we assume that type descriptors are mapped into more readable names: the leading `L` and trailing `;` are dropped from class names, and the base type characters used for primitive types are mapped to the names of those types).

So, a `getfield` instruction whose operand was an index into the constant pool that refers to a field `foo` of type `F` in class `Bar` would be represented as `getfield(field('Bar', 'foo', 'F'))`.

Constant pool entries that refer to constant values, such as `CONSTANT_String`, `CONSTANT_Integer`, `CONSTANT_Float`, `CONSTANT_Long`, `CONSTANT_Double` and `CONSTANT_Class`, are encoded via the functors whose names are `string`, `int`, `float`, `long`, `double` and `classConstant` respectively.

So an `ldc` instruction for loading the integer 91 would be encoded as `ldc(int(91))`.

The instructions as a whole are represented as a list of terms of the form `instruction(Offset, AnInstruction)`.

For example `instruction(21, aload(1))`.

The order of instructions in this list must be the same as in the class file.

Stack frame maps are represented as a list of terms of the form `stackMap(Offset, TypeState)` where `Offset` is an integer indicating the offset of the instruction the frame map applies to, and `TypeState` is the expected incoming type state for that instruction. The order of instructions in this list must be the same as in the class file.

`TypeState` has the form `frame(Locals, OperandStack, Flags)`.

Locals is a list of verification types, such that the *N*th element of the list (with 0 based indexing) represents the type of local variable *N*.

If any local variable in Locals has the type uninitializedThis, Flags is [flagThisUninit], otherwise it is an empty list.

OperandStack is a list of types, such that the first element represents the type of the top of the operand stack, and the elements below the top follow in the appropriate order.

However, note again that types of size 2 are represented by two entries, with the first entry being top and the second one being the type itself.

So a stack with a `double`, an `int` and a `long` would be represented as `[top, double, int, top, long]`.

Array types are represented by applying the functor `arrayOf` to an argument denoting the component type of the array. Other reference types are represented using the functor `class`. Hence `class(N, L)` represents the class whose binary name is *N* as loaded by the loader *L*.

Thus, *L* is an initiating loader of the class represented by `class(N, L)`. It may, or may not, be its defining loader.

The type `uninitialized(offset)` is represented by applying the functor `uninitialized` to an argument representing the numerical value of the *offset*. Other verification types are represented by Prolog atoms whose name denotes the verification type in question.

So, the class `Object` would be represented as `class('java/lang/Object', BL)`, where *BL* is the bootstrap loader. The types `int[]` and `Object[]` would be represented by `arrayOf(int)` and `arrayOf(class('java/lang/Object', BL))` respectively.

Flags is a list which may either be empty or have the single element `flagThisUninit`.

This flag is used in constructors, to mark type states where initialization of `this` has not yet been completed. In such type states, it is illegal to return from the method.

4.11.1.1 The Type hierarchy

The typechecker enforces a type system based upon a hierarchy of verification types, illustrated in figure 1 below. Most verifier types have a direct correspondence with Java virtual machine field type descriptors as given in Table 4.2. The only exceptions are the field descriptors B, C, S and Z all of which correspond to the verifier type `int`.

4.11.1.2 Subtyping Rules

Subtyping is reflexive

`isAssignable(X, X).`

`isAssignable(oneWord, top).`

`isAssignable(twoWord, top).`

`isAssignable(int, X) :- isAssignable(oneWord, X).`

`isAssignable(float, X) :- isAssignable(oneWord, X).`

`isAssignable(long, X) :- isAssignable(twoWord, X).`

`isAssignable(double, X) :- isAssignable(twoWord, X).`

`isAssignable(reference, X) :- isAssignable(oneWord, X).`

`isAssignable(uninitialized, X) :- isAssignable(reference, X).`

These subtype rules are not necessarily the most obvious formulation of subtyping. There is a clear split between subtyping rules for reference

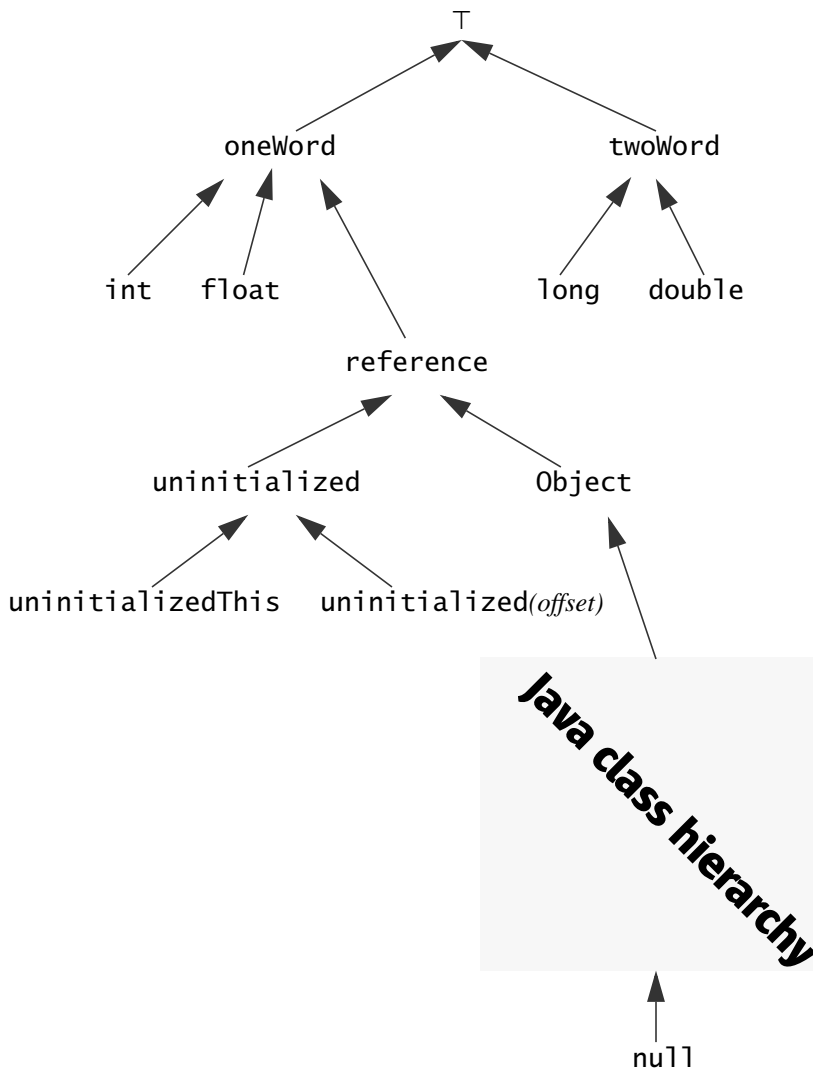


Figure 1: The verification type Hierarchy

[types in the Java programming language, and rules for the remaining verification types.](#)

Subtype rules for the reference types in the Java programming language are specified recursively in the obvious way. The remaining verification types have subtypes rules of the form:

subtype(v, X) :- subtype(the_direct_supertype_of_v, X).

That is, v is a subtype of X if the direct supertype of v is a subtype of X.

Note that the actual subtype rule is called isAssignable in the Prolog code below.

We also have a rule that says subtyping is reflexive, so together these rules cover most verification types that are not reference types in the Java programming language.

The aforementioned split allows us to state general subtyping relations between the Java programming language types and other verification types.

These relations hold independently of the Java type's position in the hierarchy. These rules are useful for the reference implementation, where they help prevent duplicate answers and excessive class loading.

For example, we don't want to start climbing up the class hierarchy in response to a query of the form class(foo, L) <: twoWord. If we use the same predicates for the entire hierarchy, we cannot restrict the climb to cases where we compare two classes.

It would be nicer to have more uniform rules for the specification, but they are not well suited for the reference implementation. We'd like the reference implementation to be as close to the specification as possible, so this is a reasonable compromise.

isAssignable(class(,), X) :- isAssignable(reference, X).

isAssignable(arrayOf(), X) :- isAssignable(reference, X).

isAssignable(uninitializedThis, X) :- isAssignable(uninitialized, X).

isAssignable(uninitialized(), X) :- isAssignable(uninitialized, X).

isAssignable(null, class(,)).

isAssignable(null, arrayOf()).

isAssignable(null, X) :-

isAssignable(class('java/lang/Object', BL), X).

isBootstrapLoader(BL).

isAssignable(class(X, Lx), class(Y, Ly)) :-

isJavaAssignable(class(X, Lx), class(Y, Ly)).

isAssignable(arrayOf(X), class(Y, L)) :-

isJavaAssignable(arrayOf(X), class(Y, L)).

isAssignable(arrayOf(X), arrayOf(Y)) :-

isJavaAssignable(arrayOf(X), arrayOf(Y)).

For assignments, interfaces are treated like Object.

isJavaAssignable(class(_), class(To, L)) :-

loadedClass(To, L, ToClass).

___ classIsInterface(ToClass).

isJavaAssignable(From, To) :-

___ isJavaSubclassOf(From, To).

Arrays are subtypes of Object.

isJavaAssignable(arrayOf(_), class('java/lang/Object', BL)) :-

isBootstrapLoader(BL).

The intent here is that array types are subtypes of Cloneable and java.io.Serializable.

isJavaAssignable(arrayOf(_), X) :-

isArrayInterface(X).

The subtyping relation between arrays of primitive type is the identity relation.

isJavaAssignable(arrayOf(X), arrayOf(Y)) :-
atom(X), atom(Y), X = Y.

Subtyping between arrays of reference type is covariant.

isJavaAssignable(arrayOf(X), arrayOf(Y)) :-
compound(X), compound(Y), isJavaAssignable(X, Y).

isArrayInterface(class('java/lang/Cloneable', BL)) :-
isBootstrapLoader(BL).

isArrayInterface(class('java/io/Serializable', BL)) :-
isBootstrapLoader(BL).

Subclassing is reflexive.

isJavaSubclassOf(class(SubClassName, L), class(SubClassName, L)).

isJavaSubclassOf(class(SubClassName, LSub),
class(SuperClassName, LSuper)) :-
superclassChain(SubClassName, LSub, Chain),
member(class(SuperclassName, L), Chain),
loadedClass(SuperClassName, L, Sup),
loadedClass(SuperClassName, LSuper, Sup).

sizeOf(X, 2) :- isAssignable(X, twoWord).

sizeOf(X, 1) :- isAssignable(X, oneWord).

sizeOf(top, 1).

Subtyping is extended pointwise to type states.

```

framesAssignable(frame(Locals1, StackMap1, Flags1),
    frame(Locals2, StackMap2, Flags2)) :-
    length(StackMap1, StackMapLength),
    length(StackMap2, StackMapLength),
    maplist(isAssignable, Locals1, Locals2),
    maplist(isAssignable, StackMap1, StackMap2),
    subset(Flags1, Flags2).

```

4.11.1.3 Typechecking Rules

4.11.1.3.1 Accessors

Stipulated Accessors: Throughout this specification, we assume the existence of certain Prolog predicates whose formal definitions are not given in the specification. In this section, we list these predicates and specify their expected behavior.

parseFieldDescriptor(Descriptor, Type)

Converts a field descriptor, **Descriptor**, into the corresponding verification type **Type** (see the beginning of Section 4.11.1.1 for the specification of this correspondence).

parseMethodDescriptor(Descriptor, ArgTypeList, ReturnType)

Converts a method descriptor, **Descriptor**, into a list of verification types, **ArgTypeList**, corresponding (Section 4.11.1.1) to the method argument types, and a verification type, **ReturnType**, corresponding to the return type.

parseCodeAttribute(Class, Method, FrameSize, MaxStack, ParsedCode, Handlers, StackMap)

Extracts the instruction stream, **ParsedCode**, of the method **Method** in **Class**, as well as the maximum operand stack size, **MaxStack**, the maximal number of local variables, **FrameSize**, the exception handlers, **Handlers**, and the stack map **StackMap**. The representation of the instruction stream and stack map attribute must be as specified in the beginning of §4.11.1 Each exception handler is represented by a functor application of the form **handler(Start, End, Target, ClassName)** whose arguments are, respectively, the start and end of the range of

instructions covered by the handler, the first instruction of the handler code, and the name of the exception class that this handler is designed to handle.

classClassName(Class, ClassName)

Extracts the name, ClassName, of the class Class.

classIsInterface(Class)

True iff the class, Class, is an interface.

classIsNotFinal(Class)

True iff the class, Class, is not a final class.

classSuperClassName(Class, SuperClassName)

Extracts the name, SuperClassName, of the superclass of class Class.

classInterfaces(Class, Interfaces)

Extracts a list, Interfaces, of the direct superinterfaces of the class Class.

classMethods(Class, Methods)

Extracts a list, Methods, of the methods declared in the class Class.

classAttributes(Class, Attributes)

Extracts a list, Attributes, of the attributes of the class Class. Each attribute is represented as a functor application of the form attribute(AttributeName, AttributeContents), where AttributeName is the name of the attribute. The format of the attributes contents is unspecified.

classDefiningLoader(Class, Loader)

Extracts the defining class loader, Loader, of the class Class.

isBootstrapLoader(Loader)

True iff the class loader Loader is the bootstrap class loader.

methodName(Method, Name)

Extracts the name, Name, of the method Method.

methodAccessFlags(Method, AccessFlags)

Extracts the access flags, AccessFlags, of the method Method.

methodDescriptor(Method, Descriptor)

Extracts the descriptor, Descriptor, of the method Method.

methodAttributes(Method, Attributes)

Extracts a list, Attributes, of the attributes of the method Method.

isNotFinal(Method, Class)

True iff Method in class Class is not final.

isProtected(MemberClass, MemberName, MemberDescriptor)

True iff there is a member named MemberName with descriptor MemberDescriptor in the class MemberClass and it is protected.

isNotProtected(MemberClass, MemberName, MemberDescriptor)

True iff there is a member named MemberName with descriptor MemberDescriptor in the class MemberClass and it is not protected.

samePackageName(Class1, Class2)

True iff the package names of Class1 and Class2 are tthe same.

differentPackageName(Class1, Class2)

True iff the package names of Class1 and Class2 are different.

There is a principle guiding the determination as to which accessors are fully specified and which are stipulated. We do not want to over-specify the representation of the class file. Providing specific accessors to the class or method term would force us to completely specify a format for the Prolog term representing the class file.

Specified Accessors and Utilities: We also provide accessor and utility rules that extract necessary information from the representation of the class and its methods.

An environment is a six-tuple consisting of a class, a method, the declared return type of the method, the instructions in a method, the maximal size of the operand stack, and a list of exception handlers.

maxOperandStackLength(Environment, MaxStack) :-

Environment = environment(Class, Method, ReturnType, Instructions, MaxStack, Handlers).

exceptionHandlers(Environment, Handlers) :-

Environment = environment(Class, Method, ReturnType,
Instructions, , Handlers).

thisMethodReturnType(Environment, ReturnType) :-

Environment = environment(Class, Method, ReturnType,
Instructions, ,).

thisClass(Environment, class(ClassName, L)) :-

Environment = environment(Class, Method, ReturnType,
Instructions, ,).

classDefiningLoader(Class, L).

classClassName(Class, ClassName).

allInstructions(Environment, Instructions) :-

Environment = environment(Class, Method, ReturnType,
Instructions, ,).

offsetStackFrame(Environment, Offset, StackFrame) :-

allInstructions(Environment, Instructions).

member(stackMap(Offset, StackFrame), Instructions).

currentClassLoader(Environment, Loader) :-

thisClass(Environment, class(, Loader)).

notMember(, []).

notMember(X, [A | More]) :- X \= A, notMember(X, More).

differentRuntimePackage(Class1, Class2) :-

sameRuntimePackage(Class1, Class2) :-

classDefiningLoader(Class1, L).

classDefiningLoader(Class2, L).

samePackageName(Class1, Class2).

differentRuntimePackage(Class1, Class2) :-
classDefiningLoader(Class1, L1),
classDefiningLoader(Class2, L2),
L1 \= L2.

differentRuntimePackage(Class1, Class2) :-
differentPackageName(Class1, Class2).

4.11.1.3.2 Abstract & Native Methods

Abstract methods are considered to be type safe if they do not override a final method.

methodIsTypeSafe(Class, Method) :-
doesNotOverrideFinalMethod(Class, Method),
 — methodAccessFlags(Method, AccessFlags),
member(abstract, AccessFlags).

Native methods are considered to be type safe if they do not override a final method.

methodIsTypeSafe(Class, Method) :-
doesNotOverrideFinalMethod(Class, Method),
methodAccessFlags(Method, AccessFlags),
member(native, AccessFlags).

doesNotOverrideFinalMethod(class('java/lang/Object', L), Method) :-
isBootstrapLoader(L).

doesNotOverrideFinalMethod(Class, Method) :-

classSuperClassName(Class, SuperclassName),
classDefiningLoader(Class, L),
loadedClass(SuperclassName, L, Superclass),
classMethods(Superclass, MethodList),
finalMethodNotOverridden(Method, Superclass, MethodList).

finalMethodNotOverridden(Method, Superclass, MethodList) :-
methodName(Method, Name),
methodDescriptor(Method, Descriptor),
member(method(, Name, Descriptor), MethodList),
isNotFinal(Method, Superclass).

finalMethodNotOverridden(Method, Superclass, MethodList) :-
methodName(Method, Name),
methodDescriptor(Method, Descriptor),
notMember(method(, Name, Descriptor), MethodList),
doesNotOverrideFinalMethod(Superclass, Method).

4.11.1.3.3 Checking Code

Non-abstract, non-native methods are type correct if they have code and the code is type correct.

methodsTypeSafe(Class, Method) :-
doesNotOverrideFinalMethod(Class, Method),
methodAccessFlags(Method, AccessFlags),
methodAttributes(Method, Attributes),
notMember(native, AccessFlags),
notMember(abstract, AccessFlags),
member(attribute('Code', _), Attributes),

— methodWithCodeIsTypeSafe(Class, Method).

A method with code is type safe if it is possible to merge the code and the stack frames into a single stream such that each stack map precedes the instruction it corresponds to, and the resulting merged stream is type correct.

methodWithCodeIsTypeSafe(Class, Method) :-

— parseCodeAttribute(Class, Method, FrameSize, MaxStack, ParsedCode, Handlers, StackMap).

— mergeStackMapAndCode(StackMap, ParsedCode, MergedCode).

— methodInitialStackFrame(Class, Method, FrameSize, StackFrame, ReturnType).

— Environment =

environment(Class, Method, ReturnType, MergedCode, MaxStack, Handlers).

— handlersAreLegal(Environment).

— mergedCodeIsTypeSafe(Environment, MergedCode, StackFrame).

The initial type state of a method consists of an empty operand stack and local variable types derived from the type of `this` and the arguments, as well as the appropriate flag, depending on whether this is an `<init>` method.

methodInitialStackFrame(Class, Method, FrameSize,

frame(Locals, [], Flags), ReturnType):-

— methodDescriptor(Method, Descriptor).

— parseMethodDescriptor(Descriptor, RawArgs, ReturnType).

— expandTypeList(RawArgs, Args).

— methodInitialThisType(Class, Method, ThisList).

— flags(ThisList, Flags).

— append(ThisList, Args, ThisArgs).

```

__   expandToLength(ThisArgs, FrameSize, top, Locals).

flags([uninitializedThis], [flagThisUninit]).
flags(X, []) :- X \= [uninitializedThis].

expandToLength(List, Size, Filler, List) :- length(List, Size).
expandToLength(List, Size, Filler, Result) :-
__   length(List, ListLength).
__   ListLength < Size,
__   Delta is Size - ListLength,
__   length(Extra, Delta).
__   checklist(=(Filler), Extra).
__   append(List, Extra, Result).

```

For a static method `this` is irrelevant; the list is empty. For an instance method, we get the type of `this` and put it in a list.

```

methodInitialThisType( Class, Method, []) :-
__   methodAccessFlags(Method, AccessFlags).
__   member(static, AccessFlags).
__   methodName(Method, MethodName).
__   MethodName \= '<init>'.
methodInitialThisType(Class, Method, [This]) :-
__   methodAccessFlags(Method, AccessFlags).
__   notMember(static, AccessFlags).
__   instanceMethodInitialThisType(Class, Method, This).

```

In the `<init>` method of Object the type of `this` is Object. In other `<init>` methods, the type of `this` is `uninitializedThis`. Otherwise, the type of `this` in an instance method is `class(N, L)`, where N is the name of the class containing the method and L is its defining class loader.

instanceMethodInitialThisType(Class, Method, class(ClassName, L)) :-

— methodName(Method, MethodName),
 — MethodName \= '<init>',
classDefiningLoader(Class, L),
 — classClassName(Class, ClassName).

instanceMethodInitialThisType(Class, Method,
class('java/lang/Object', L)) :-

— methodName(Method, '<init>'),
classDefiningLoader(Class, L),
isBootstrapLoader(L),
 — classClassName(Class, 'java/lang/Object').

instanceMethodInitialThisType(Class, Method, uninitializedThis) :-

— methodName(Method, '<init>'),
 — classClassName(Class, ClassName),
classDefiningLoader(Class, CurrentLoader),
superclassChain(ClassName, CurrentLoader, Chain),
 — Chain \= [].

Below are the rules for iterating through the code stream. The assumption is that the stream is a well formed mixture of instructions and stack maps, such that the stack map for bytecode index N appears just before instruction N .

The rules for building this mixed stream are given later, by the predicate `mergeStackMapAndCode`.

The special marker `aftergoto` is used to indicate an unconditional branch.

If we have an unconditional branch at the end of the code, stop.

`mergedCodeIsTypeSafe(Environment, [endOfCode(Offset)], afterGoto).`

After an unconditional branch, if we have a stack map giving the type state for the following instructions, we can proceed and typecheck them using the type state provided by the stack map.

mergedCodeIsTypeSafe(Environment, [stackMap(Offset, MapFrame) | MoreCode], afterGoto):-

— mergedCodeIsTypeSafe(Environment, MoreCode, MapFrame).

If we have a stack map and an incoming type state, the type state must be assignable to the one in the stack map. We may then proceed to type check the rest of the stream with the type state given in the stack map.

mergedCodeIsTypeSafe(Environment, [stackMap(Offset, MapFrame) | MoreCode], frame(Locals, OperandStack, Flags)) :-

— frameIsAssignable(frame(Locals, OperandStack, Flags), MapFrame).

— mergedCodeIsTypeSafe(Environment, MoreCode, MapFrame).

It is illegal to have code after an unconditional branch without a stack frame map being provided for it.

mergedCodeIsTypeSafe(Environment, [instruction(,) | MoreCode], afterGoto) :-

— write_In('No stack frame after unconditional branch').

— fail.

A merged code stream is type safe relative to an incoming type state T, if it begins with an instruction I that is type safe relative to T, I satisfies its exception handlers, and the tail of the stream is type safe given the type state following that execution of I.

mergedCodeIsTypeSafe(Environment, [instruction(Offset, Parse) | MoreCode], frame(Locals, OperandStack, Flags)) :-

Verify the instruction. `NextStackFrame` indicates what falls through to the following instruction. `ExceptionStackFrame` indicates what is passed to exception handlers.

— `instructionIsTypeSafe(Parse, Environment, Offset, frame(Locals, OperandStack, Flags), NextStackFrame, ExceptionStackFrame).`

— `instructionSatisfiesHandlers(Environment, Offset, ExceptionStackFrame).`

— `mergedCodeIsTypeSafe(Environment, MoreCode, NextStackFrame).`

Branching to a target is type safe if the target has an associated stack frame, `Frame`, and the current stack frame, `StackFrame`, is assignable to `Frame`.

`targetIsTypeSafe(Environment, StackFrame, Target) :-`

 `offsetStackFrame(Environment, Target, Frame).`

 `framesAssignable(StackFrame, Frame).`

4.11.1.3.4 Combining Stack Maps and Instruction Streams

The merge of a stream of stack frames and a stream of instructions is defined in this section.

Merging an empty `StackMap` and a list of instructions yields the original list of instructions.

`mergeStackMapAndCode([], CodeList, CodeList).`

Given a list of stack frame maps beginning with the type state for the instruction at `Offset`, and a list of instructions beginning at `Offset`, the merged list consists of the head of the stack frame list, followed by the head of the instruction list, followed by the merge of the tails of the two lists.

mergeStackMapAndCode([stackMap(Offset, Map) | RestMap],
 [instruction(Offset, Parse) | RestCode],
 [stackMap(Offset, Map),
 instruction(Offset, Parse) | RestMerge]) :-
 mergeStackMapAndCode(RestMap, RestCode, RestMerge).

Otherwise, given a list of stack frames beginning with the type state for
 the instruction at OffsetM, and a list of instructions beginning at
 OffsetP, then, if OffsetP < OffsetM then the merged list consists of the
 head of the instruction list, followed by the merge of the stack frame list
 and the tail of the instruction list.

mergeStackMapAndCode([stackMap(OffsetM, Map) | RestMap],
 [instruction(OffsetP, Parse) | RestCode],
 [instruction(OffsetP, Parse) | RestMerge]) :-
 OffsetP < OffsetM,
 mergeStackMapAndCode([stackMap(OffsetM, Map) | RestMap],
 RestCode, RestMerge).

Otherwise, the merge of the two lists is undefined.

Since the instruction list has monotonically increasing offsets, the merge
 of the two lists is not defined unless:

- Every stackmap offset has a corresponding instruction offset.
 - The stackmaps are in monotonically increasing order.
-

4.11.1.3.5 Exception Handling

An instruction *satisfies its exception handlers* if it satisfies every
 exception handler that is applicable to the instruction.

instructionSatisfiesHandlers(Environment, Offset, ExceptionStackFrame)
 :-

```

__  exceptionHandlers(Environment, Handlers),
__  sublist(isApplicableHandler(Offset), Handlers, ApplicableHandlers),
__  checklist(instructionSatisfiesHandler(Environment,
ExceptionStackFrame), ApplicableHandlers).

```

An exception handler is *applicable* to an instruction if the offset of the instruction is greater or equal to the start of the handler's range and less than the end of the handler's range.

isApplicableHandler(Offset, handler(Start, End, Target, ClassName)) :-

```

__  Offset >= Start,
__  Offset < End.

```

An instruction *satisfies* an exception handler if its incoming type state is **StackFrame**, and the handler's target (the initial instruction of the handler code) is type safe assuming an incoming type state T. The type state T is derived from **StackFrame** by replacing the operand stack with a stack whose sole element is the handler's exception class.

instructionSatisfiesHandler(Environment, StackFrame, Handler) :-

```

    Handler = handler( , , Target, ),
    currentClassLoader(Environment, CurrentLoader),
__  handlerExceptionClass(Handler, ExceptionClass, CurrentLoader),
__  /* The stack consists of just the exception. */
__  StackFrame = frame(Locals, , Flags),
__  targetIsTypeSafe(Environment, frame(Locals, [ ExceptionClass ],
Flags), Target).

```

The exception class of a handler is **Throwable** if the handlers class entry is 0, otherwise it is the class named in the handler.

```

handlerExceptionClass(handler( , , , 0),
class('java/lang/Throwable', BL), ) :-

```

isBootstrapLoader(BL).

handlerExceptionClass(handler(, , , Name), class(Name, L), L) :-

Name \= 0.

handlersAreLegal(Environment) :-

exceptionHandlers(Environment, Handlers),

___ checklist(handlerIsLegal(Environment), Handlers).

An exception handler is legal if its start (Start) is less than its end (End), there exists an instruction whose offset is equal to Start, there exists an instruction whose offset equals End and the handler's exception class is assignable to the class Throwable.

handlerIsLegal(Environment, Handler) :-

Handler = handler(Start, End, Target, _).

___ Start < End.

___ allInstructions(Environment, Instructions).

___ member(instruction(Start, _), Instructions).

___ offsetStackFrame(Environment, Target, _).

___ instructionsIncludeEnd(Instructions, End).

currentClassLoader(Environment, CurrentLoader).

___ handlerExceptionClass(Handler, ExceptionClass, CurrentLoader).

isBootstrapLoader(BL).

___ isAssignable(ExceptionClass, class('java/lang/Throwable', BL)).

instructionsIncludeEnd(Instructions, End) :-

member(instruction(End, _), Instructions).

instructionsIncludeEnd(Instructions, End) :-

member(endOfCode(End), Instructions).

4.11.1.4 Instructions

4.11.1.4.1 Isomorphic Instructions

Many bytecodes have type rules that are completely isomorphic to the rules for other bytecodes. If a bytecode $b1$ is isomorphic to another bytecode $b2$, then the type rule for $b1$ is the same as the type rule for $b2$.

instructionIsTypeSafe(Instruction, Environment, Offset, StackFrame,
NextStackFrame, ExceptionStackFrame) :-

instructionHasEquivalentTypeRule(Instruction, IsomorphicInstruction),
instructionIsTypeSafe(IsomorphicInstruction, Environment, Offset,
StackFrame, NextStackFrame, ExceptionStackFrame).

4.11.1.4.2 Manipulating the Operand Stack

This section defines the rules for legally manipulating the type state's operand stack. Manipulation of the operand stack is complicated by the fact that some types occupy two entries on the stack. The predicates given in this section take this into account, allowing the rest of the specification to abstract from this issue.

canPop(frame(Locals, OperandStack, Flags), Types,
frame(Locals, PoppedOperandStack, Flags)) :-
popMatchingList(OperandStack, Types, PoppedOperandStack).

popMatchingList(OperandStack, [], OperandStack).
popMatchingList(OperandStack, [P | Rest], NewOperandStack) :-
popMatchingType(OperandStack, P, TempOperandStack,
ActualType),
popMatchingList(TempOperandStack, Rest, NewOperandStack).

Pop an individual type off the stack. More precisely, if the logical top of the stack is some subtype of the specified type, `Type`, then pop it. If a type occupies two stack slots, the logical top of stack type is really the type just below the top, and the top of stack is the unusable type `top`.

`popMatchingType([ActualType | OperandStack], Type, OperandStack, ActualType) :-`

— `sizeOf(Type, 1).`
 — `isAssignable(ActualType, Type).`

`popMatchingType([top, ActualType | OperandStack], Type, OperandStack, ActualType) :-`

— `sizeOf(Type, 2).`
 — `isAssignable(ActualType, Type).`

Push a logical type onto the stack. The exact behavior varies with the size of the type. If the pushed type is of size 1, we just push it onto the stack. If the pushed type is of size 2, we push it, and then push `top`.

`pushOperandStack(OperandStack, 'void', OperandStack).`

`pushOperandStack(OperandStack, Type, [Type | OperandStack]) :-`

— `sizeOf(Type, 1).`

`pushOperandStack(OperandStack, Type, [top, Type | OperandStack]) :-`

— `sizeOf(Type, 2).`

The length of the operand stack must not exceed the declared maximum stack length.

`operandStackHasLegalLength(Environment, OperandStack) :-`

— `length(OperandStack, Length).`
 — `maxOperandStackLength(Environment, MaxStack).`
 — `Length <= MaxStack.`

Category 1 types occupy a single stack slot. Popping a logical type of category 1, *Type*, off the stack is possible if the top of the stack is *Type* and *Type* is not *top* (otherwise it could denote the upper half of a category 2 type). The result is the incoming stack, with the top slot popped off.

popCategory1([*Type* | *Rest*], *Type*, *Rest*) :-

— *Type* \= *top*.
 — sizeOf(*Type*, 1).

Category 2 types occupy two stack slots. Popping a logical type of category 2, *Type*, off the stack is possible if the top of the stack is type *top*, and the slot directly below it is *Type*. The result is the incoming stack, with the top 2 slots popped off.

popCategory2([*top*, *Type* | *Rest*], *Type*, *Rest*) :-

— sizeOf(*Type*, 2).

canSafelyPush(Environment, InputOperandStack, *Type*,
OutputOperandStack) :-

— pushOperandStack(InputOperandStack, *Type*,
OutputOperandStack).
 — operandStackHasLegalLength(Environment,
OutputOperandStack).

canSafelyPushList(Environment, InputOperandStack, *Types*,
OutputOperandStack) :-

— canPushList(InputOperandStack, *Types*, OutputOperandStack).
 — operandStackHasLegalLength(Environment,
OutputOperandStack).

canPushList(InputOperandStack, [*Type* | *Rest*], OutputOperandStack) :-

— pushOperandStack(InputOperandStack, *Type*,
InterimOperandStack).

— canPushList(InterimOperandStack, Rest, OutputOperandStack).
canPushList(InputOperandStack, [], InputOperandStack).

4.11.1.4.3 Loads

All load instructions are variations on a common pattern, varying the type of the value that the instruction loads.

Loading a value of type *Type* from local variable *Index* is type safe, if the type of that local variable is *ActualType*, *ActualType* is assignable to *Type*, and pushing *ActualType* onto the incoming operand stack is a valid type transition that yields a new type state *NextStackFrame*. After execution of the load instruction, the type state will be *NextStackFrame*.

loadIsTypeSafe(Environment, Index, Type, StackFrame, NextStackFrame)

:-

— StackFrame = frame(Locals, _OperandStack, _Flags).
 — nth0(Index, Locals, ActualType).
 — isAssignable(ActualType, Type).
 — validTypeTransition(Environment, [], ActualType, StackFrame, NextStackFrame).

4.11.1.4.4 Stores

All store instructions are variations on a common pattern, varying the type of the value that the instruction stores.

In general, a store instruction is type safe if the local variable it references is of a type that is a supertype of *Type*, and the top of the operand stack is of a subtype of *Type*, where *Type* is the type the instruction is designed to store.

More precisely, the store is type safe if one can pop a type *ActualType* that “matches” *Type* (i.e., is a subtype of *Type*) off the operand stack, and then legally assign that type the local variable *LIndex*.

```
storeIsTypeSafe( Environment, Index, Type, frame(Locals,
OperandStack, Flags),
frame(NextLocals, NextOperandStack, Flags)) :-
popMatchingType(OperandStack, Type, NextOperandStack,
ActualType),
modifyLocalVariable(Index, ActualType, Locals, NextLocals).
```

Given local variables *Locals*, modifying *LIndex* to have type *Type* results in the local variable list *NewLocals*. The modifications are somewhat involved, because some values (and their corresponding types) occupy two local variables. Hence, modifying *LN* may require modifying *LN+1* (because the type will occupy both the *N* and *N+1* slots) or *LN-1* (because local *N* used to be the upper half of the two word value/type starting at local *N-1*, and so local *N-1* must be invalidated), or both. This is described further below. We start at *L0* and count up.

```
modifyLocalVariable(Index, Type, Locals, NewLocals) :-
modifyLocalVariable(0, Index, Type, Locals, NewLocals).
```

Given the suffix of the local variable list starting at index *I*, *LocalsRest*, modifying local variable *Index* to have type *Type* results in the local variable list suffix *NewLocalsRest*.

If $I < \text{Index}-1$, just copy the input to the output and recurse forward. If $I \equiv \text{Index}-1$, the type of local *I* may change. This can occur if *I* has a type of size 2. Once we set *I+1* to the new type (and the corresponding value), the type/value of *I* will be invalidated, as its upper half will be trashed. Then we recurse forward.

When we find the variable, and it only occupies one word, we change it to *Type* and we’re done.

When we find the variable, and it occupies two words, we change its type to *Type* and the next word to *top*.

```

modifyLocalVariable(I, Index, Type, [Locals1 | LocalsRest],
    [Locals1 | NextLocalsRest] ) :-
__ I < Index - 1,
__ I1 is I + 1,
__ modifyLocalVariable(I1, Index, Type, LocalsRest, NextLocalsRest).

```

```

modifyLocalVariable(I, Index, Type, [Locals1 | LocalsRest],
    [NextLocals1 | NextLocalsRest] ) :-
__ I := Index - 1,
__ modifyPreIndexVariable(Locals1, NextLocals1),
__ modifyLocalVariable(Index, Index, Type, LocalsRest,
NextLocalsRest).

```

```

modifyLocalVariable(Index, Index, Type, [ _ | LocalsRest],
    [Type | LocalsRest] ) :-
__ sizeOf(Type, 1).

```

```

modifyLocalVariable(Index, Index, Type, [ _ , _ | LocalsRest],
    [Type, top | LocalsRest]) :-
__ sizeOf(Type, 2).

```

We refer to a local whose index immediately precedes a local whose type will be modified as a *pre-index variable*. The future type of a pre-index variable of type InputType is Result. If the type, Value, of the pre-index local is of size 1, it doesn't change. If the type of the pre-index local, Value, is 2, we need to mark the lower half of its two word value as unusable, by setting its type to **top**.

```

modifyPreIndexVariable(Type, Type) :- sizeOf(Type, 1).
modifyPreIndexVariable(Type, top) :- sizeOf(Type, 2).

```

Given a list of types, this clause produces a list where every type of size 2 has been substituted by two entries: one for itself, and one **top** entry.

The result then corresponds to the representation of the list as 32 bit words in the Java virtual machine.

expandTypeList([], []).

expandTypeList([Item | List], [Item | Result]) :-

— sizeOf(Item, 1).

— expandTypeList(List, Result).

expandTypeList([Item | List], [Item, top | Result]) :-

— sizeOf(Item, 2).

— expandTypeList(List, Result).

4.11.1.4.5 List of all Instructions

In general, the type rule for an instruction is given relative to an environment Environment that defines the class and method in which the instruction occurs, and the offset Offset within the method at which the instruction occurs. The rule states that if the incoming type state StackFrame fulfills certain requirements, then

- The instruction is type safe.
- It is provable that the type state after the instruction completes normally has a particular form given by NextStackFrame, and that the type state after the instruction completes abruptly is given by ExceptionStackFrame.

The natural language description of the rule is intended to be readable, intuitive and concise. As such, the description avoids repeating all the contextual assumptions given above. In particular:

- We do not explicitly mention the environment.
- When we speak of the operand stack or local variables in the following, we are referring to the operand stack and local variable components of a type state: either the incoming type state or the outgoing one.
- The type state after the instruction completes abruptly is almost always identical to the incoming type state. We only discuss the type state after the instruction completes abruptly when that is not the case.
- We speak of popping and pushing types onto the operand stack. We do not explicitly discuss issues of stack underflow or overflow, but assume that these operations can be

completed successfully. The formal rules for operand stack manipulation ensure that the necessary checks are made.

- Similarly, the text discusses only the manipulation of logical types. In practice, some types take more than one word. We abstract from these representation details in our discussion, but the logical rules that manipulate data do not.

Any ambiguities can be resolved by referring to the formal Prolog rules.

aaload:

An `aaload` instruction is type safe iff one can validly replace types matching `int` and an array type with component type `ComponentType` where `ComponentType` is a subtype of `Object`, with `ComponentType` yielding the outgoing type state.

```

instructionIsTypeSafe(aaload, Environment, _Offset, StackFrame,
                     _NextStackFrame, ExceptionStackFrame) :-
    nth1OperandStackIs(2, StackFrame, ArrayType),
    arrayComponentType(ArrayType, ComponentType),
    isBootstrapLoader(BL),
    validTypeTransition(Environment,
                        [int, arrayOf(class('java/lang/Object', BL))],
                        ComponentType,
                        StackFrame, NextStackFrame),
    exceptionStackFrame(StackFrame, ExceptionStackFrame).

```

The component type of an array of `X` is `X`.

```
arrayComponentType(arrayOf(X), X).
```

We define the component type of `null` to be `null`.

```
arrayComponentType(null, null).
```

aastore:

An `aastore` instruction is type safe iff one can validly pop types matching `Object`, `int`, and an array of `Object` off the incoming operand stack yielding the outgoing type state.

`instructionIsTypeSafe(aastore, Environment, Offset, StackFrame, NextStackFrame, ExceptionStackFrame) :-`
`isBootstrapLoader(BL),`
`canPop(StackFrame, [class('java/lang/Object', BL), int,`
`arrayOf(class('java/lang/Object', BL))], NextStackFrame),`
`exceptionStackFrame(StackFrame, ExceptionStackFrame).`

`aconst null:`

An `aconst null` instruction is type safe if one can validly push the type `null` onto the incoming operand stack yielding the outgoing type state.

`instructionIsTypeSafe(aconst null, Environment, Offset, StackFrame, NextStackFrame, ExceptionStackFrame) :-`
`validTypeTransition(Environment, [], null, StackFrame,`
`NextStackFrame),`
`exceptionStackFrame(StackFrame, ExceptionStackFrame).`

`aload:`

An `aload` instruction with operand `Index` is type safe and yields an outgoing type state `NextStackFrame`, if a load instruction with operand `Index` and type `reference` is type safe and yields an outgoing type state `NextStackFrame`.

`instructionIsTypeSafe(aload(Index), Environment, Offset, StackFrame, NextStackFrame, ExceptionStackFrame) :-`
`loadIsTypeSafe(Environment, Index, reference, StackFrame,`
`NextStackFrame),`
`exceptionStackFrame(StackFrame, ExceptionStackFrame).`

aload <n>:

The instructions `aload <n>`, for $0 \leq n \leq 3$, are typesafe iff the equivalent `aload` instruction is type safe.

`instructionHasEquivalentTypeRule(aload 0, aload(0)).`

`instructionHasEquivalentTypeRule(aload 1, aload(1)).`

`instructionHasEquivalentTypeRule(aload 2, aload(2)).`

`instructionHasEquivalentTypeRule(aload 3, aload(3)).`

anewarray:

An `anewarray` instruction with operand `CP` is type safe iff `CP` refers to a constant pool entry denoting either a class type or an array type, and one can legally replace a type matching `int` on the incoming operand stack with an array with component type `CP` yielding the outgoing type state.

`instructionIsTypeSafe(anewarray(CP), Environment, _Offset, StackFrame, NextStackFrame, ExceptionStackFrame) :-`

___ `(CP = class(,) ; CP = arrayOf()).`

___ `validTypeTransition(Environment, [int], arrayOf(CP),`
_____ `StackFrame, NextStackFrame).`

___ `exceptionStackFrame(StackFrame, ExceptionStackFrame).`

areturn:

An `areturn` instruction is type safe iff the enclosing method has a declared return type, `ReturnType`, that is a reference type, and one can validly pop a type matching `ReturnType` off the incoming operand stack.

`instructionIsTypeSafe(areturn, Environment, _Offset, StackFrame,`
_____ `afterGoto, ExceptionStackFrame) :-`

___ `thisMethodReturnType(Environment, ReturnType).`

___ `isAssignable(ReturnType, reference).`

```

—   canPop(StackFrame, [ReturnType], _ PoppedStackFrame),
—   exceptionStackFrame(StackFrame, ExceptionStackFrame).

```

arraylength:

An `arraylength` instruction is type safe iff one can validly replace an array type on the incoming operand stack with the type `int` yielding the outgoing type state.

```

instructionIsTypeSafe(arraylength, Environment, _ Offset, StackFrame,
                      _ NextStackFrame, ExceptionStackFrame) :-
    nth1OperandStackIs(1, StackFrame, ArrayType),
    —   arrayComponentType(ArrayType, _), % ensure that it is an Array
    —   validTypeTransition(Environment, [top], int, StackFrame,
                          NextStackFrame),
    —   _ exceptionStackFrame(StackFrame,
                          ExceptionStackFrame).

```

astore:

An `astore` instruction with operand `Index` is type safe and yields an outgoing type state `NextStackFrame`, if a store instruction with operand `Index` and type `reference` is type safe and yields an outgoing type state `NextStackFrame`.

```

instructionIsTypeSafe(astore(Index), Environment, _ Offset, StackFrame,
                      _ NextStackFrame, ExceptionStackFrame) :-
    storeIsTypeSafe(Environment, Index, reference, StackFrame,
                    NextStackFrame),
    —   exceptionStackFrame(StackFrame, ExceptionStackFrame).

```

astore <n>:

The instructions `astore_<n>`, for $0 \leq n \leq 3$, are typesafe iff the equivalent `astore` instruction is type safe.

instructionHasEquivalentTypeRule(astore 0, astore(0)).
instructionHasEquivalentTypeRule(astore 1, astore(1)).
instructionHasEquivalentTypeRule(astore 2, astore(2)).
instructionHasEquivalentTypeRule(astore 3, astore(3)).

athrow:

An **athrow** instruction is type safe iff the top of the operand stack matches **Throwable**.

instructionIsTypeSafe(athrow, Environment, Offset, StackFrame,
afterGoto, ExceptionStackFrame) :-
isBootstrapLoader(BL),
canPop(StackFrame, [class('java/lang/Throwable', BL)],
PoppedStackFrame),
exceptionStackFrame(StackFrame, ExceptionStackFrame).

baload:

A **baload** instruction is type safe iff one can validly replace types matching **int** and a small array type on the incoming operand stack with **int** yielding the outgoing type state.

instructionIsTypeSafe(baload, Environment, Offset, StackFrame,
NextStackFrame, ExceptionStackFrame) :-
nth1OperandStackIs(2, StackFrame, Array),
isSmallArray(Array),
validTypeTransition(Environment, [int, top], int,
StackFrame, NextStackFrame),
exceptionStackFrame(StackFrame, ExceptionStackFrame).

An array type is a *small array type* if it is an array of **byte**, an array of **boolean**, or a subtype thereof (**null**).

[isSmallArray\(arrayOf\(byte\)\).](#)
[isSmallArray\(arrayOf\(boolean\)\).](#)
[isSmallArray\(null\).](#)

bastore:

[A **bastore** instruction is type safe iff one can validly pop types matching **int**, **int** and a small array type off the incoming operand stack yielding the outgoing type state.](#)

[instructionIsTypeSafe\(bastore, Environment, Offset, StackFrame, _____ NextStackFrame, ExceptionStackFrame\) :-](#)
[____ nth1OperandStackIs\(3, StackFrame, Array\).](#)
[____ isSmallArray\(Array\).](#)
[____ canPop\(StackFrame, \[int, int, top\], NextStackFrame\).](#)
[exceptionStackFrame\(StackFrame, ExceptionStackFrame\).](#)

bipush:

[A **bipush** instruction is type safe iff the equivalent **sipush** instruction is type safe](#)

[instructionHasEquivalentTypeRule\(bipush\(Value\), sipush\(Value\)\).](#)

caload:

[A **caload** instruction is type safe iff one can validly replace types matching **int** and array of **char** on the incoming operand stack with **int** yielding the outgoing type state.](#)

[instructionIsTypeSafe\(caload, Environment, Offset, StackFrame, _____ NextStackFrame, ExceptionStackFrame\) :-](#)
[____ validTypeTransition\(Environment, \[int, arrayOf\(char\)\], int,](#)

StackFrame, NextStackFrame).
 — exceptionStackFrame(StackFrame, ExceptionStackFrame).

castore:

A **castore** instruction is type safe iff one can validly pop types matching **int**, **int** and array of **char** off the incoming operand stack yielding the outgoing type state.

instructionIsTypeSafe(castore, Environment, Offset, StackFrame,
NextStackFrame, ExceptionStackFrame) :-
 — canPop(StackFrame, [int, int, arrayOf(char)], NextStackFrame).
 — exceptionStackFrame(StackFrame, ExceptionStackFrame).

checkcast:

A **checkcast** instruction with operand **CP** is type safe iff **CP** refers to a constant pool entry denoting either a class or an array, and one can validly replace the type **Object** on top of the incoming operand stack with the type denoted by **CP** yielding the outgoing type state.

instructionIsTypeSafe(checkcast(CP), Environment, Offset, StackFrame,
NextStackFrame, ExceptionStackFrame) :-
 — (CP = class(,) ; CP = arrayOf()),
isBootstrapLoader(BL),
 — validTypeTransition(Environment, [class('java/lang/Object', BL)], CP,
StackFrame, NextStackFrame).
 — exceptionStackFrame(StackFrame, ExceptionStackFrame).

d2f:

A **d2f** instruction is type safe if one can validly pop **double** off the incoming operand stack and replace it with **float**, yielding the outgoing type state.

instructionIsTypeSafe(d2f, Environment, Offset, StackFrame,
NextStackFrame, ExceptionStackFrame) :-


```

__  validTypeTransition(Environment, [double], float,
      StackFrame, NextStackFrame),
__  exceptionStackFrame(StackFrame, ExceptionStackFrame).

```

d2i:

A d2i instruction is type safe if one can validly pop double off the incoming operand stack and replace it with int, yielding the outgoing type state.

```

instructionIsTypeSafe(d2i, Environment, Offset, StackFrame,
      NextStackFrame, ExceptionStackFrame) :-
__  validTypeTransition(Environment, [double], int,
      StackFrame, NextStackFrame),
__  exceptionStackFrame(StackFrame, ExceptionStackFrame).

```

d2l:

A d2l instruction is type safe if one can validly pop double off the incoming operand stack and replace it with long, yielding the outgoing type state.

```

instructionIsTypeSafe(d2l, Environment, Offset, StackFrame,
      NextStackFrame, ExceptionStackFrame) :-
__  validTypeTransition(Environment, [double], long,
      StackFrame, NextStackFrame),
__  exceptionStackFrame(StackFrame, ExceptionStackFrame).

```

dadd:

A dadd instruction is type safe iff one can validly replace types matching double and double on the incoming operand stack with double yielding the outgoing type state.

instructionIsTypeSafe(dadd, Environment, Offset, StackFrame,
NextStackFrame, ExceptionStackFrame) :-
validTypeTransition(Environment, [double, double], double,
StackFrame, NextStackFrame),
exceptionStackFrame(StackFrame, ExceptionStackFrame).

daload:

A **daload** instruction is type safe iff one can validly replace types
matching **int** and array of **double** on the incoming operand stack
with **double** yielding the outgoing type state.

instructionIsTypeSafe(daload, Environment, Offset, StackFrame,
NextStackFrame, ExceptionStackFrame) :-
validTypeTransition(Environment, [int, arrayOf(double)], double,
StackFrame, NextStackFrame),
exceptionStackFrame(StackFrame, ExceptionStackFrame).

dastore:

A **dastore** instruction is type safe iff one can validly pop types
matching **double**, **int** and array of **double** off the incoming
operand stack yielding the outgoing type state.

instructionIsTypeSafe(dastore, Environment, Offset, StackFrame,
NextStackFrame, ExceptionStackFrame) :-
canPop(StackFrame, [double, int, arrayOf(double)],
NextStackFrame),
exceptionStackFrame(StackFrame, ExceptionStackFrame).

dcmp<op>:

A **dcmpg** instruction is type safe iff one can validly replace types
matching **double** and **double** on the incoming operand stack with
int yielding the outgoing type state.

instructionIsTypeSafe(dcmpeg, Environment, Offset, StackFrame,
NextStackFrame, ExceptionStackFrame) :-
validTypeTransition(Environment, [double, double], int,
StackFrame, NextStackFrame),
exceptionStackFrame(StackFrame, ExceptionStackFrame).

A dcmpl instruction is type safe iff the equivalent dcmpeg instruction
is type safe.

instructionHasEquivalentTypeRule(dcmpl, dcmpeg).

dconst <d>:

A dconst_0 instruction is type safe if one can validly push the type
double onto the incoming operand stack yielding the outgoing type
state.

instructionIsTypeSafe(dconst_0, Environment, Offset, StackFrame,
NextStackFrame, ExceptionStackFrame) :-
validTypeTransition(Environment, [], double, StackFrame,
NextStackFrame),
exceptionStackFrame(StackFrame, ExceptionStackFrame).

A dconst_1 instruction is type safe iff the equivalent dconst_0
instruction is type safe.

instructionHasEquivalentTypeRule(dconst_1, dconst_0).

ddiv:

A ddiv instruction is type safe iff the equivalent dadd instruction is
type safe.

instructionHasEquivalentTypeRule(ddiv, dadd).

dload:

A dload instruction with operand Index is type safe and yields an outgoing type state NextStackFrame, if a load instruction with operand Index and type double is type safe and yields an outgoing type state NextStackFrame.

instructionIsTypeSafe(dload(Index), Environment, Offset, StackFrame, NextStackFrame, ExceptionStackFrame) :-
 loadIsTypeSafe(Environment, Index, double, StackFrame, NextStackFrame),
 exceptionStackFrame(StackFrame, ExceptionStackFrame).

dload <n>:

The instructions dload <n>, for $0 \leq n \leq 3$, are typesafe iff the equivalent dload instruction is type safe.

instructionHasEquivalentTypeRule(dload 0, dload(0)).
instructionHasEquivalentTypeRule(dload 1, dload(1)).
instructionHasEquivalentTypeRule(dload 2, dload(2)).
instructionHasEquivalentTypeRule(dload 3, dload(3)).

dmul:

A dmul instruction is type safe iff the equivalent dadd instruction is type safe.

instructionHasEquivalentTypeRule(dmul, dadd).

dneg:

A **dneg** instruction is type safe iff there is a type matching **double** on the incoming operand stack. The **dneg** instruction does not alter the type state.

```
instructionIsTypeSafe(dneg, Environment, _Offset, StackFrame,  
NextStackFrame, ExceptionStackFrame) :-  
validTypeTransition(Environment, [double], double, StackFrame,  
NextStackFrame),  
exceptionStackFrame(StackFrame, ExceptionStackFrame).
```

drem:

A **drem** instruction is type safe iff the equivalent **dadd** instruction is type safe.

instructionHasEquivalentTypeRule(drem, dadd).

```

dreturn:

```

A **dreturn** instruction is type safe if the enclosing method has a declared return type of **double**, and one can validly pop a type matching **double** off the incoming operand stack.

```

instructionIsTypeSafe(dreturn, Environment, Offset, StackFrame,
_____ afterGoto, ExceptionStackFrame) :-
____ thisMethodReturnType(Environment, double),
____ canPop(StackFrame, [double], PoppedStackFrame),
____ exceptionStackFrame(StackFrame, ExceptionStackFrame).

```

dstore:

A `dstore` instruction with operand `Index` is type safe and yields an outgoing type state `NextStackFrame`, if a store instruction with

operand Index and type double is type safe and yields an outgoing type state NextStackFrame.

instructionIsTypeSafe(dstore(Index), Environment, Offset, StackFrame, NextStackFrame, ExceptionStackFrame) :-
— storeIsTypeSafe(Environment, Index, double, StackFrame, NextStackFrame),
— exceptionStackFrame(StackFrame, ExceptionStackFrame).

dstore <n>:

The instructions dstore_<n>, for $0 \leq n \leq 3$, are type safe iff the equivalent dstore instruction is type safe.

instructionHasEquivalentTypeRule(dstore_0, dstore(0)).
instructionHasEquivalentTypeRule(dstore_1, dstore(1)).
instructionHasEquivalentTypeRule(dstore_2, dstore(2)).
instructionHasEquivalentTypeRule(dstore_3, dstore(3)).

dsub:

A dsub instruction is type safe iff the equivalent dadd instruction is type safe.

instructionHasEquivalentTypeRule(dsub, dadd).

dup:

A dup instruction is type safe iff one can validly replace a category 1 type, Type, with the types Type, Type, yielding the outgoing type state.

instructionIsTypeSafe(dup, Environment, Offset, StackFrame, NextStackFrame, ExceptionStackFrame) :-
StackFrame = frame(Locals, InputOperandStack, Flags),

```

    popCategory1(InputOperandStack, Type, _).
  __  canSafelyPush(Environment, InputOperandStack, Type,
    OutputOperandStack).
  __  NextStackFrame = frame(Locals, OutputOperandStack, Flags).
  __  exceptionStackFrame(StackFrame, ExceptionStackFrame).

```

dup x1:

A **dup_x1** instruction is type safe iff one can validly replace two category 1 types, **Type1**, and **Type2**, on the incoming operand stack with the types **Type1**, **Type2**, **Type1**, yielding the outgoing type state.

```

instructionIsTypeSafe(dup_x1, Environment, _Offset, StackFrame,
    NextStackFrame, ExceptionStackFrame) :-
  __  StackFrame = frame(Locals, InputOperandStack, Flags).
  __  popCategory1(InputOperandStack, Type1, Stack1).
  __  popCategory1(Stack1, Type2, Rest).
  __  canSafelyPushList(Environment, Rest, [Type1, Type2, Type1],
    OutputOperandStack).
  __  NextStackFrame = frame(Locals, OutputOperandStack, Flags).
  __  exceptionStackFrame(StackFrame, ExceptionStackFrame).

```

dup x2:

A **dup_x2** instruction is type safe iff it is a type safe form of the **dup_x2** instruction.

```

instructionIsTypeSafe(dup_x2, Environment, _Offset, StackFrame,
    NextStackFrame, ExceptionStackFrame) :-
  __  StackFrame = frame(Locals, InputOperandStack, Flags).
  __  dup_x2SomeFormsIsTypeSafe(Environment, InputOperandStack,
    OutputOperandStack).
  __  NextStackFrame = frame(Locals, OutputOperandStack, Flags).
  __  exceptionStackFrame(StackFrame, ExceptionStackFrame).

```

A `dup_x2` instruction is a type safe form of the `dup_x2` instruction iff it is a type safe form 1 `dup_x2` instruction or a type safe form 2 `dup_x2` instruction.

`dup_x2SomeForm1IsTypeSafe(Environment, InputOperandStack, OutputOperandStack) :-`

— `dup_x2Form1IsTypeSafe(Environment, InputOperandStack, OutputOperandStack).`

`dup_x2SomeForm1IsTypeSafe(Environment, InputOperandStack, OutputOperandStack) :-`

— `dup_x2Form2IsTypeSafe(Environment, InputOperandStack, OutputOperandStack).`

A `dup_x2` instruction is a type safe form 1 `dup_x2` instruction iff one can validly replace three category 1 types, `Type1`, `Type2`, `Type3` on the incoming operand stack with the types `Type1`, `Type2`, `Type3`, `Type1`, yielding the outgoing type state.

`dup_x2Form1IsTypeSafe(Environment, InputOperandStack, OutputOperandStack) :-`

— `popCategory1(InputOperandStack, Type1, Stack1).`

— `popCategory1(Stack1, Type2, Stack2).`

— `popCategory1(Stack2, Type3, Rest).`

— `canSafelyPushList(Environment, Rest, [Type1, Type3, Type2, Type1], OutputOperandStack).`

A `dup_x2` instruction is a type safe form 2 `dup_x2` instruction iff one can validly replace a category 1 type, `Type1`, and a category 2 type, `Type2`, on the incoming operand stack with the types `Type1`, `Type2`, `Type1`, yielding the outgoing type state.

`dup_x2Form2IsTypeSafe(Environment, InputOperandStack, OutputOperandStack) :-`

— `popCategory1(InputOperandStack, Type1, Stack1).`

— [popCategory2\(Stack1, Type2, Rest\).](#)
 — [canSafelyPushList\(Environment, Rest, \[Type1, Type2, Type1\],
 OutputOperandStack\).](#)

dup2:

[A **dup2** instruction is type safe iff it is a type safe form of the **dup2** instruction.](#)

[instructionIsTypeSafe\(dup2, Environment, Offset, StackFrame,
 NextStackFrame, ExceptionStackFrame\) :-](#)
 — [StackFrame = frame\(Locals, InputOperandStack, Flags\).](#)
 — [dup2SomeFormIsTypeSafe\(Environment, InputOperandStack,
 OutputOperandStack\).](#)
 — [NextStackFrame = frame\(Locals, OutputOperandStack, Flags\).](#)
 — [exceptionStackFrame\(StackFrame, ExceptionStackFrame\).](#)

[A **dup2** instruction is a type safe form of the **dup2** instruction iff it is a type safe form 1 **dup2** instruction or a type safe form 2 **dup2** instruction.](#)

[dup2SomeFormIsTypeSafe\(Environment, InputOperandStack,
 OutputOperandStack\) :-](#)
 — [dup2Form1IsTypeSafe\(Environment, InputOperandStack,
 OutputOperandStack\).](#)
[dup2SomeFormIsTypeSafe\(Environment, InputOperandStack,
 OutputOperandStack\) :-](#)
 — [dup2Form2IsTypeSafe\(Environment, InputOperandStack,
 OutputOperandStack\).](#)

[A **dup2** instruction is a type safe form 1 **dup2** instruction iff one can validly replace two category 1 types, **Type1** and **Type2** on the incoming operand stack with the types **Type1, Type2, Type1, Type2**, yielding the outgoing type state.](#)

dup2Form1IsTypeSafe(Environment, InputOperandStack, OutputOperandStack):-

```

__  popCategory1(InputOperandStack, Type1, TempStack),
    popCategory1(TempStack, Type2, _),
__  canSafelyPushList(Environment, InputOperandStack,
                      [Type1, Type2], OutputOperandStack).
```

A **dup2** instruction is a type safe form 2 **dup2** instruction iff one can validly replace a category 2 type, **Type** on the incoming operand stack with the types **Type, Type**, yielding the outgoing type state.

dup2Form2IsTypeSafe(Environment, InputOperandStack, OutputOperandStack):-

```

__  popCategory2(InputOperandStack, Type, _),
__  canSafelyPush(Environment, InputOperandStack, Type,
                  OutputOperandStack).
```

dup2 x1:

A **dup2_x1** instruction is type safe iff it is a type safe form of the **dup2_x1** instruction.

instructionIsTypeSafe(dup2_x1, Environment, Offset, StackFrame, NextStackFrame, ExceptionStackFrame) :-

```

__  StackFrame = frame(Locals, InputOperandStack, Flags),
__  dup2_x1SomeFormIsTypeSafe(Environment, InputOperandStack,
                              OutputOperandStack),
__  NextStackFrame = frame(Locals, OutputOperandStack, Flags),
__  exceptionStackFrame(StackFrame, ExceptionStackFrame).
```

A **dup2_x1** instruction is a type safe form of the **dup2_x1** instruction iff it is a type safe form 1 **dup2_x1** instruction or a type safe form 2 **dup_x2** instruction.

dup2_x1SomeFormIsTypeSafe(Environment, InputOperandStack, OutputOperandStack) :-

— dup2_x1Form1IsTypeSafe(Environment, InputOperandStack, OutputOperandStack).

dup2_x1SomeFormIsTypeSafe(Environment, InputOperandStack, OutputOperandStack) :-

— dup2_x1Form2IsTypeSafe(Environment, InputOperandStack, OutputOperandStack).

A dup2_x1 instruction is a type safe form 1 dup2_x1 instruction iff one can validly replace three category 1 types, Type1, Type2, Type3, on the incoming operand stack with the types Type1, Type2, Type3, Type1, Type2, yielding the outgoing type state.

dup2_x1Form1IsTypeSafe(Environment, InputOperandStack, OutputOperandStack) :-

— popCategory1(InputOperandStack, Type1, Stack1),

— popCategory1(Stack1, Type2, Stack2),

— popCategory1(Stack2, Type3, Rest),

— canSafelyPushList(Environment, Rest, [Type2, Type1, Type3, Type2, Type1], OutputOperandStack).

A dup2_x1 instruction is a type safe form 2 dup2_x1 instruction iff one can validly replace a category 2 type, Type1, and a category 1 type, Type2, on the incoming operand stack with the types Type1, Type2, Type1, yielding the outgoing type state.

dup2_x1Form2IsTypeSafe(Environment, InputOperandStack, OutputOperandStack) :-

— popCategory2(InputOperandStack, Type1, Stack1),

— popCategory1(Stack1, Type2, Rest),

— canSafelyPushList(Environment, Rest, [Type1, Type2, Type1], OutputOperandStack).

dup2 x2:

A dup2_x2 instruction is type safe iff it is a type safe form of the dup2_x2 instruction.

instructionIsTypeSafe(dup2_x2, Environment, Offset, StackFrame, NextStackFrame, ExceptionStackFrame) :-

— StackFrame = frame(Locals, InputOperandStack, Flags),
— dup2_x2SomeFormIsTypeSafe(Environment, InputOperandStack, OutputOperandStack),
— NextStackFrame = frame(Locals, OutputOperandStack, Flags),
— exceptionStackFrame(StackFrame, ExceptionStackFrame).

A dup2_x2 instruction is a type safe form of the dup2_x2 instruction iff one of the following holds:

- it is a type safe form 1 dup2_x2 instruction.
- it is a type safe form 2 dup2_x2 instruction.
- it is a type safe form 3 dup2_x2 instruction.
- it is a type safe form 4 dup2_x2 instruction.

dup2_x2SomeFormIsTypeSafe(Environment, InputOperandStack, OutputOperandStack) :-

— dup2_x2Form1IsTypeSafe(Environment, InputOperandStack, OutputOperandStack).

dup2_x2SomeFormIsTypeSafe(Environment, InputOperandStack, OutputOperandStack) :-

— dup2_x2Form2IsTypeSafe(Environment, InputOperandStack, OutputOperandStack).

dup2_x2SomeFormIsTypeSafe(Environment, InputOperandStack, OutputOperandStack) :-

— dup2_x2Form3IsTypeSafe(Environment, InputOperandStack, OutputOperandStack).

dup2_x2SomeFormIsTypeSafe(Environment, InputOperandStack, OutputOperandStack) :-

category 2 type, Type3, on the incoming operand stack with the types Type1, Type2, Type3, Type1, Type2, yielding the outgoing type state.

dup2_x2Form3IsTypeSafe(Environment, InputOperandStack, OutputOperandStack) :-

```

—   popCategory1(InputOperandStack, Type1, Stack1),
—   popCategory1(Stack1, Type2, Stack2),
—   popCategory2(Stack2, Type3, Rest),
—   canSafelyPushList(Environment, Rest, [Type2, Type1, Type3,
Type2, Type1], OutputOperandStack).

```

A dup2_x2 instruction is a type safe form 4 dup2_x2 instruction iff one can validly replace two category 2 types, Type1, Type2, on the incoming operand stack with the types Type1, Type2, Type1, yielding the outgoing type state.

dup2_x2Form4IsTypeSafe(Environment, InputOperandStack, OutputOperandStack) :-

```

—   popCategory2(InputOperandStack, Type1, Stack1),
—   popCategory2(Stack1, Type2, Rest),
—   canSafelyPushList(Environment, Rest, [Type1, Type2, Type1],
OutputOperandStack).

```

f2d:

An f2d instruction is type safe if one can validly pop float off the incoming operand stack and replace it with double, yielding the outgoing type state.

instructionIsTypeSafe(f2d, Environment, Offset, StackFrame,
NextStackFrame, ExceptionStackFrame) :-

```

—   validTypeTransition(Environment, [float], double, StackFrame,
NextStackFrame),
—   exceptionStackFrame(StackFrame, ExceptionStackFrame).

```

f2i:

An f2i instruction is type safe if one can validly pop float off the incoming operand stack and replace it with int, yielding the outgoing type state.

instructionIsTypeSafe(f2i, Environment, Offset, StackFrame,
NextStackFrame, ExceptionStackFrame) :-

— validTypeTransition(Environment, [float], int, StackFrame,
NextStackFrame),

— exceptionStackFrame(StackFrame, ExceptionStackFrame).

f2l:

An f2l instruction is type safe if one can validly pop float off the incoming operand stack and replace it with long, yielding the outgoing type state.

instructionIsTypeSafe(f2l, Environment, Offset, StackFrame,
NextStackFrame, ExceptionStackFrame) :-

— validTypeTransition(Environment, [float], long, StackFrame,
NextStackFrame),

— exceptionStackFrame(StackFrame, ExceptionStackFrame).

fadd:

An fadd instruction is type safe iff one can validly replace types matching float and float on the incoming operand stack with float yielding the outgoing type state.

instructionIsTypeSafe(fadd, Environment, Offset, StackFrame,
NextStackFrame, ExceptionStackFrame) :-

— validTypeTransition(Environment, [float, float], float, StackFrame,
NextStackFrame),

— exceptionStackFrame(StackFrame, ExceptionStackFrame).

faload:

An `faload` instruction is type safe iff one can validly replace types matching `int` and array of `float` on the incoming operand stack with `float` yielding the outgoing type state.

`instructionIsTypeSafe(faload, Environment, Offset, StackFrame, NextStackFrame, ExceptionStackFrame) :-`

— `validTypeTransition(Environment, [int, arrayOf(float)], float, StackFrame, NextStackFrame),`

— `exceptionStackFrame(StackFrame, ExceptionStackFrame).`

fastore:

An `fastore` instruction is type safe iff one can validly pop types matching `float`, `int` and array of `float` off the incoming operand stack yielding the outgoing type state.

`instructionIsTypeSafe(fastore, Environment, Offset, StackFrame, NextStackFrame, ExceptionStackFrame) :-`

— `canPop(StackFrame, [float, int, arrayOf(float)], NextStackFrame),`

— `exceptionStackFrame(StackFrame, ExceptionStackFrame).`

fcmp<op>:

An `fcmpg` instruction is type safe iff one can validly replace types matching `float` and `float` on the incoming operand stack with `int` yielding the outgoing type state.

`instructionIsTypeSafe(fcmpg, Environment, Offset, StackFrame, NextStackFrame, ExceptionStackFrame) :-`

— `validTypeTransition(Environment, [float, float], int, StackFrame, NextStackFrame),`

— `exceptionStackFrame(StackFrame, ExceptionStackFrame).`

An `fcmpl` instruction is type safe iff the equivalent `fcmpg` instruction is type safe.

instructionHasEquivalentTypeRule(fcimpl, fcmpg).

fconst <f>:

An fconst_0 instruction is type safe if one can validly push the type float onto the incoming operand stack yielding the outgoing type state.

```

instructionIsTypeSafe(fconst 0, Environment, _Offset, StackFrame,
                    NextStackFrame, ExceptionStackFrame) :-
    __ validTypeTransition(Environment, [], float, StackFrame,
NextStackFrame),
    __ exceptionStackFrame(StackFrame, ExceptionStackFrame).

```

The rules for the other variants of `fconst` are equivalent:

instructionHasEquivalentTypeRule(fconst 1, fconst 0).
instructionHasEquivalentTypeRule(fconst 2, fconst 0).

fdiv:

An `fdiv` instruction is type safe iff the equivalent `fadd` instruction is type safe.

instructionHasEquivalentTypeRule(fdiv, fadd).

fload:

An `fload` instruction with operand `Index` is type safe and yields an outgoing type state `NextStackFrame`, if a load instruction with operand `Index` and type `float` is type safe and yields an outgoing type state `NextStackFrame`.

instructionIsTypeSafe(float(Index), Environment, _Offset, StackFrame,
 _____ NextStackFrame, ExceptionStackFrame) :-
_____ loadIsTypeSafe(Environment, Index, float, StackFrame,
 NextStackFrame),
exceptionStackFrame(StackFrame, ExceptionStackFrame).

fload <n>:

The instructions **fload <n>**, for $0 \leq n \leq 3$, are typesafe iff the
 equivalent **fload** instruction is type safe.

instructionHasEquivalentTypeRule(fload 0, fload(0)).
instructionHasEquivalentTypeRule(fload 1, fload(1)).
instructionHasEquivalentTypeRule(fload 2, fload(2)).
instructionHasEquivalentTypeRule(fload 3, fload(3)).

fmul:

An **fmul** instruction is type safe iff the equivalent **fadd** instruction is
 type safe.

instructionHasEquivalentTypeRule(fmul, fadd).

fneg:

An **fneg** instruction is type safe iff there is a type matching **float** on
 the incoming operand stack. The **fneg** instruction does not alter the
 type state.

instructionIsTypeSafe(fneg, Environment, _Offset, StackFrame,
 _____ NextStackFrame, ExceptionStackFrame) :-
_____ validTypeTransition(Environment, [float], float, StackFrame,
 NextStackFrame),
_____ exceptionStackFrame(StackFrame, ExceptionStackFrame).

frem:

An frem instruction is type safe iff the equivalent fadd instruction is type safe.

instructionHasEquivalentTypeRule(frem, fadd).

freturn:

An freturn instruction is type safe if the enclosing method has a declared return type of float, and one can validly pop a type matching float off the incoming operand stack.

instructionIsTypeSafe(freturn, Environment, Offset, StackFrame,
afterGoto, ExceptionStackFrame) :-
 — thisMethodReturnType(Environment, float),
 — canPop(StackFrame, [float], PoppedStackFrame),
 — exceptionStackFrame(StackFrame, ExceptionStackFrame).

fstore:

An fstore instruction with operand Index is type safe and yields an outgoing type state NextStackFrame, if a store instruction with operand Index and type float is type safe and yields an outgoing type state NextStackFrame.

instructionIsTypeSafe(fstore(Index), Environment, Offset, StackFrame,
NextStackFrame, ExceptionStackFrame) :-
storeIsTypeSafe(Environment, Index, float, StackFrame,
NextStackFrame),
exceptionStackFrame(StackFrame, ExceptionStackFrame).

fstore <n>:

The protected check applies only to members of superclasses of the current class. Other cases will be caught by the access checking done at resolution time. If the name of a class is not the name of any superclass, it cannot be a superclass, and so it can safely be ignored.

passesProtectedCheck(Environment, MemberClassName, MemberName, MemberDescriptor, StackFrame) :-

thisClass(Environment, class(CurrentClassName, CurrentLoader)),
superclassChain(CurrentClassName, CurrentLoader, Chain),
notMember(class(MemberClassName, _), Chain).

Using a superclass member that is not protected is trivially correct.

If the MemberClassName is the same as the name of a superclass, the class being resolved may indeed be a superclass. In this case, if no superclass named MemberClassName in a different runtime package has a protected member named MemberName with descriptor MemberDescriptor, the the protected check need not apply. This is because the actual class being resolved will either be one of these superclasses, in which case we know that it is either in the same runtime package, and the access is legal, or the member in question is not protected and the check does not apply; or it will be a subclass, in which case the check would succeed anyway; or it will be some other class in the same run time package ,in which case the access is legal and the check need not take place; or the verifier need not flag this as a problem, since it will be caught anyway because resolution will per force fail.

passesProtectedCheck(Environment, MemberClassName, MemberName, MemberDescriptor, StackFrame) :-

thisClass(Environment, class(CurrentClassName, CurrentLoader)),
superclassChain(CurrentClassName, CurrentLoader, Chain),
member(class(MemberClassName, _), Chain),
classesInOtherPkgWithProtectedMember
class(CurrentClassName, CurrentLoader), MemberName,
MemberDescriptor, MemberClassName, Chain, []).

If there does exist a protected superclass member in a different run time package, then load MemberClassName; if the member in question is not protected, the check does not apply.

passesProtectedCheck(Environment, MemberClassName, MemberName, MemberDescriptor, frame(Locals, [Target | Rest], _Flags)) :-

```
thisClass(Environment, class(CurrentClassName, CurrentLoader)),
superclassChain(CurrentClassName, CurrentLoader, Chain),
member(class(MemberClassName, _), Chain),
classesInOtherPkgWithProtectedMember(
    class(CurrentClassName, CurrentLoader), MemberName,
    MemberDescriptor, MemberClassName, Chain, List),
List /= [],
loadedClass(MemberClassName, CurrentLoader,
    ReferencedClass),
isNotProtected(ReferencedClass, MemberName,
    MemberDescriptor).
```

Otherwise, use of a member of an object of type Target requires that Target be assignable to the type of the current class.

passesProtectedCheck(Environment, MemberClassName, MemberName, MemberDescriptor, frame(Locals, [Target | Rest], _Flags)) :-

```
thisClass(Environment, class(CurrentClassName, CurrentLoader)),
superclassChain(CurrentClassName, CurrentLoader, Chain),
member(class(MemberClassName, _), Chain),
classesInOtherPkgWithProtectedMember(
    class(CurrentClassName, CurrentLoader),
```

MemberName, MemberDescriptor, MemberClassName,
Chain, List),
List /= [],
loadedClass(MemberClassName, CurrentLoader,
ReferencedClass),
isProtected(ReferencedClass, MemberName, MemberDescriptor),
isAssignable(Target, class(CurrentClassName, CurrentLoader)).

superclassChain(ClassName, L, [class(SuperclassName, Ls) | Rest]) :-
loadedClass(ClassName, L, Class),
classSuperclassName(Class, SuperclassName),
classDefiningLoader(Class, Ls),
superclassChain(SuperclassName, Ls, Rest).

superclassChain('java/lang/Object', L, []) :-
loadedClass('java/lang/Object', L, Class),
classDefiningLoader(Class, BL),
isBootstrapLoader(BL).

The predicate `classesInOtherPkgWithProtectedMember(Class,`
`MemberName, MemberDescriptor, MemberClassName, Chain,`
`List)` is true if `List` is the set of classes in `Chain` with name
`MemberClassName` that are in a different runtime package then
`Class` which have a protected member named `MemberName` with
descriptor `MemberDescriptor`

classesInOtherPkgWithProtectedMember(, , , , [], []).
classesInOtherPkgWithProtectedMember(Class, MemberName,
MemberDescriptor, MemberClassName, [class(MemberClassName, L) |
Tail], [class(MemberClassName, L) | T]) :-
differentRuntimePackage(Class, class(MemberClassName, L)),
loadedClass(MemberClassName, L, Super),

isProtected(Super, MemberName, MemberDescriptor),
classesInOtherPkgWithProtectedMember(Class, MemberName,
MemberDescriptor, MemberClassName, Tail, T).

classesInOtherPkgWithProtectedMember(Class, MemberName,
MemberDescriptor, MemberClassName, [class(MemberClassName, L) |
Tail], T) :-
differentRuntimePackage(Class, class(MemberClassName, L)),
loadedClass(MemberClassName, L, Super),
isNotProtected(Super, MemberName, MemberDescriptor),
classesInOtherPkgWithProtectedMember(Class, MemberName,
MemberDescriptor, MemberClassName, Tail, T).

classesInOtherPkgWithProtectedMember(Class, MemberName,
MemberDescriptor, MemberClassName, [class(MemberClassName, L) |
Tail], T) :-
sameRuntimePackage(Class, class(MemberClassName, L)),
classesInOtherPkgWithProtectedMember(Class,
MemberName, MemberDescriptor, MemberClassName, Tail, T).

getstatic:

A `getstatic` instruction with operand CP is type safe iff CP refers
to a constant pool entry denoting a field whose declared type is
FieldType, and one can validly push FieldType on the incoming
operand stack yielding the outgoing type state.

instructionIsTypeSafe(getstatic(CP), Environment, _Offset, StackFrame,
_____NextStackFrame, ExceptionStackFrame) :-
___ CP = field(_FieldClass, _FieldName, FieldDescriptor),

```

—   parseFieldDescriptor(FieldDescriptor, FieldType).
—   validTypeTransition(Environment, [], FieldType, StackFrame,
NextStackFrame).
—   exceptionStackFrame(StackFrame, ExceptionStackFrame).

```

goto:

A goto instruction is type safe iff its target operand is a valid branch target.

```

instructionIsTypeSafe(goto(Target), Environment, Offset, StackFrame,
_____ afterGoto, ExceptionStackFrame) :-
    targetIsTypeSafe(Environment, StackFrame, Target),
—   exceptionStackFrame(StackFrame, ExceptionStackFrame).

```

i2b:

An i2b instruction is type safe iff the equivalent ineg instruction is type safe.

```

instructionHasEquivalentTypeRule(i2b, ineg).

```

i2c:

An i2c instruction is type safe iff the equivalent ineg instruction is type safe.

```

instructionHasEquivalentTypeRule(i2c, ineg).

```

i2d:

An i2d instruction is type safe if one can validly pop int off the incoming operand stack and replace it with double, yielding the outgoing type state.

instructionIsTypeSafe(i2d, Environment, Offset, StackFrame,
NextStackFrame, ExceptionStackFrame) :-
 — validTypeTransition(Environment, [int], double, StackFrame,
NextStackFrame),
 — exceptionStackFrame(StackFrame, ExceptionStackFrame).

i2f:

An i2f instruction is type safe if one can validly pop int off the
incoming operand stack and replace it with float, yielding the
outgoing type state.

instructionIsTypeSafe(i2f, Environment, Offset, StackFrame,
NextStackFrame, ExceptionStackFrame) :-
 — validTypeTransition(Environment, [int], float, StackFrame,
NextStackFrame),
 — exceptionStackFrame(StackFrame, ExceptionStackFrame).

i2l:

An i2l instruction is type safe if one can validly pop int off the
incoming operand stack and replace it with long, yielding the outgoing
type state.

instructionIsTypeSafe(i2l, Environment, Offset, StackFrame,
NextStackFrame, ExceptionStackFrame) :-
 — validTypeTransition(Environment, [int], long, StackFrame,
NextStackFrame),
 — exceptionStackFrame(StackFrame, ExceptionStackFrame).

i2s:

An i2s instruction is type safe iff the equivalent ineg instruction is
type safe.

instructionHasEquivalentTypeRule(i2s, ineg).

iadd:

An iadd instruction is type safe iff one can validly replace types matching `int` and `int` on the incoming operand stack with `int` yielding the outgoing type state.

instructionIsTypeSafe(iadd, Environment, Offset, StackFrame, NextStackFrame, ExceptionStackFrame) :-

— validTypeTransition(Environment, [int, int], int, StackFrame, NextStackFrame),

— exceptionStackFrame(StackFrame, ExceptionStackFrame).

iaload:

An iaload instruction is type safe iff one can validly replace types matching `int` and array of `int` on the incoming operand stack with `int` yielding the outgoing type state.

instructionIsTypeSafe(iaload, Environment, Offset, StackFrame, NextStackFrame, ExceptionStackFrame) :-

— validTypeTransition(Environment, [int, arrayOf(int)], int, StackFrame, NextStackFrame),

— exceptionStackFrame(StackFrame, ExceptionStackFrame).

iand:

An iand instruction is type safe iff the equivalent iadd instruction is type safe.

instructionHasEquivalentTypeRule(iand, iadd).

iastore:

An iastore instruction is type safe iff one can validly pop types matching `int`, `int` and array of `int` off the incoming operand stack yielding the outgoing type state.

instructionIsTypeSafe(iastore, Environment, Offset, StackFrame,
NextStackFrame, ExceptionStackFrame) :-
 — canPop(StackFrame, [int, int, arrayOf(int)], NextStackFrame),
 — exceptionStackFrame(StackFrame, ExceptionStackFrame).

if_acmp<cond>:

An if_acmpeq instruction is type safe iff one can validly pop types
matching reference and reference on the incoming operand
stack yielding the outgoing type state NextStackFrame, and the
operand of the instruction, Target, is a valid branch target assuming an
incoming type state of NextStackFrame.

instructionIsTypeSafe(if_acmpeq(Target), Environment, Offset,
StackFrame,
NextStackFrame, ExceptionStackFrame) :-
 — canPop(StackFrame, [reference, reference], NextStackFrame),
 — targetIsTypeSafe(Environment, NextStackFrame, Target),
exceptionStackFrame(StackFrame, ExceptionStackFrame).

The rule for if_acmp_ne is identical.

instructionHasEquivalentTypeRule(if_acmpne(Target),
if_acmpeq(Target)).

if_icmp<cond>:

An if_icmpeq instruction is type safe iff one can validly pop types
matching int and int on the incoming operand stack yielding the
outgoing type state NextStackFrame, and the operand of the
instruction, Target, is a valid branch target assuming an incoming type
state of NextStackFrame.

instructionIsTypeSafe(if_icmpeq(Target), Environment, Offset,
StackFrame,
NextStackFrame, ExceptionStackFrame) :-
 — canPop(StackFrame, [int, int], NextStackFrame),

```

__  targetIsTypeSafe(Environment, NextStackFrame, Target),
__  exceptionStackFrame(StackFrame, ExceptionStackFrame).

```

The rules for all other variants of the `if_icmp` instruction are identical

```

instructionHasEquivalentTypeRule(if_icmpge(Target), if_icmpeq(Target)).
instructionHasEquivalentTypeRule(if_icmpgt(Target), if_icmpeq(Target)).
instructionHasEquivalentTypeRule(if_icmple(Target), if_icmpeq(Target)).
instructionHasEquivalentTypeRule(if_icmplt(Target), if_icmpeq(Target)).
instructionHasEquivalentTypeRule(if_icmpne(Target), if_icmpeq(Target)).

```

if <cond>:

An `if_eq` instruction is type safe iff one can validly pop a type matching `int` off the incoming operand stack yielding the outgoing type state `NextStackFrame`, and the operand of the instruction, `Target`, is a valid branch target assuming an incoming type state of `NextStackFrame`.

```

instructionIsTypeSafe(ifeq(Target), Environment, _Offset, StackFrame,
__  NextStackFrame, ExceptionStackFrame) :-
__  canPop(StackFrame, [int], NextStackFrame),
__  targetIsTypeSafe(Environment, NextStackFrame, Target),
__  exceptionStackFrame(StackFrame, ExceptionStackFrame).

```

The rules for all other variations of the `if<cond>` instruction are identical

```

instructionHasEquivalentTypeRule(ifge(Target), ifeq(Target)).
instructionHasEquivalentTypeRule(ifgt(Target), ifeq(Target)).
instructionHasEquivalentTypeRule(ifle(Target), ifeq(Target)).
instructionHasEquivalentTypeRule(iflt(Target), ifeq(Target)).
instructionHasEquivalentTypeRule(ifne(Target), ifeq(Target)).

```

ifnonnull:

An `ifnonnull` instruction is type safe iff one can validly pop a type matching `reference` off the incoming operand stack yielding the outgoing type state `NextStackFrame`, and the operand of the instruction, `Target`, is a valid branch target assuming an incoming type state of `NextStackFrame`.

`instructionIsTypeSafe(ifnonnull(Target), Environment, Offset, StackFrame,`
`NextStackFrame, ExceptionStackFrame) :-`
`canPop(StackFrame, [reference], NextStackFrame),`
`targetIsTypeSafe(Environment, NextStackFrame, Target),`
`exceptionStackFrame(StackFrame, ExceptionStackFrame).`

ifnull:

An `ifnull` instruction is type safe iff the equivalent `ifnonnull` instruction is type safe.

`instructionHasEquivalentTypeRule(ifnull(Target), ifnonnull(Target)).`

iinc:

An `iinc` instruction with first operand `Index` is type safe iff `LIndex` has type `int`. The `iinc` instruction does not change the type state.

`instructionIsTypeSafe(iinc(Index, Value), Environment, Offset, StackFrame,`
`StackFrame, ExceptionStackFrame) :-`
`StackFrame = frame(Locals, OperandStack, Flags),`
`nth0(Index, Locals, int),`
`exceptionStackFrame(StackFrame, ExceptionStackFrame).`

iload:

An `iload` instruction with operand `Index` is type safe and yields an outgoing type state `NextStackFrame`, if a load instruction with operand `Index` and type `int` is type safe and yields an outgoing type state `NextStackFrame`.

`instructionIsTypeSafe(ildoad(Index), Environment, Offset, StackFrame, NextStackFrame, ExceptionStackFrame) :-`

`loadIsTypeSafe(Environment, Index, int, StackFrame, NextStackFrame),`

`exceptionStackFrame(StackFrame, ExceptionStackFrame).`

iload <n>:

The instructions `iload <n>`, for $0 \leq n \leq 3$, are typesafe iff the equivalent `iload` instruction is type safe.

`instructionHasEquivalentTypeRule(ildoad 0, iload(0)).`

`instructionHasEquivalentTypeRule(ildoad 1, iload(1)).`

`instructionHasEquivalentTypeRule(ildoad 2, iload(2)).`

`instructionHasEquivalentTypeRule(ildoad 3, iload(3)).`

imul:

An `imul` instruction is type safe iff the equivalent `iadd` instruction is type safe.

`instructionHasEquivalentTypeRule(imul, iadd).`

ineg:

An `ineg` instruction is type safe iff there is a type matching `int` on the incoming operand stack. The `ineg` instruction does not alter the type state.

`instructionIsTypeSafe(ineg, Environment, Offset, StackFrame,`

NextStackFrame, ExceptionStackFrame) :-
 — validTypeTransition(Environment, [int], int, StackFrame,
NextStackFrame).
 — exceptionStackFrame(StackFrame, ExceptionStackFrame).

instanceof:

An instanceof instruction with operand CP is type safe iff CP
refers to a constant pool entry denoting either a class or an array, and
one can validly replace the type Object on top of the incoming
operand stack with type int yielding the outgoing type state.

instructionIsTypeSafe(instanceof(CP), Environment, _Offset,
StackFrame, NextStackFrame, ExceptionStackFrame) :-
 — (CP = class(,) : CP = arrayOf(),
isBootstrapLoader(BL),
 — validTypeTransition(Environment, [class('java/lang/Object'), BL], int,
StackFrame, NextStackFrame).
 — exceptionStackFrame(StackFrame, ExceptionStackFrame).

invokeinterface:

An invokeinterface instruction is type safe iff all of the following
conditions hold:

- Its first operand, CP, refers to a constant pool entry denoting an
interface method named MethodName with descriptor Descriptor
that is a member of an interface MethodClassName.
- MethodName is not <init>.
- MethodName is not <clinit>.
- Its second operand, Count, is a valid count operand (see below).
- One can validly replace types matching the type
MethodClassName and the argument types given in Descriptor on
the incoming operand stack with the return type given in Descriptor,
yielding the outgoing type state.

instructionIsTypeSafe(invokeinterface(CP, Count, 0), Environment,
_Offset,

```

      StackFrame, NextStackFrame, ExceptionStackFrame) :-
__   CP = imethod(MethodClassName, MethodName, Descriptor),
__   MethodName \= '<init>',
__   MethodName \= '<clinit>',
__   parseMethodDescriptor(Descriptor, OperandArgList, ReturnType),
      currentClassLoader(Environment, L),
__   reverse([class(MethodClassName, L) | OperandArgList],
StackArgList),
__   canPop(StackFrame, StackArgList, TempFrame),
__   validTypeTransition(Environment, [], ReturnType, TempFrame,
NextStackFrame),
__   countsValid(Count, StackFrame, TempFrame),
exceptionStackFrame(StackFrame, ExceptionStackFrame).

```

The count operand of an `invokeinterface` instruction is valid if it equals the size of the arguments to the instruction. This is equal to the difference between the size of `InputFrame` and `OutputFrame`.

```

countsValid(Count, InputFrame, OutputFrame) :-
__   InputFrame = frame( Locals1, OperandStack1, Flags1),
__   OutputFrame = frame( Locals2, OperandStack2, Flags2),
__   length(OperandStack1, Length1),
__   length(OperandStack2, Length2),
__   Count =:= Length1 - Length2.

```

invokespecial:

An `invokespecial` instruction is type safe iff all of the following conditions hold:

- Its first operand, CP, refers to a constant pool entry denoting a method named `MethodName` with descriptor `Descriptor` that is a member of a class `MethodClassName`.

Either

- `MethodName` is not `<init>`.
- `MethodName` is not `<clinit>`.

- One can validly replace types matching the current class and the argument types given in Descriptor on the incoming operand stack with the return type given in Descriptor, yielding the outgoing type state.
- One can validly replace types matching the class `MethodClassName` and the argument types given in Descriptor on the incoming operand stack with the return type given in Descriptor.

```

instructionIsTypeSafe(invokespecial(CP), Environment, _Offset,
StackFrame, NextStackFrame, ExceptionStackFrame) :-
    __ CP = method(MethodClassName, MethodName, Descriptor),
    __ MethodName \= '<init>',
    __ MethodName \= '<clinit>',
    __ parseMethodDescriptor(Descriptor, OperandArgList, ReturnType),
    __ thisClass(Environment, CurrentClass),
    __ reverse([CurrentClass | OperandArgList], StackArgList),
    __ validTypeTransition(Environment, StackArgList, ReturnType,
StackFrame, NextStackFrame),
        currentClassLoader(Environment, L),
        reverse([class(MethodClassName, L) | OperandArgList],
StackArgList2),
    __ validTypeTransition(Environment, StackArgList2, ReturnType,
StackFrame, _ResultStackFrame),
        isAssignable(class(CurrentClassName, L),
class(MethodClassName, L)),
        exceptionStackFrame(StackFrame, ExceptionStackFrame).

```

Or

- MethodName is `<init>`.
- Descriptor specifies a `void` return type.
- One can validly pop types matching the argument types given in Descriptor and an uninitialized type, `UninitializedArg`, off the incoming operand stack, yielding `OperandStack`.

- The outgoing type state is derived from the incoming type state by first replacing the incoming operand stack with `OperandStack` and then replacing all instances of `UninitializedArg` with the type of instance being initialized.

`instructionIsTypeSafe(invokespecial(CP), Environment, Offset, StackFrame, NextStackFrame, ExceptionStackFrame) :-`

```

—   CP = method(MethodClassName, '<init>', Descriptor),
—   parseMethodDescriptor(Descriptor, OperandArgList, void),
—   reverse(OperandArgList, StackArgList),
—   canPop(StackFrame, StackArgList, TempFrame),
—   TempFrame = frame(Locals, FullOperandStack, Flags),
—   FullOperandStack = [UninitializedArg | OperandStack],
—   currentClassLoader(Environment, CurrentLoader),
—   rewrittenUninitializedType(UninitializedArg, Environment,
—   class(MethodClassName, CurrentLoader), This),
—   rewrittenInitializationFlags(UninitializedArg, Flags, NextFlags),
—   substitute(UninitializedArg, This, OperandStack,
—   NextOperandStack),
—   substitute(UninitializedArg, This, Locals, NextLocals),
—   NextStackFrame = frame(NextLocals, NextOperandStack,
—   NextFlags),
—   ExceptionStackFrame = frame(NextLocals, [], Flags),
—   passesProtectedCheck(Environment, MethodClassName, '<init>',
—   Descriptor, StackFrame).

```

Special rule for `invokespecial` of an `<init>` method.

This rule is the sole motivation for passing back a distinct exception stack frame. The concern is that `invokespecial` can cause a superclass `<init>` method to be invoked, and that invocation could fail, leaving `this` uninitialized. This situation cannot be created using Java programming language source code, but can be created through JVM assembly programming.

The original frame holds an uninitialized object in a local and has flag `uninitializedThis`. Normal termination of `invokespecial` initializes the uninitialized object and turns off the `uninitializedThis` flag. But if the invocation of an `<init>` method throws an exception, the uninitialized object might be left in a partially initialized state, and needs to be made permanently unusable. This is represented by an exception frame containing the broken object (the new value of the local) and the `uninitializedThis` flag (the old flag). There is no way to get from an apparently-initialized object bearing the `uninitializedThis` flag to a properly initialized object, so the object is permanently unusable. If not for this case, the exception stack frame could be the same as the input stack frame.

`rewrittenUninitializedType(uninitializedThis, Environment, MethodClass, This) :-`

— `thisClass(Environment, This).`

`rewrittenUninitializedType(uninitialized(Address), Environment, MethodClass, MethodClass) :-`

— `allInstructions(Environment, Instructions).`

— `member(instruction(Address, new(MethodClass)), Instructions).`

Computes what type the uninitialized argument's type needs to be rewritten to.

There are 2 cases.

If we are initializing an object within its constructor, its type is initially `uninitializedThis`. This type will be rewritten to the type of the class of the `<init>` method.

The second case arises from initialization of an object created by `new`. The uninitialized arg type is rewritten to `MethodClass`, the type of the method holder of `<init>`. We check whether there really is a `new` instruction at `Address`.

`rewrittenInitializationFlags(uninitializedThis, Flags, []).`

`rewrittenInitializationFlags(uninitialized(), Flags, Flags).`

substitute(Old, New, [], []).

substitute(Old, New, [Old | FromRest], [New | ToRest]) :- substitute(Old, New, FromRest, ToRest).

substitute(Old, New, [From1 | FromRest], [From1 | ToRest]) :-

From1 \= Old,

substitute(Old, New, FromRest, ToRest).

invokestatic:

An `invokestatic` instruction is type safe iff all of the following conditions hold:

- Its first operand, CP, refers to a constant pool entry denoting a method named `MethodName` with descriptor `Descriptor`.
- `MethodName` is not `<clinit>`.
- One can validly replace types matching the argument types given in `Descriptor` on the incoming operand stack with the return type given in `Descriptor`, yielding the outgoing type state.

instructionIsTypeSafe(invokestatic(CP), Environment, Offset, StackFrame,

NextStackFrame, ExceptionStackFrame) :-

— CP = method(MethodClassName, MethodName, Descriptor),

— MethodName \= '<clinit>',

— parseMethodDescriptor(Descriptor, OperandArgList, ReturnType),

— reverse(OperandArgList, StackArgList).

— validTypeTransition(Environment, StackArgList, ReturnType, StackFrame, NextStackFrame),

— exceptionStackFrame(StackFrame, ExceptionStackFrame).

invokevirtual:

An `invokevirtual` instruction is type safe iff all of the following conditions hold:

- Its first operand, CP, refers to a constant pool entry denoting a method named `MethodName` with descriptor `Descriptor` that is a member of an class `MethodClassName`.
- `MethodName` is not `<init>`.
- `MethodName` is not `<clinit>`.
- One can validly replace types matching the class `MethodClassName` and the argument types given in `Descriptor` on the incoming operand stack with the return type given in `Descriptor`, yielding the outgoing type state.
- If the method is protected, the usage conforms to the special rules governing access to protected members.

instructionIsTypeSafe(invokevirtual(CP), Environment, Offset, StackFrame, NextStackFrame, ExceptionStackFrame) :-

```

—      CP = method(MethodClassName, MethodName, Descriptor),
—      MethodName \= '<init>',
—      MethodName \= '<clinit>',
—      parseMethodDescriptor(Descriptor, OperandArgList,
Return Type),
      reverse(OperandArgList, ArgList),
      currentClassLoader(Environment, L),
—      reverse([class(MethodClassName, L) | OperandArgList],
StackArgList),
—      validTypeTransition(Environment, StackArgList, Return Type,
StackFrame, NextStackFrame),
      canPop(StackFrame, ArgList, PoppedFrame),
      passesProtectedCheck(Environment, MethodClassName,
MethodName, Descriptor, PoppedFrame),
—      exceptionStackFrame(StackFrame, ExceptionStackFrame).
```

ior:

An `ior` instruction is type safe iff the equivalent `iadd` instruction is type safe.

instructionHasEquivalentTypeRule(ior, iadd).

irem:

An *irem* instruction is type safe iff the equivalent *iadd* instruction is type safe.

instructionHasEquivalentTypeRule(irem, iadd).

ireturn:

An *ireturn* instruction is type safe if the enclosing method has a declared return type of *int*, and one can validly pop a type matching *int* off the incoming operand stack.

instructionIsTypeSafe(ireturn, Environment, Offset, StackFrame,
afterGoto, ExceptionStackFrame) :-
— thisMethodReturnType(Environment, int),
— canPop(StackFrame, [int], PoppedStackFrame),
— exceptionStackFrame(StackFrame, ExceptionStackFrame).

ishl:

An *ishl* instruction is type safe iff the equivalent *iadd* instruction is type safe.

instructionHasEquivalentTypeRule(ishl, iadd).

ishr:

An *ishr* instruction is type safe iff the equivalent *iadd* instruction is type safe.

instructionHasEquivalentTypeRule(ishr, iadd).

istore:

An **istore** instruction with operand **Index** is type safe and yields an outgoing type state **NextStackFrame**, if a store instruction with operand **Index** and type **int** is type safe and yields an outgoing type state **NextStackFrame**.

instructionIsTypeSafe(istore(Index), Environment, Offset, StackFrame, NextStackFrame, ExceptionStackFrame) :-
storeIsTypeSafe(Environment, Index, int, StackFrame, NextStackFrame).
— exceptionStackFrame(StackFrame, ExceptionStackFrame).

istore <n>:

The instructions **istore <n>**, for $0 \leq n \leq 3$, are typesafe iff the equivalent **istore** instruction is type safe.

instructionHasEquivalentTypeRule(istore 0, istore(0)).
instructionHasEquivalentTypeRule(istore 1, istore(1)).
instructionHasEquivalentTypeRule(istore 2, istore(2)).
instructionHasEquivalentTypeRule(istore 3, istore(3)).

isub:

An **isub** instruction is type safe iff the equivalent **iadd** instruction is type safe.

instructionHasEquivalentTypeRule(isub, iadd).

iushr:

An **iushr** instruction is type safe iff the equivalent **iadd** instruction is type safe.

instructionHasEquivalentTypeRule(iushr, iadd).

ixor:

An **ixor** instruction is type safe iff the equivalent **iadd** instruction is type safe.

instructionHasEquivalentTypeRule(ixor, iadd).

l2d:

An **l2d** instruction is type safe if one can validly pop **long** off the incoming operand stack and replace it with **double**, yielding the outgoing type state.

instructionIsTypeSafe(l2d, Environment, Offset, StackFrame,
NextStackFrame, ExceptionStackFrame) :-
__ validTypeTransition(Environment, [long], double, StackFrame,
NextStackFrame),
__ exceptionStackFrame(StackFrame, ExceptionStackFrame).

l2f:

An **l2f** instruction is type safe if one can validly pop **long** off the incoming operand stack and replace it with **float**, yielding the outgoing type state.

instructionIsTypeSafe(l2f, Environment, Offset, StackFrame,
NextStackFrame, ExceptionStackFrame) :-
__ validTypeTransition(Environment, [long], float, StackFrame,
NextStackFrame),
__ exceptionStackFrame(StackFrame, ExceptionStackFrame).

l2i:

An **l2i** instruction is type safe if one can validly pop **long** off the incoming operand stack and replace it with **int**, yielding the outgoing type state.

instructionIsTypeSafe(l2i, Environment, Offset, StackFrame,

```

NextStackFrame, ExceptionStackFrame) :-
    validTypeTransition(Environment, [long], int, StackFrame,
NextStackFrame),
    exceptionStackFrame(StackFrame, ExceptionStackFrame).

```

ladd:

An **ladd** instruction is type safe iff one can validly replace types matching **long** and **long** on the incoming operand stack with **long** yielding the outgoing type state.

```

instructionIsTypeSafe(ladd, Environment, Offset, StackFrame,
NextStackFrame, ExceptionStackFrame) :-
    validTypeTransition(Environment, [long, long], long, StackFrame,
NextStackFrame),
    exceptionStackFrame(StackFrame, ExceptionStackFrame).

```

laload:

An **laload** instruction is type safe iff one can validly replace types matching **int** and array of **long** on the incoming operand stack with **long** yielding the outgoing type state.

```

instructionIsTypeSafe(laload, Environment, Offset, StackFrame,
NextStackFrame, ExceptionStackFrame) :-
    validTypeTransition(Environment, [int, arrayOf(long)], long,
StackFrame, NextStackFrame),
    exceptionStackFrame(StackFrame, ExceptionStackFrame).

```

land:

An **land** instruction is type safe iff the equivalent **ladd** instruction is type safe.

```

instructionHasEquivalentTypeRule(land, ladd).

```

lastore:

A `lastore` instruction is type safe iff one can validly pop types matching `long`, `int` and array of `long` off the incoming operand stack yielding the outgoing type state.

`instructionIsTypeSafe(lastore, Environment, Offset, StackFrame, NextStackFrame, ExceptionStackFrame) :-`
— `canPop(StackFrame, [long, int, arrayOf(long)], NextStackFrame).`
— `exceptionStackFrame(StackFrame, ExceptionStackFrame).`

lcmp:

A `lcmp` instruction is type safe iff one can validly replace types matching `long` and `long` on the incoming operand stack with `int` yielding the outgoing type state.

`instructionIsTypeSafe(lcmp, Environment, Offset, StackFrame, NextStackFrame, ExceptionStackFrame) :-`
— `validTypeTransition(Environment, [long, long], int, StackFrame, NextStackFrame).`
— `exceptionStackFrame(StackFrame, ExceptionStackFrame).`

lconst <l>:

An `lconst_0` instruction is type safe if one can validly push the type `long` onto the incoming operand stack yielding the outgoing type state.

`instructionIsTypeSafe(lconst 0, Environment, Offset, StackFrame, NextStackFrame, ExceptionStackFrame) :-`
— `validTypeTransition(Environment, [], long, StackFrame, NextStackFrame).`
— `exceptionStackFrame(StackFrame, ExceptionStackFrame).`

An `lconst_1` instruction is type safe iff the equivalent `lconst_0` instruction is type safe.

instructionHasEquivalentTypeRule(lconst 1,lconst 0).

ldc:

An ldc instruction with operand CP is type safe iff CP refers to a constant pool entry denoting an entity of type Type, where Type is either int, float, String or Class and one can validly push Type onto the incoming operand stack yielding the outgoing type state.

instructionIsTypeSafe(ldc(CP), Environment, Offset, StackFrame, NextStackFrame, ExceptionStackFrame) :-

— functor(CP, Tag, _).
isBootstrapLoader(BL).
— member([Tag, Type], [[int, int], [float, float], [string, class('java/lang/String', BL)], [classConst, class('java/lang/Class', BL)]]).
— validTypeTransition(Environment, [], Type, StackFrame, NextStackFrame).
— exceptionStackFrame(StackFrame, ExceptionStackFrame).

ldc w:

An ldc_w instruction is type safe iff the equivalent ldc instruction is type safe.

instructionHasEquivalentTypeRule(ldc_w(CP), ldc(CP))

ldc2 w:

An ldc2_w instruction with operand CP is type safe iff CP refers to a constant pool entry denoting an entity of type Tag, where Tag is either Long or double, and one can validly push Tag onto the incoming operand stack yielding the outgoing type state.

instructionIsTypeSafe(ldc2_w(CP), Environment, Offset, StackFrame, NextStackFrame, ExceptionStackFrame) :-

— functor(CP, Tag, _).

— [member\(Tag, \[long, double\]\).](#)
 — [validTypeTransition\(Environment, \[\], Tag, StackFrame,
NextStackFrame\).](#)
 — [exceptionStackFrame\(StackFrame, ExceptionStackFrame\).](#)

ldiv:

[An `ldiv` instruction is type safe iff the equivalent `ladd` instruction is type safe.](#)

[instructionHasEquivalentTypeRule\(ldiv, ladd\).](#)

lload:

[An `lload` instruction with operand `Index` is type safe and yields an outgoing type state `NextStackFrame`, if a load instruction with operand `Index` and type `long` is type safe and yields an outgoing type state `NextStackFrame`.](#)

[instructionIsTypeSafe\(lload\(Index\), Environment, Offset, StackFrame,
NextStackFrame, ExceptionStackFrame\) :-](#)

— [loadIsTypeSafe\(Environment, Index, long, StackFrame,
NextStackFrame\).](#)
 — [exceptionStackFrame\(StackFrame, ExceptionStackFrame\).](#)

lload <n>:

[The instructions `lload_<n>`, for \$0 \leq n \leq 3\$, are typesafe iff the equivalent `lload` instruction is type safe.](#)

[instructionHasEquivalentTypeRule\(lload_0, lload\(0\)\).](#)
[instructionHasEquivalentTypeRule\(lload_1, lload\(1\)\).](#)
[instructionHasEquivalentTypeRule\(lload_2, lload\(2\)\).](#)
[instructionHasEquivalentTypeRule\(lload_3, lload\(3\)\).](#)

lmul:

An `lor` instruction is type safe iff the equivalent `ladd` instruction is type safe.

`instructionHasEquivalentTypeRule(lor, ladd).`

lrem:

An `lrem` instruction is type safe iff the equivalent `ladd` instruction is type safe.

`instructionHasEquivalentTypeRule(lrem, ladd).`

lreturn:

An `lreturn` instruction is type safe if the enclosing method has a declared return type of `long`, and one can validly pop a type matching `long` off the incoming operand stack.

`instructionIsTypeSafe(lreturn, Environment, Offset, StackFrame,`
`afterGoto, ExceptionStackFrame) :-`

— `thisMethodReturnType(Environment, long),`
 — `canPop(StackFrame, [long], PoppedStackFrame),`
 — `exceptionStackFrame(StackFrame, ExceptionStackFrame).`

lshl:

An `lshl` instruction is type safe if one can validly replace the types `int` and `long` on the incoming operand stack with the type `long` yielding the outgoing type state.

`instructionIsTypeSafe(lshl, Environment, Offset, StackFrame,`
`NextStackFrame, ExceptionStackFrame) :-`

— `validTypeTransition(Environment, [int, long], long, StackFrame,`
`NextStackFrame),`
 — `exceptionStackFrame(StackFrame, ExceptionStackFrame).`

lshr:

An `lshr` instruction is type safe iff the equivalent `lshl` instruction is type safe.

`instructionHasEquivalentTypeRule(lshr, lshl).`

lstore:

An `lstore` instruction with operand `Index` is type safe and yields an outgoing type state `NextStackFrame`, if a store instruction with operand `Index` and type `long` is type safe and yields an outgoing type state `NextStackFrame`.

`instructionIsTypeSafe(lstore(Index), Environment, _Offset, StackFrame, _NextStackFrame, ExceptionStackFrame) :-`

 `storeIsTypeSafe(Environment, Index, long, StackFrame, NextStackFrame),`

 `exceptionStackFrame(StackFrame, ExceptionStackFrame).`

lstore <n>:

The instructions `lstore_<n>`, for $0 \leq n \leq 3$, are typesafe iff the equivalent `lstore` instruction is type safe.

`instructionHasEquivalentTypeRule(lstore_0, lstore(0)).`

`instructionHasEquivalentTypeRule(lstore_1, lstore(1)).`

`instructionHasEquivalentTypeRule(lstore_2, lstore(2)).`

`instructionHasEquivalentTypeRule(lstore_3, lstore(3)).`

lsub:

An `lsub` instruction is type safe iff the equivalent `ladd` instruction is type safe.

instructionHasEquivalentTypeRule(lsub, ladd).

lxor:

An lxor instruction is type safe iff the equivalent ladd instruction is type safe.

instructionHasEquivalentTypeRule(lxor, ladd).

lushr:

An lushr instruction is type safe iff the equivalent lshl instruction is type safe.

instructionHasEquivalentTypeRule(lushr, lshl).

monitorenter:

A monitorenter instruction is type safe iff one can validly pop a type matching reference off the incoming operand stack yielding the outgoing type state.

instructionIsTypeSafe(monitorenter, Environment, Offset, StackFrame, NextStackFrame, ExceptionStackFrame) :-
canPop(StackFrame, [reference], NextStackFrame),
exceptionStackFrame(StackFrame, ExceptionStackFrame).

monitorexit:

A monitorexit instruction is type safe iff the equivalent monitorenter instruction is type safe.

instructionHasEquivalentTypeRule(monitorexit, monitorenter).

multinewarray:

A `multianewarray` instruction with operands `CP` and `Dim` is type safe iff `CP` refers to a constant pool entry denoting an array type whose dimension is greater or equal to `Dim`, `Dim` is strictly positive, and one can validly replace `Dim` `int` types on the incoming operand stack with the type denoted by `CP` yielding the outgoing type state.

`instructionIsTypeSafe(multianewarray(CP, Dim), Environment, _Offset, StackFrame,`

NextStackFrame, ExceptionStackFrame) :-

```

__ CP = arrayOf( _).
__ classDimension(CP, Dimension).
__ Dimension >= Dim.
__ Dim > 0.
__ /* Make a list of Dim ints */
__ findall(int, between(1, Dim, _), IntList).
__ validTypeTransition(Environment, IntList, CP, StackFrame,
NextStackFrame).
__ exceptionStackFrame(StackFrame, ExceptionStackFrame).

```

The dimension of an array type whose component type is also an array type is 1 more than the dimension of its component type.

`classDimension(arrayOf(X), Dimension) :-`

```

__ classDimension(X, Dimension1).
__ Dimension is Dimension1 + 1.

```

`classDimension(_ , Dimension) :- Dimension = 0.`

new:

A `new` instruction with operand `CP` at offset `Offset` is type safe iff `CP` refers to a constant pool entry denoting a class type, the type `uninitialized(Offset)` does not appear in the incoming operand stack, and one can validly push `uninitialized(Offset)`

onto the incoming operand stack and replace uninitialized(Offset) with top in the incoming local variables yielding the outgoing type state.

```

instructionIsTypeSafe(new(CP), Environment, Offset, StackFrame,
    NextStackFrame, ExceptionStackFrame) :-
    StackFrame = frame(Locals, OperandStack, Flags),
    CP = class( , ),
    NewItem = uninitialized(Offset),
    notMember(NewItem, OperandStack),
    substitute(NewItem, top, Locals, NewLocals),
    validTypeTransition(Environment, [], NewItem,
        frame(NewLocals, OperandStack, Flags),
        NextStackFrame),
    exceptionStackFrame(StackFrame, ExceptionStackFrame).

```

newarray:

A newarray instruction with operand TypeCode is type safe iff TypeCode corresponds to the primitive type ElementType, and one can validly replace the type int on the incoming operand stack with the type 'array of ElementType', yielding the outgoing type state.

```

instructionIsTypeSafe(newarray(TypeCode), Environment, Offset,
    StackFrame,
    NextStackFrame, ExceptionStackFrame) :-
    primitiveArrayInfo(TypeCode, TypeChar, ElementType,
        VerifierType),
    validTypeTransition(Environment, [int], arrayOf(ElementType),
        StackFrame, NextStackFrame),
    exceptionStackFrame(StackFrame, ExceptionStackFrame).

```

The correspondence between type codes and primitive types is specified by the following predicate:

[primitiveArrayInfo\(4, 0'Z, boolean, int\).](#)
[primitiveArrayInfo\(5, 0'C, char, int\).](#)
[primitiveArrayInfo\(6, 0'F, float, float\).](#)
[primitiveArrayInfo\(7, 0'D, double, double\).](#)
[primitiveArrayInfo\(8, 0'B, byte, int\).](#)
[primitiveArrayInfo\(9, 0'S, short, int\).](#)
[primitiveArrayInfo\(10, 0'I, int, int\).](#)
[primitiveArrayInfo\(11, 0'J, long, long\).](#)

nop:

[A **nop** instruction is always type safe. The **nop** instruction does not affect the type state.](#)

[instructionIsTypeSafe\(nop, Environment, Offset, StackFrame,](#)
[StackFrame, ExceptionStackFrame\) :-](#)
[exceptionStackFrame\(StackFrame, ExceptionStackFrame\).](#)

pop:

[A **pop** instruction is type safe iff one can validly pop a category 1 type off the incoming operand stack yielding the outgoing type state.](#)

[instructionIsTypeSafe\(pop, Environment, Offset, StackFrame,](#)
[NextStackFrame, ExceptionStackFrame\) :-](#)
[StackFrame = frame\(Locals, \[Type | Rest\], Flags\).](#)
[Type \= top.](#)
[sizeof\(Type, 1\).](#)
[NextStackFrame = frame\(Locals, Rest, Flags\).](#)
[exceptionStackFrame\(StackFrame, ExceptionStackFrame\).](#)

pop2:

A **pop2** instruction is type safe iff it is a type safe form of the **pop2** instruction.

instructionIsTypeSafe(pop2, Environment, Offset, StackFrame,
NextStackFrame, ExceptionStackFrame) :-

- StackFrame = frame(Locals, InputOperandStack, Flags).
- pop2SomeFormIsTypeSafe(InputOperandStack,
OutputOperandStack).
- NextStackFrame = frame(Locals, OutputOperandStack, Flags).
- exceptionStackFrame(StackFrame, ExceptionStackFrame).

A **pop2** instruction is a type safe form of the **pop2** instruction iff it is a type safe form 1 **pop2** instruction or a type safe form 2 **pop2** instruction.

pop2SomeFormIsTypeSafe(InputOperandStack, OutputOperandStack) :-

- pop2Form1IsTypeSafe(InputOperandStack, OutputOperandStack).

pop2SomeFormIsTypeSafe(InputOperandStack, OutputOperandStack) :-

- pop2Form2IsTypeSafe(InputOperandStack, OutputOperandStack).

A **pop2** instruction is a type safe form 1 **pop2** instruction iff one can validly pop two types of size 1 off the incoming operand stack yielding the outgoing type state.

pop2Form1IsTypeSafe([Type1, Type2 | Rest], Rest) :-

- sizeOf(Type1, 1).
- sizeOf(Type2, 1).

A **pop2** instruction is a type safe form 2 **pop2** instruction iff one can validly pop a type of size 2 off the incoming operand stack yielding the outgoing type state.

pop2Form2IsTypeSafe([top, Type | Rest], Rest) :-
 — sizeOf(Type, 2).

putfield:

A **putfield** instruction with operand CP is type safe iff CP refers to a constant pool entry denoting a field whose declared type is **FieldType**, declared in a class **FieldClass**, and one can validly pop types matching **FieldType** and **FieldClass** off the incoming operand stack yielding the outgoing type state.

instructionIsTypeSafe(putfield(CP), Environment, Offset, StackFrame,
 NextStackFrame, ExceptionStackFrame) :-
 — CP = field(FieldClass, FieldName, FieldDescriptor),
 — parseFieldDescriptor(FieldDescriptor, FieldType),
canPop(StackFrame, [FieldType], PoppedFrame),
passesProtectedCheck(Environment, FieldClass, FieldName,
FieldDescriptor, PoppedFrame),
currentClassLoader(Environment, CurrentLoader),
canPop(StackFrame, [FieldType, class(FieldClass, CurrentLoader)],
NextStackFrame),
 — exceptionStackFrame(StackFrame, ExceptionStackFrame).

putstatic:

A **putstatic** instruction with operand CP is type safe iff CP refers to a constant pool entry denoting a field whose declared type is **FieldType**, and one can validly pop a type matching **FieldType** off the incoming operand stack yielding the outgoing type state.

```

instructionIsTypeSafe(putstatic(CP), Environment, Offset, StackFrame,
NextStackFrame, ExceptionStackFrame) :-
__ CP = field( FieldClass, FieldName, FieldDescriptor),
__ parseFieldDescriptor(FieldDescriptor, FieldType),
__ canPop(StackFrame, [FieldType], NextStackFrame),
__ exceptionStackFrame(StackFrame, ExceptionStackFrame).

```

return:

A **return** instruction is type safe if the enclosing method declares a **void** return type, and either:

- The enclosing method is not an **<init>** method, **or**
- **this** has already been completely initialized at the point where the instruction occurs.

```

instructionIsTypeSafe(return, Environment, Offset, StackFrame,
afterGoto, ExceptionStackFrame) :-
__ thisMethodReturnType(Environment, void),
__ StackFrame = frame( Locals, OperandStack, Flags),
__ notMember(flagThisUninit, Flags),
__ exceptionStackFrame(StackFrame, ExceptionStackFrame).

```

saload:

An **saload** instruction is type safe iff one can validly replace types matching **int** and array of **short** on the incoming operand stack with **int** yielding the outgoing type state.

```

instructionIsTypeSafe(saload, Environment, Offset, StackFrame,
NextStackFrame, ExceptionStackFrame) :-
__ validTypeTransition(Environment, [int, arrayOf(short)], int,
StackFrame, NextStackFrame),
__ exceptionStackFrame(StackFrame, ExceptionStackFrame).

```

sastore:

An **sastore** instruction is type safe iff one can validly pop types matching **int**, **int** and array of **short** off the incoming operand stack yielding the outgoing type state.

```

instructionIsTypeSafe(sastore, Environment, Offset, StackFrame,
                     NextStackFrame, ExceptionStackFrame) :-
    canPop(StackFrame, [int, int, arrayOf(short)], NextStackFrame),
    exceptionStackFrame(StackFrame, ExceptionStackFrame).

```

sipush:

An **sipush** instruction is type safe iff one can validly push the type **int** onto the incoming operand stack yielding the outgoing type state.

```

instructionIsTypeSafe(sipush( Value), Environment, Offset, StackFrame,
                     NextStackFrame, ExceptionStackFrame) :-
    validTypeTransition(Environment, [], int, StackFrame,
                       NextStackFrame),
    exceptionStackFrame(StackFrame, ExceptionStackFrame).

```

swap:

A **swap** instruction is type safe iff one can validly replace two category 1 types, **Type1** and **Type2**, on the incoming operand stack with the types **Type2** and **Type1** yielding the outgoing type state.

```

instructionIsTypeSafe(swap, Environment, Offset, StackFrame,
                     NextStackFrame, ExceptionStackFrame) :-
    StackFrame = frame( Locals, [Type1, Type2 | Rest], Flags),
    sizeOf(Type1, 1),
    sizeOf(Type2, 1),
    NextStackFrame = frame( Locals, [Type2, Type1 | Rest], Flags).

```

— exceptionStackFrame(StackFrame, ExceptionStackFrame).

tableswitch:

A **tableswitch** instruction is type safe if its keys are sorted, one can validly pop **int** off the incoming operand stack yielding a new type state **BranchStackFrame**, and all of the instructions targets are valid branch targets assuming **BranchStackFrame** as their incoming type state.

instructionIsTypeSafe(tableswitch(Targets, Keys), Environment, _Offset, _StackFrame, afterGoto, ExceptionStackFrame) :-

— sort(Keys, Keys).

— canPop(StackFrame, [int], BranchStackFrame).

— checklist(targetIsTypeSafe(Environment, BranchStackFrame), Targets).

— exceptionStackFrame(StackFrame, ExceptionStackFrame).

wide:

The **wide** instructions follow the same rules as the instructions they widen.

instructionHasEquivalentTypeRule(wide(WidenedInstruction), WidenedInstruction).

The type state after an instruction completes abruptly is the same as the incoming type state, except that the operand stack is
empty.exceptionStackFrame(StackFrame, ExceptionStackFrame) :-

— StackFrame = frame(Locals, _OperandStack, Flags).

— ExceptionStackFrame = frame(Locals, [], Flags).

Most of the type rules in this specification depend on the notion of a valid type transition

A type transition is valid if one can pop a list of expected types off the incoming type state's operand stack and replace them with an expected result type, resulting in a new valid type state. In particular, the size of the operand stack in the new type state must not exceed its maximum declared size.

```
validTypeTransition(Environment, ExpectedTypesOnStack, ResultType,  
_____ frame(Locals, InputOperandStack, Flags),  
_____ frame(Locals, NextOperandStack, Flags)) :-  
___ popMatchingList(InputOperandStack, ExpectedTypesOnStack,  
InterimOperandStack),  
___ pushOperandStack(InterimOperandStack, ResultType,  
NextOperandStack),  
___ operandStackHasLegalLength(Environment, NextOperandStack).
```

Access lth element of the operand stack from a type state.

```
nth1OperandStackIs(l, frame( Locals, OperandStack, Flags), Element)  
:-  
___ nth1(l, OperandStack, Element).
```

4.11.2 Verification by Type Inference

Class files that do not contain StackMapTable attributes (which necessarily have a version number of 49.0 or below) must be verified using type inference, as described below.

4.11.2.1 The Process of Verification by Type Inference

During linking, the verifier checks the code array of the Code attribute for each method of the class file by performing data-flow analysis on each method. The verifier ensures that at any given point in the program, no matter what code path is taken to reach that point, the following is true:

- The operand stack is always the same size and contains the same types of values.
- No local variable is accessed unless it is known to contain a value of an appropriate type.
- Methods are invoked with the appropriate arguments.
- Fields are assigned only using values of appropriate types.
- All opcodes have appropriate type arguments on the operand stack and in the local variable array.
- There is never an uninitialized class instance in a local variable in code protected by an exception handler. However, an uninitialized class instance may be on the operand stack in code protected by an exception handler. When an exception is thrown, the contents of the operand stack are discarded.

For further information on this pass, see Section 4.11.2.2, “The Bytecode Verifier.”

For efficiency reasons, certain tests that could in principle be performed by the verifier are delayed until the first time the code for the method is actually invoked. In so doing, the verifier avoids loading *class* files unless it has to.

For example, if a method invokes another method that returns an instance of class A, and that instance is assigned only to a field of the same type, the verifier does not bother to check if the class A actually exists. However, if it is assigned to a field of the type B, the definitions of both A and B must be loaded in to ensure that A is a subclass of B.

4.11.2.2 The Bytecode Verifier

This section looks at the verification of Java virtual machine code in more detail.

The code for each method is verified independently. First, the bytes that make up the code are broken up into a sequence of instructions, and the index into the code array of the start of each instruction is placed in an array. The verifier then goes through the code a second time and parses the instructions. During this pass a data structure is built to hold information about each Java virtual machine instruction in the method. The operands, if any, of each instruction are checked to make sure they are valid. For instance:

- Branches must be within the bounds of the code array for the method.
- The targets of all control-flow instructions are each the start of an instruction. In the case of a *wide* instruction, the *wide* opcode is considered the start of the instruction, and the opcode giving the operation modified by that *wide* instruction is not considered to start an instruction. Branches into the middle of an instruction are disallowed.
- No instruction can access or modify a local variable at an index greater than or equal to the number of local variables that its method indicates it allocates.
- All references to the constant pool must be to an entry of the appropriate type. For example: the instruction *getfield* must reference a field.
- The code does not end in the middle of an instruction.
- Execution cannot fall off the end of the code.
- For each exception handler, the starting and ending point of code protected by the handler must be at the beginning of an instruction or, in the case of the ending point, immediately past the end of the code. The starting point must be before the ending point. The exception handler code must start at a valid instruction, and it must not start at an opcode being modified by the *wide* instruction.

For each instruction of the method, the verifier records the contents of the operand stack and the contents of the local variable array prior to the execution of that instruction. For the operand stack, it needs to know the stack height and the type of each value on it. For each local variable, it needs to know either the type of the contents of that local variable or that the local variable contains an unusable or unknown value (it might be uninitialized). The bytecode verifier does not need to distinguish between the integral types (e.g., *byte*, *short*, *char*) when determining the value types on the operand stack.

Next, a data-flow analyzer is initialized. For the first instruction of the method, the local variables that represent parameters initially contain values of the types indicated by the method's type descriptor; the operand stack is empty. All other local variables contain an illegal value. For the other instructions, which have not been examined yet, no information is available regarding the operand stack or local variables.

Finally, the data-flow analyzer is run. For each instruction, a "changed" bit indicates whether this instruction needs to be looked at. Initially, the "changed"

bit is set only for the first instruction. The data-flow analyzer executes the following loop:

1. Select a virtual machine instruction whose “changed” bit is set. If no instruction remains whose “changed” bit is set, the method has successfully been verified. Otherwise, turn off the “changed” bit of the selected instruction.
2. Model the effect of the instruction on the operand stack and local variable array by doing the following:
 - If the instruction uses values from the operand stack, ensure that there are a sufficient number of values on the stack and that the top values on the stack are of an appropriate type. Otherwise, verification fails.
 - If the instruction uses a local variable, ensure that the specified local variable contains a value of the appropriate type. Otherwise, verification fails.
 - If the instruction pushes values onto the operand stack, ensure that there is sufficient room on the operand stack for the new values. Add the indicated types to the top of the modeled operand stack.
 - If the instruction modifies a local variable, record that the local variable now contains the new type.
3. Determine the instructions that can follow the current instruction. Successor instructions can be one of the following:
 - The next instruction, if the current instruction is not an unconditional control transfer instruction (for instance *goto*, *return*, or *athrow*). Verification fails if it is possible to “fall off” the last instruction of the method.
 - The target(s) of a conditional or unconditional branch or switch.
 - Any exception handlers for this instruction.
4. Merge the state of the operand stack and local variable array at the end of the execution of the current instruction into each of the successor instructions. In the special case of control transfer to an exception handler, the operand stack is set to contain a single object of the exception type indicated by the exception handler information.

- If this is the first time the successor instruction has been visited, record that the operand stack and local variable values calculated in steps 2 and 3 are the state of the operand stack and local variable array prior to executing the successor instruction. Set the “changed” bit for the successor instruction.
- If the successor instruction has been seen before, merge the operand stack and local variable values calculated in steps 2 and 3 into the values already there. Set the “changed” bit if there is any modification to the values.

5. Continue at step 1.

To merge two operand stacks, the number of values on each stack must be identical. The types of values on the stacks must also be identical, except that differently typed reference values may appear at corresponding places on the two stacks. In this case, the merged operand stack contains a reference to an instance of the first common superclass of the two types. Such a reference type always exists because the type `Object` is a superclass of all class and interface types. If the operand stacks cannot be merged, verification of the method fails.

To merge two local variable array states, corresponding pairs of local variables are compared. If the two types are not identical, then unless both contain reference values, the verifier records that the local variable contains an unusable value. If both of the pair of local variables contain reference values, the merged state contains a reference to an instance of the first common superclass of the two types.

If the data-flow analyzer runs on a method without reporting a verification failure, then the method has been successfully verified by the `class` file verifier.

Certain instructions and data types complicate the data-flow analyzer. We now examine each of these in more detail.

4.11.2.3 Values of Types `long` and `double`

Values of the `long` and `double` types are treated specially by the verification process.

Whenever a value of type `long` or `double` is moved into a local variable at index n , index $n + 1$ is specially marked to indicate that it has been reserved by the value at index n and must not be used as a local variable index. Any value previously at index $n + 1$ becomes unusable.

Whenever a value is moved to a local variable at index n , the index $n - 1$ is examined to see if it is the index of a value of type `long` or `double`. If so, the local variable at index $n - 1$ is changed to indicate that it now contains an unusable value. Since the local variable at index n has been overwritten, the local variable at index $n - 1$ cannot represent a value of type `long` or `double`.

Dealing with values of types `long` or `double` on the operand stack is simpler; the verifier treats them as single values on the stack. For example, the verification code for the *dadd* opcode (add two `double` values) checks that the top two items on the stack are both of type `double`. When calculating operand stack length, values of type `long` and `double` have length two.

Untyped instructions that manipulate the operand stack must treat values of type `double` and `long` as atomic (indivisible). For example, the verifier reports a failure if the top value on the stack is a `double` and it encounters an instruction such as *pop* or *dup*. The instructions *pop2* or *dup2* must be used instead.

4.11.2.4 Instance Initialization Methods and Newly Created Objects

Creating a new class instance is a multistep process. The statement

```
...
new myClass(i, j, k);
...
```

can be implemented by the following:

```
...
new #1           // Allocate uninitialized space for myClass
dup              // Duplicate object on the operand stack
iload_1          // Push i
iload_2          // Push j
iload_3          // Push k
invokespecial #5 // Invoke myClass.<init>
...
```

This instruction sequence leaves the newly created and initialized object on top of the operand stack. (Additional examples of compilation to the instruction set of the Java virtual machine are given in Chapter 77, ~~“Compiling for the Java Virtual Machine.”~~[“Compiling for the Java Virtual Machine.”](#))

The instance initialization method (§3.9)(§3.9) for class `myClass` sees the new uninitialized object as its `this` argument in local variable 0. Before that method invokes another instance initialization method of `myClass` or its direct superclass

on `this`, the only operation the method can perform on `this` is assigning fields declared within `myClass`.

When doing dataflow analysis on instance methods, the verifier initializes local variable 0 to contain an object of the current class, or, for instance initialization methods, local variable 0 contains a special type indicating an uninitialized object. After an appropriate instance initialization method is invoked (from the current class or the current superclass) on this object, all occurrences of this special type on the verifier's model of the operand stack and in the local variable array are replaced by the current class type. The verifier rejects code that uses the new object before it has been initialized or that initializes the object more than once. In addition, it ensures that every normal return of the method has invoked an instance initialization method either in the class of this method or in the direct superclass.

Similarly, a special type is created and pushed on the verifier's model of the operand stack as the result of the Java virtual machine instruction *new*. The special type indicates the instruction by which the class instance was created and the type of the uninitialized class instance created. When an instance initialization method is invoked on that class instance, all occurrences of the special type are replaced by the intended type of the class instance. This change in type may propagate to subsequent instructions as the dataflow analysis proceeds.

The instruction number needs to be stored as part of the special type, as there may be multiple not-yet-initialized instances of a class in existence on the operand stack at one time. For example, the Java virtual machine instruction sequence that implements

```
new InputStream(new Foo(), new InputStream("foo"))
```

may have two uninitialized instances of `InputStream` on the operand stack at once. When an instance initialization method is invoked on a class instance, only those occurrences of the special type on the operand stack or in the local variable array that are the *same object* as the class instance are replaced.

A valid instruction sequence must not have an uninitialized object on the operand stack or in a local variable during a backwards branch, or in a local variable in code protected by an exception handler or a `finally` clause. Otherwise, a devious piece of code might fool the verifier into thinking it had initialized a class instance when it had, in fact, initialized a class instance created in a previous pass through a loop.

4.11.2.5 Exception Handlers

Java virtual machine code produced by Sun's compiler for the Java programming language always generates exception handlers such that:

- Either the ranges of instructions protected by two different exception handlers always are completely disjoint, or else one is a subrange of the other. There is never a partial overlap of ranges.
- The handler for an exception will never be inside the code that is being protected.
- The only entry to an exception handler is through an exception. It is impossible to fall through or “goto” the exception handler.

These restrictions are not enforced by the `class` file verifier since they do not pose a threat to the integrity of the Java virtual machine. As long as every nonexceptional path to the exception handler causes there to be a single object on the operand stack, and as long as all other criteria of the verifier are met, the verifier will pass the code.

4.12 Limitations of the Java Virtual Machine

The following limitations of the Java virtual machine are implicit in the `class` file format:

- The per-class or per-interface constant pool is limited to 65535 entries by the 16-bit `constant_pool_count` field of the `ClassFile` structure (§4.2). This acts as an internal limit on the total complexity of a single class or interface.
- The greatest number of local variables in the local variables array of a frame created upon invocation of a method is limited to 65535 by the size of the `max_locals` item of the `Code` attribute (§4.8.3) giving the code of the method, and by the 16-bit local variable indexing of the Java virtual machine instruction set. Note that values of type `long` and `double` are each considered to reserve two local variables and contribute two units toward the `max_locals` value, so use of local variables of those types further reduces this limit.
- The number of fields that may be declared by a class or interface is limited to 65535 by the size of the `fields_count` item of the `ClassFile` structure (§4.2). Note that the value of the `fields_count` item of the `ClassFile` structure does not include fields that are inherited from superclasses or superinterfaces.
- The number of methods that may be declared by a class or interface is limited to 65535 by the size of the `methods_count` item of the `ClassFile` structure

(§4.2). Note that the value of the `methods_count` item of the `ClassFile` structure does not include methods that are inherited from superclasses or superinterfaces.

- The number of direct superinterfaces of a class or interface is limited to 65535 by the size of the `interfaces_count` item of the `ClassFile` structure (§4.2).
- The size of an operand stack in a frame (§3.6)(§3.6) is limited to 65535 values by the `max_stack` field of the `Code` attribute (§4.8.3). Note that values of type `long` and `double` are each considered to contribute two units toward the `max_stack` value, so use of values of these types on the operand stack further reduces this limit.
- The number of dimensions in an array is limited to 255 by the size of the *dimensions* opcode of the *multianewarray* instruction and by the constraints imposed on the *multianewarray*, *anewarray*, and *newarray* instructions by §4.10.2.
- The number of method parameters is limited to 255 by the definition of a method descriptor (§4.4.3), where the limit includes one unit for `this` in the case of instance or interface method invocations. Note that a method descriptor is defined in terms of a notion of method parameter length in which a parameter of type `long` or `double` contributes two units to the length, so parameters of these types further reduce the limit.
- The length of field and method names, field and method descriptors, and other constant string values is limited to 65535 characters by the 16-bit unsigned `length` item of the `CONSTANT_Utf8_info` structure (§4.5.7). Note that the limit is on the number of bytes in the encoding and not on the number of encoded characters. UTF-8 encodes some characters using two or three bytes. Thus, strings incorporating multibyte characters are further constrained.